

Applying RNNs to Character Prediction for Text Generation



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Language modeling is an important area of ML research

Text prediction is one of several classic language modeling problems

Multi-RNNs are a specific type of RNN that work well in language modeling

A key technique is smartly re-initialising state of the RNN during prediction

An evaluation metric called perplexity is used to assess predictive performance

Language Modeling

Two Familiar Problems

Word Embeddings

Express a word in terms of context
in numeric form

Sentiment Analysis

Classify a set of words

Types of Machine Learning Problems




Classification



Regression



Clustering



Rule-extraction

Types of Machine Learning Problems



Classification



Regression



Clustering



Rule-extraction

Sentiment Analysis as Binary Classification



Two Familiar Problems

Word Embeddings

Express a word in terms of context
in numeric form

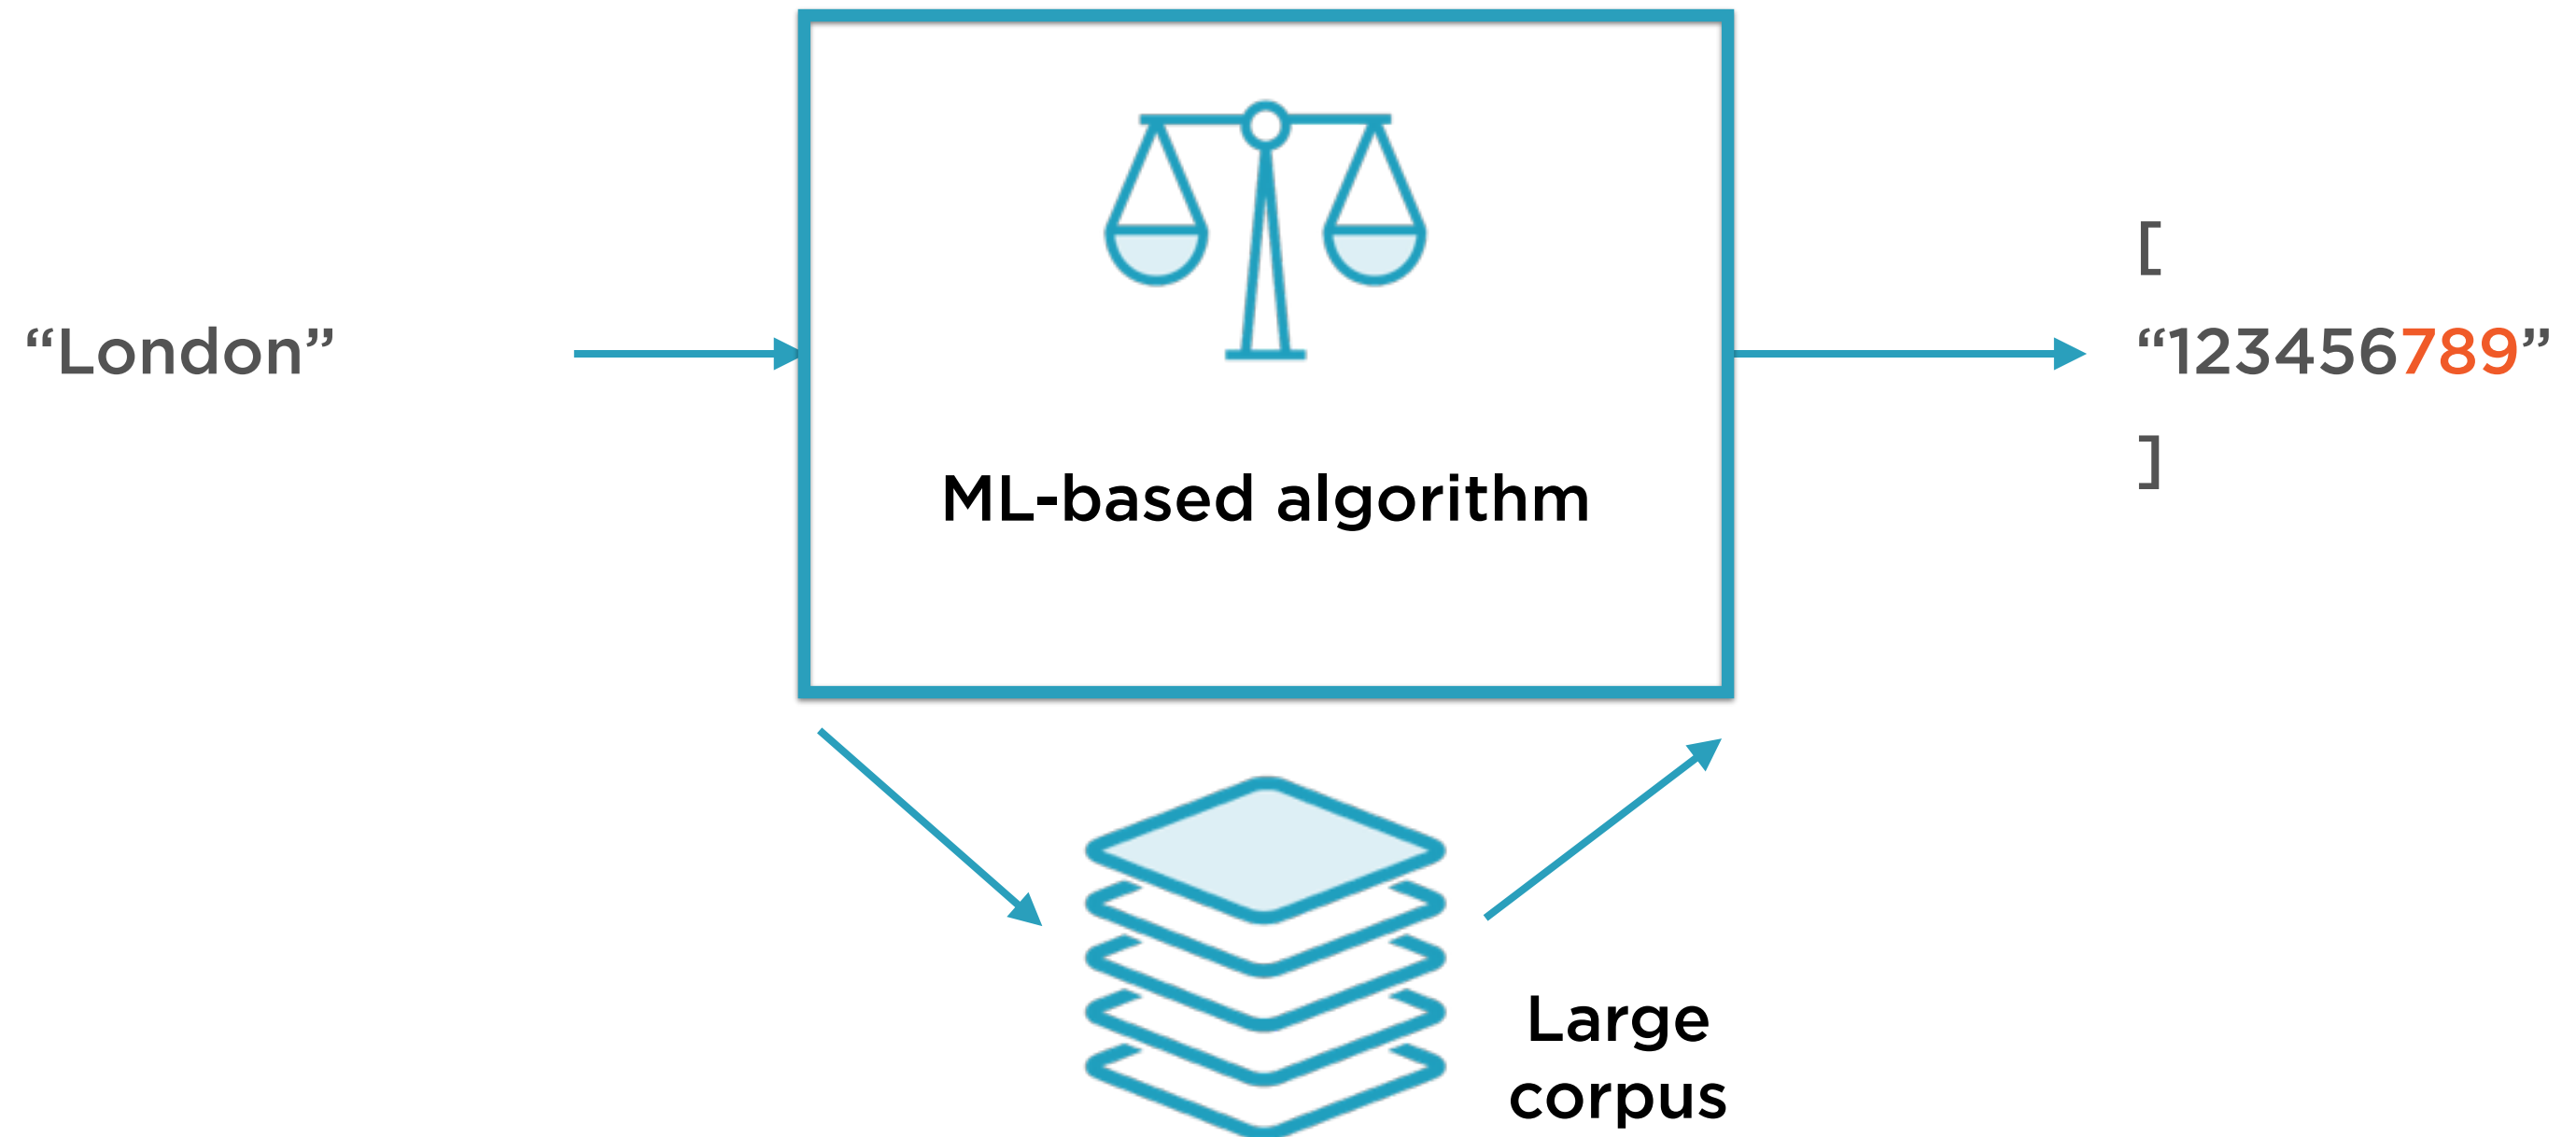
Sentiment Analysis

Classify a set of words

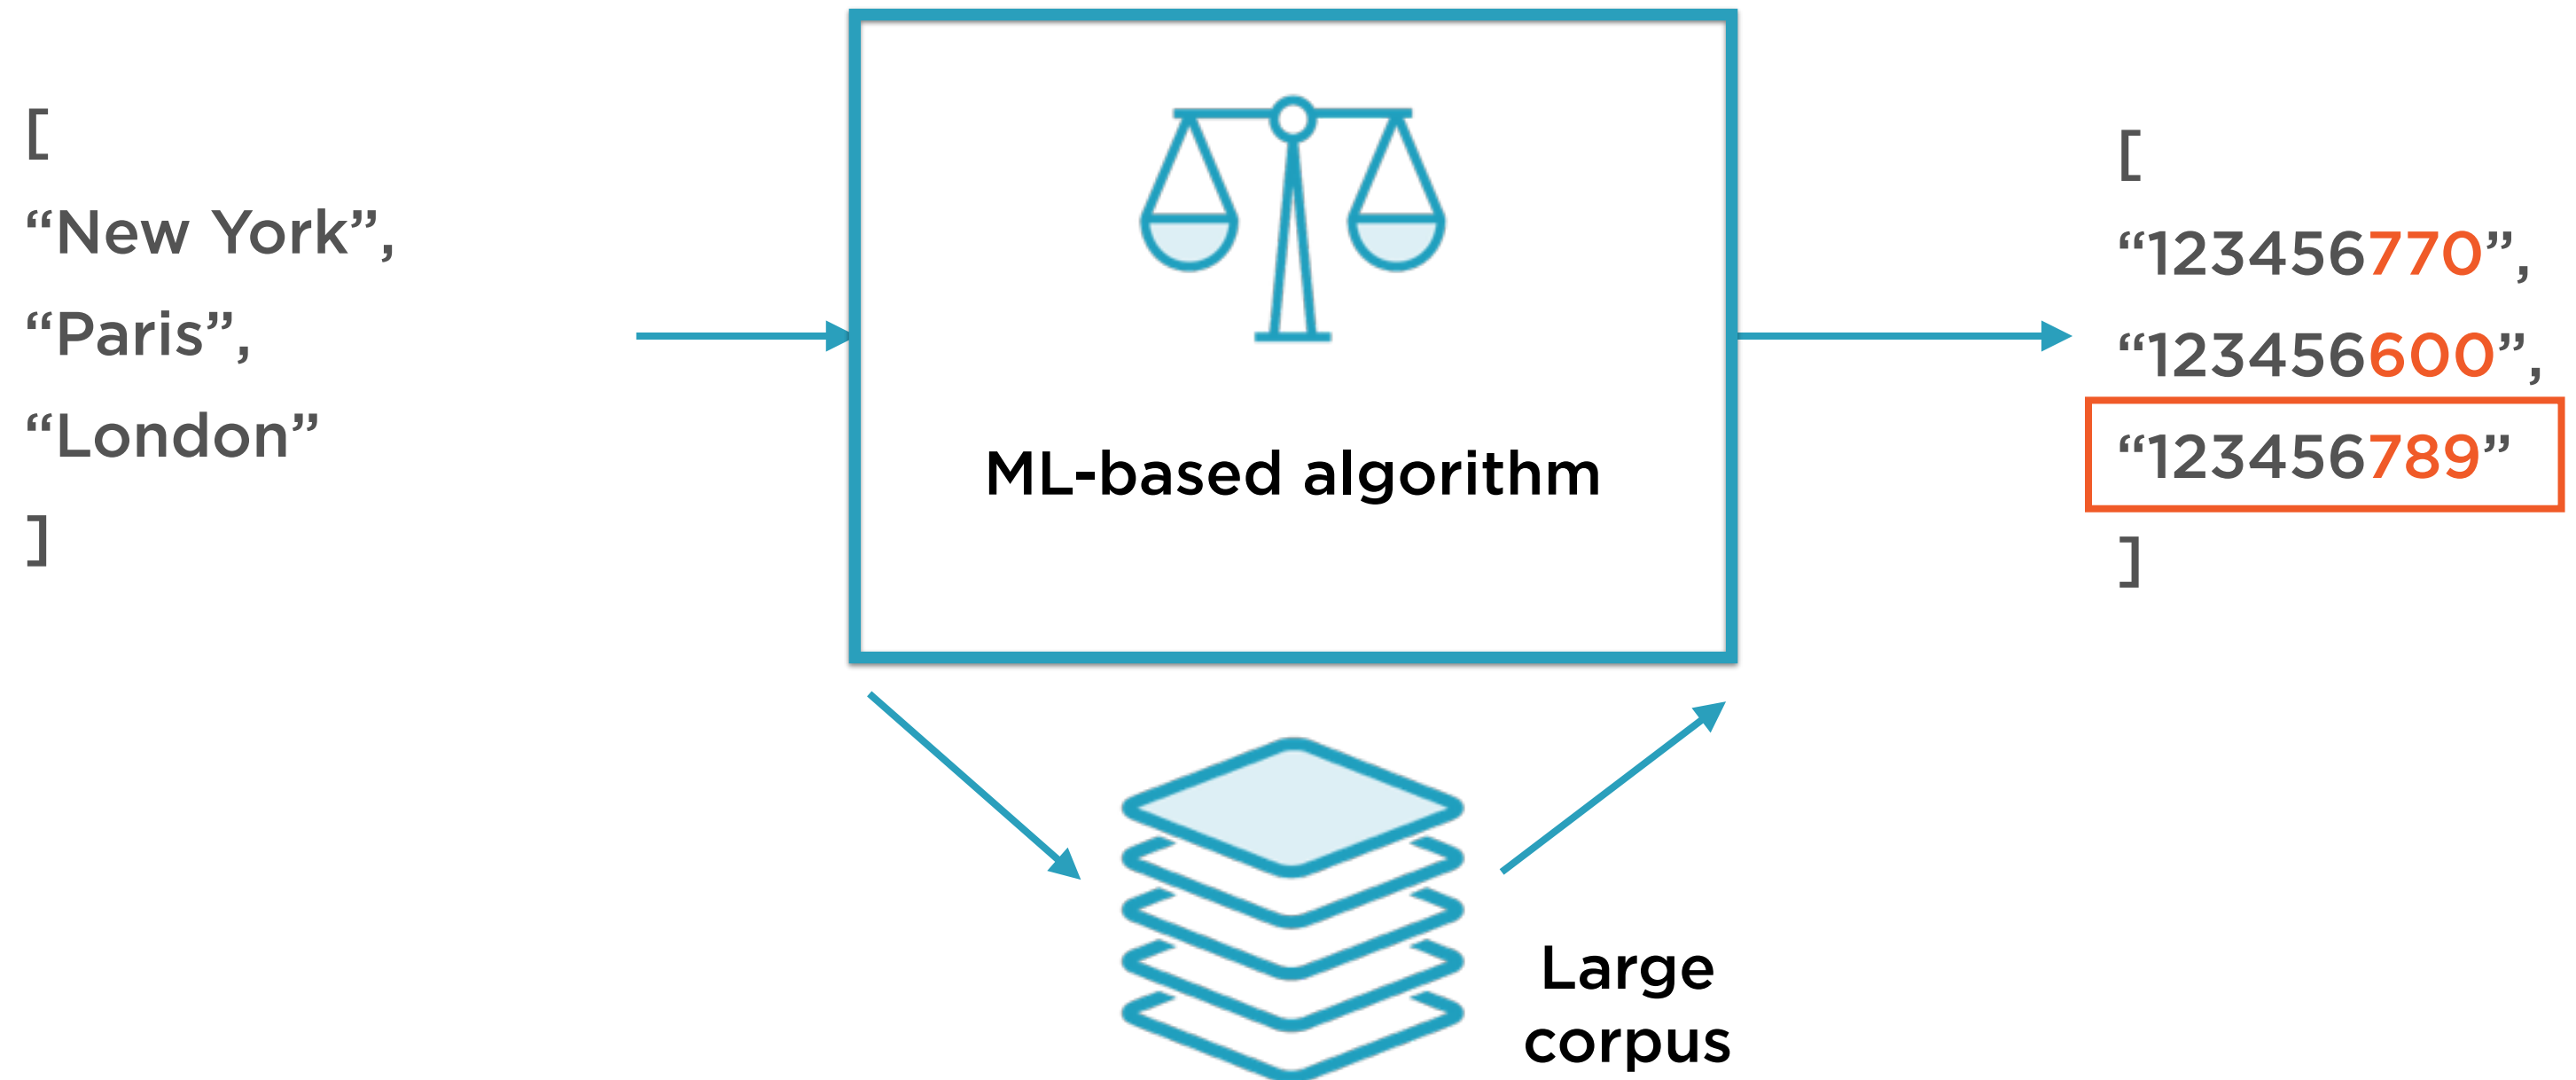
Given words from its context,
predict the word

Given a word, **predict the
words in its context**

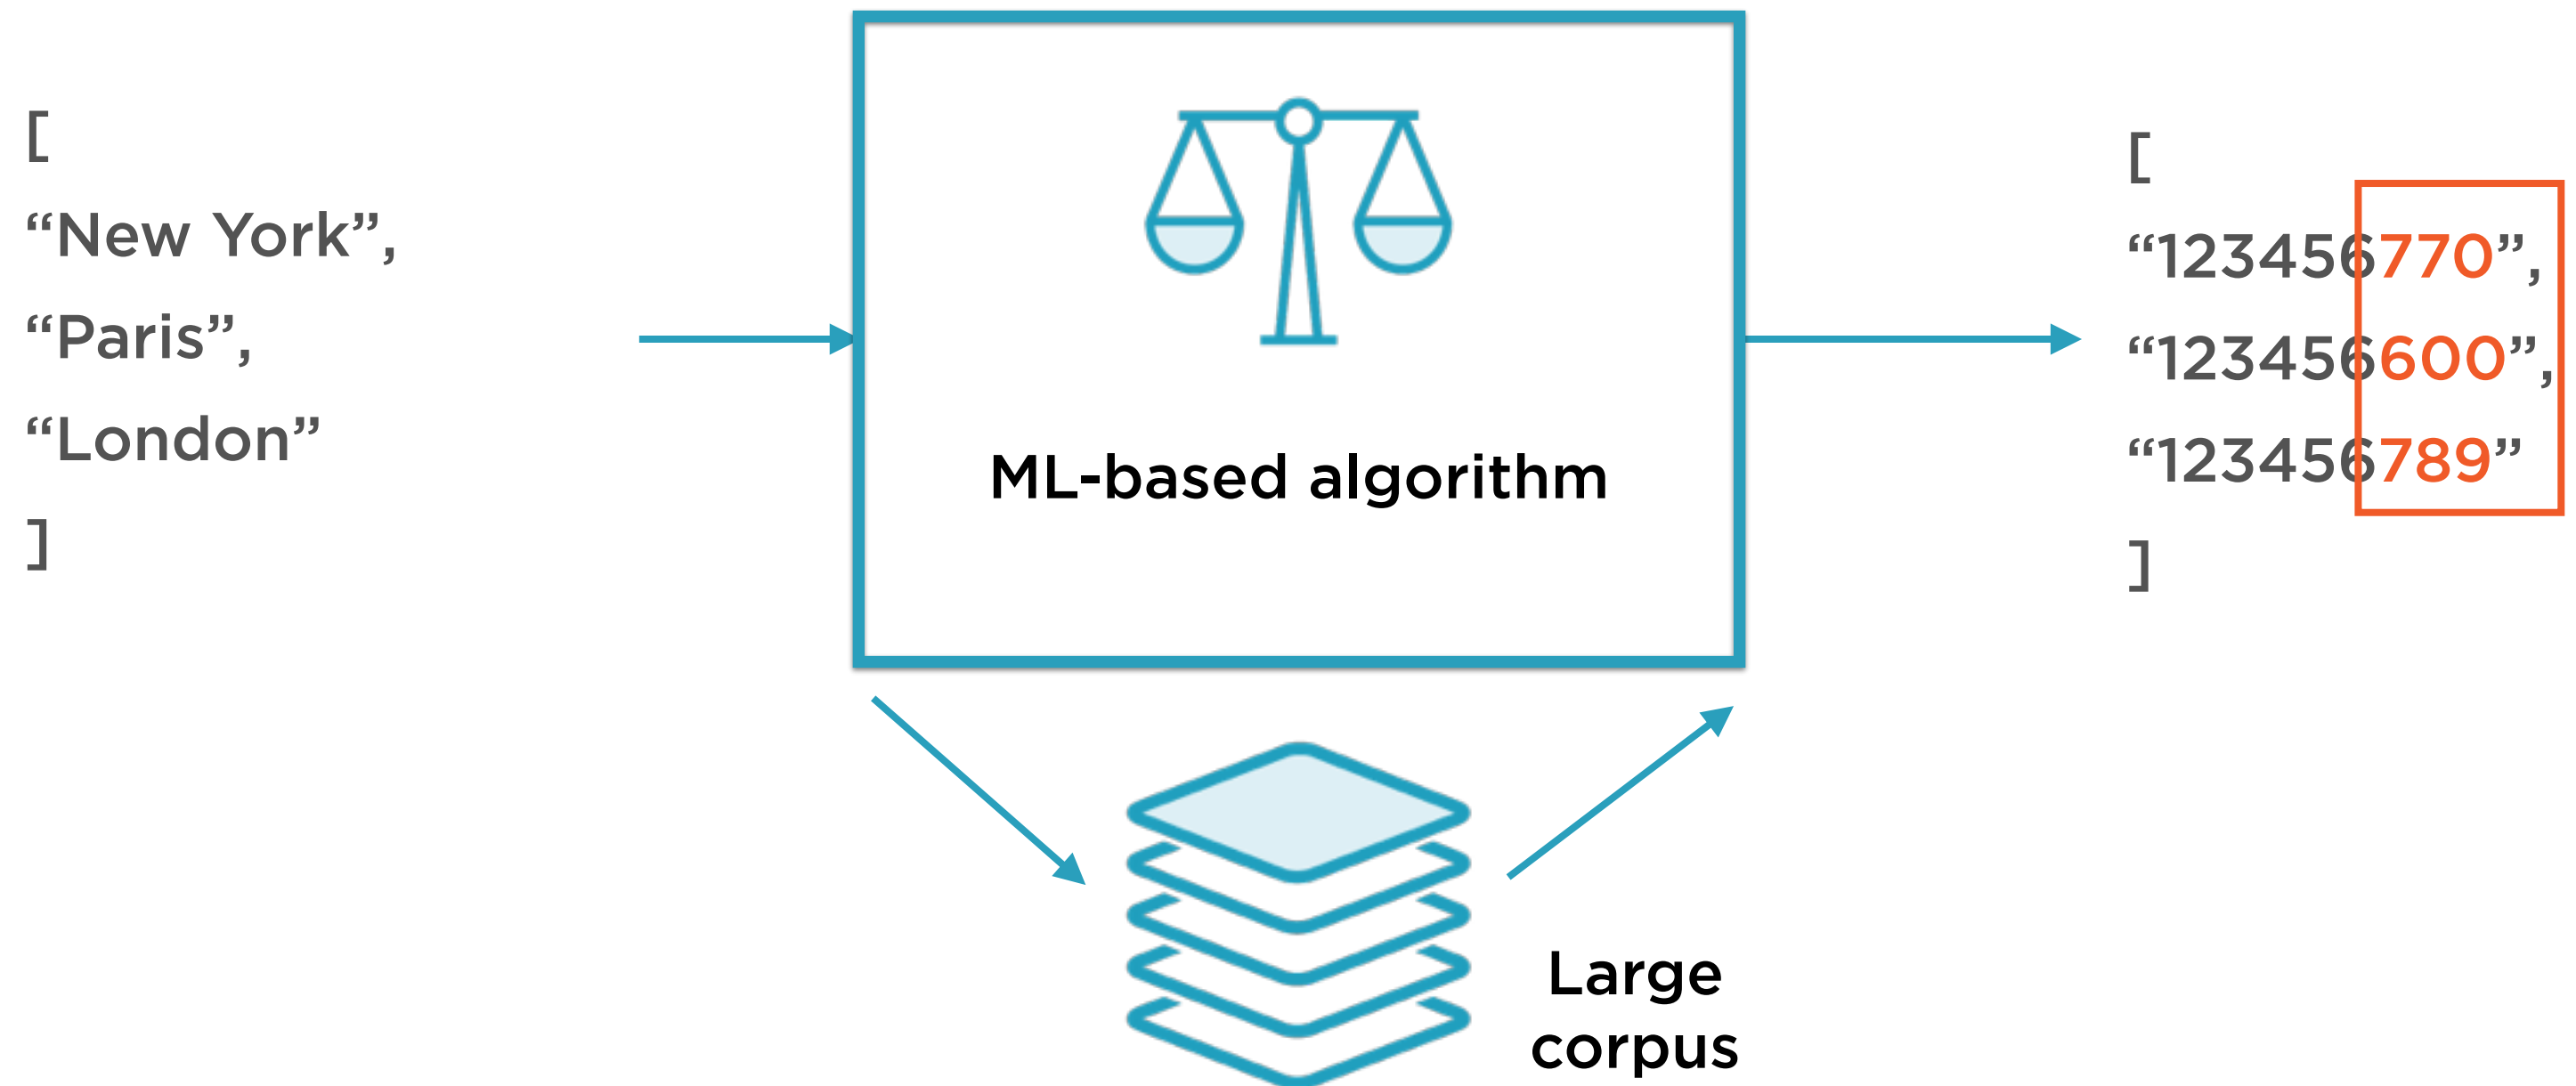
Prediction-based Word Embeddings



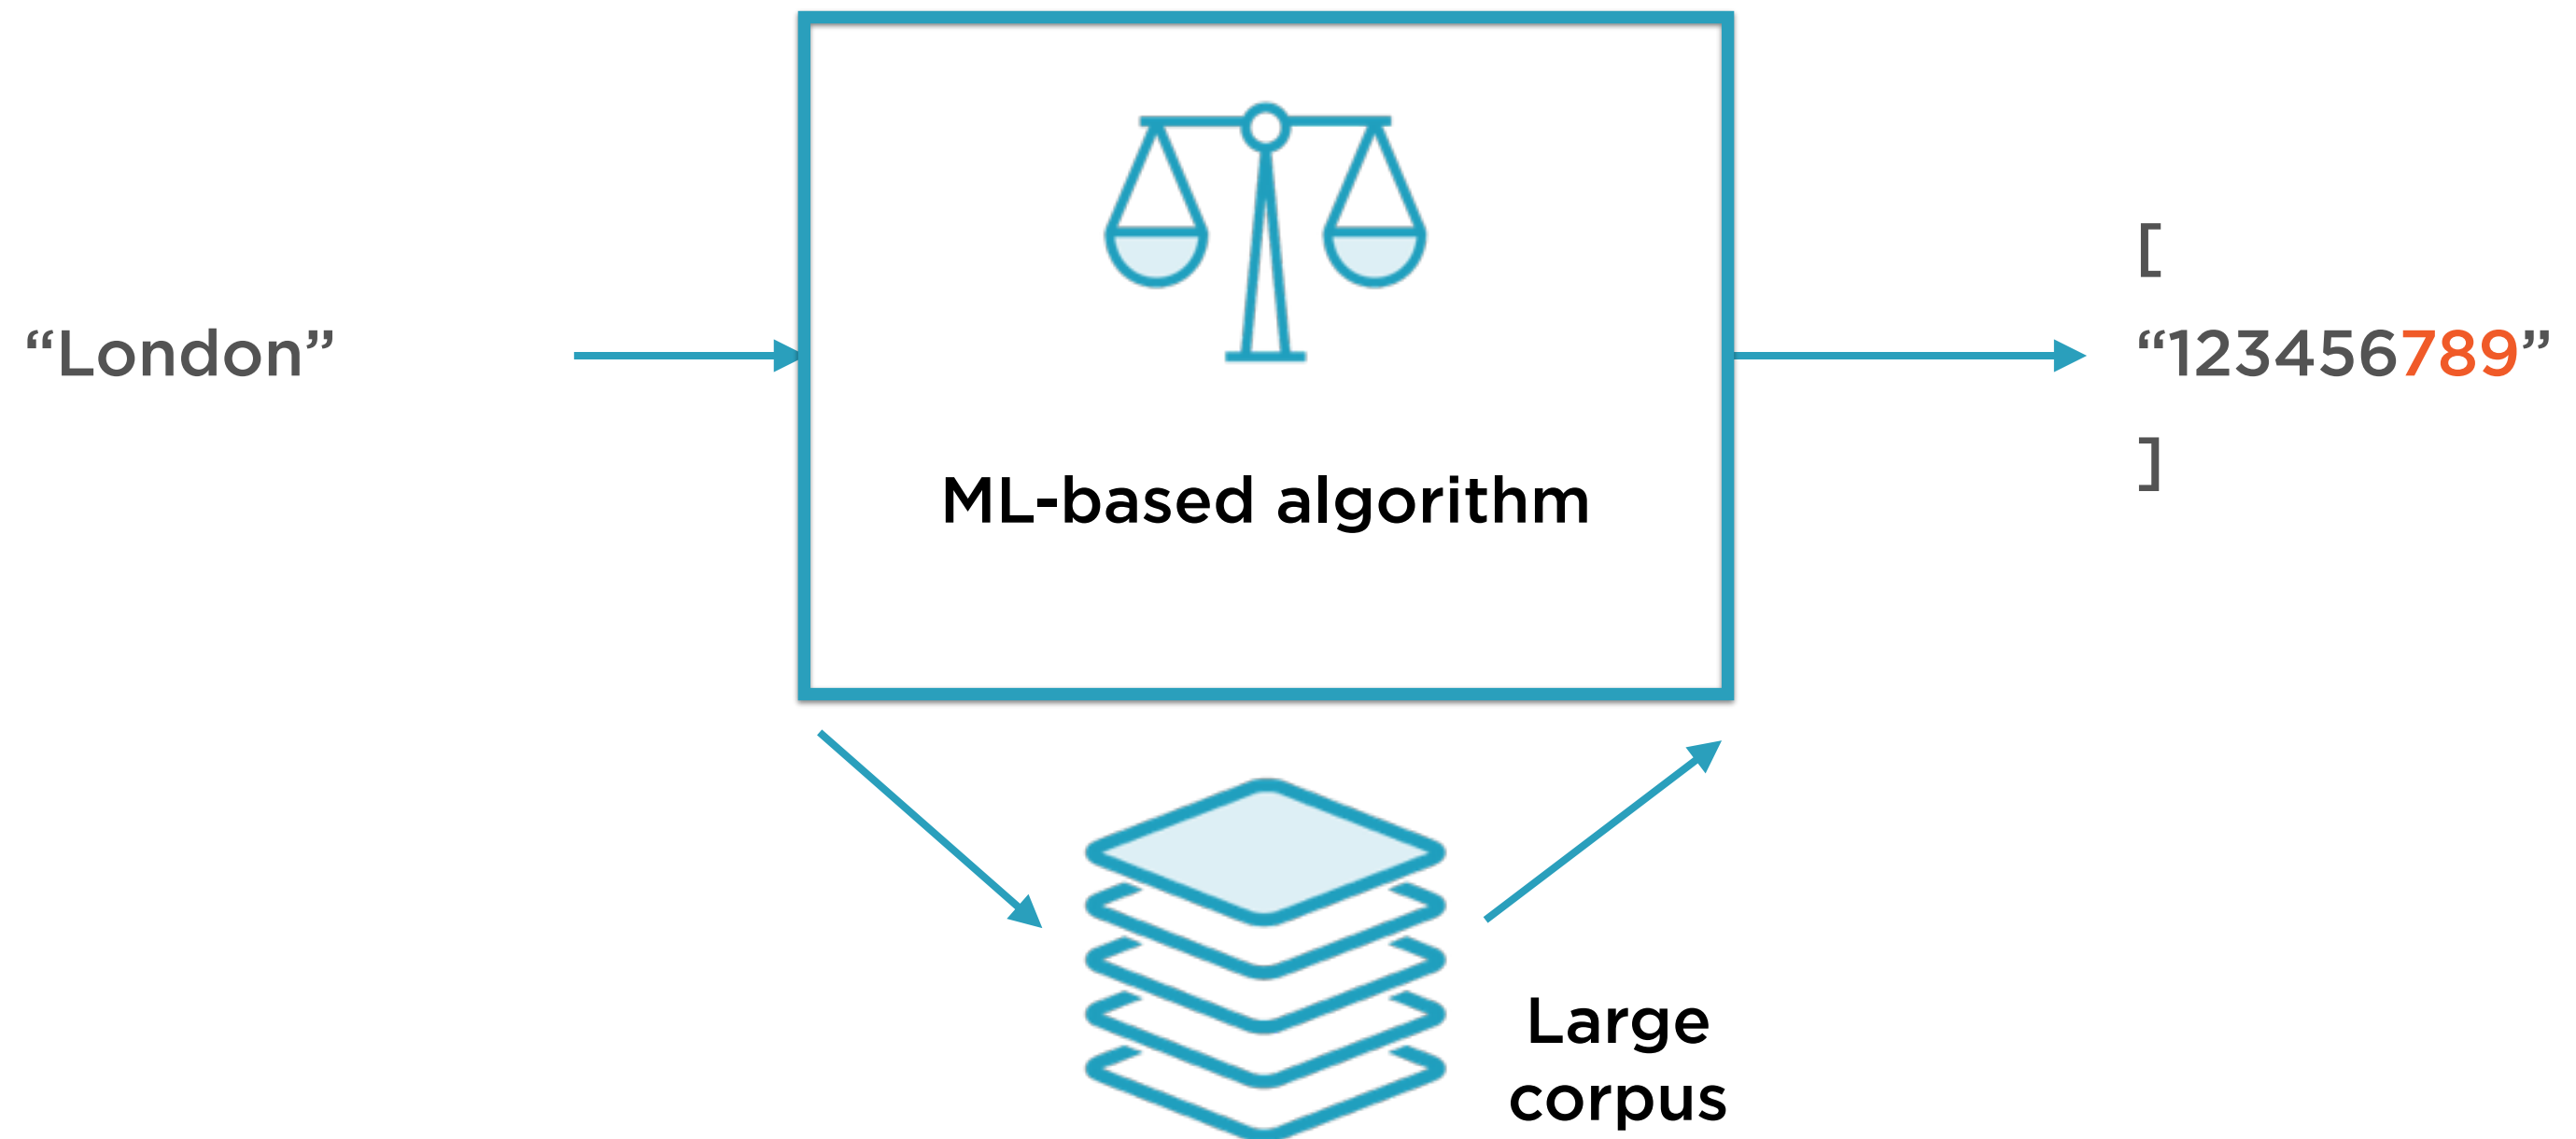
Prediction-based Word Embeddings



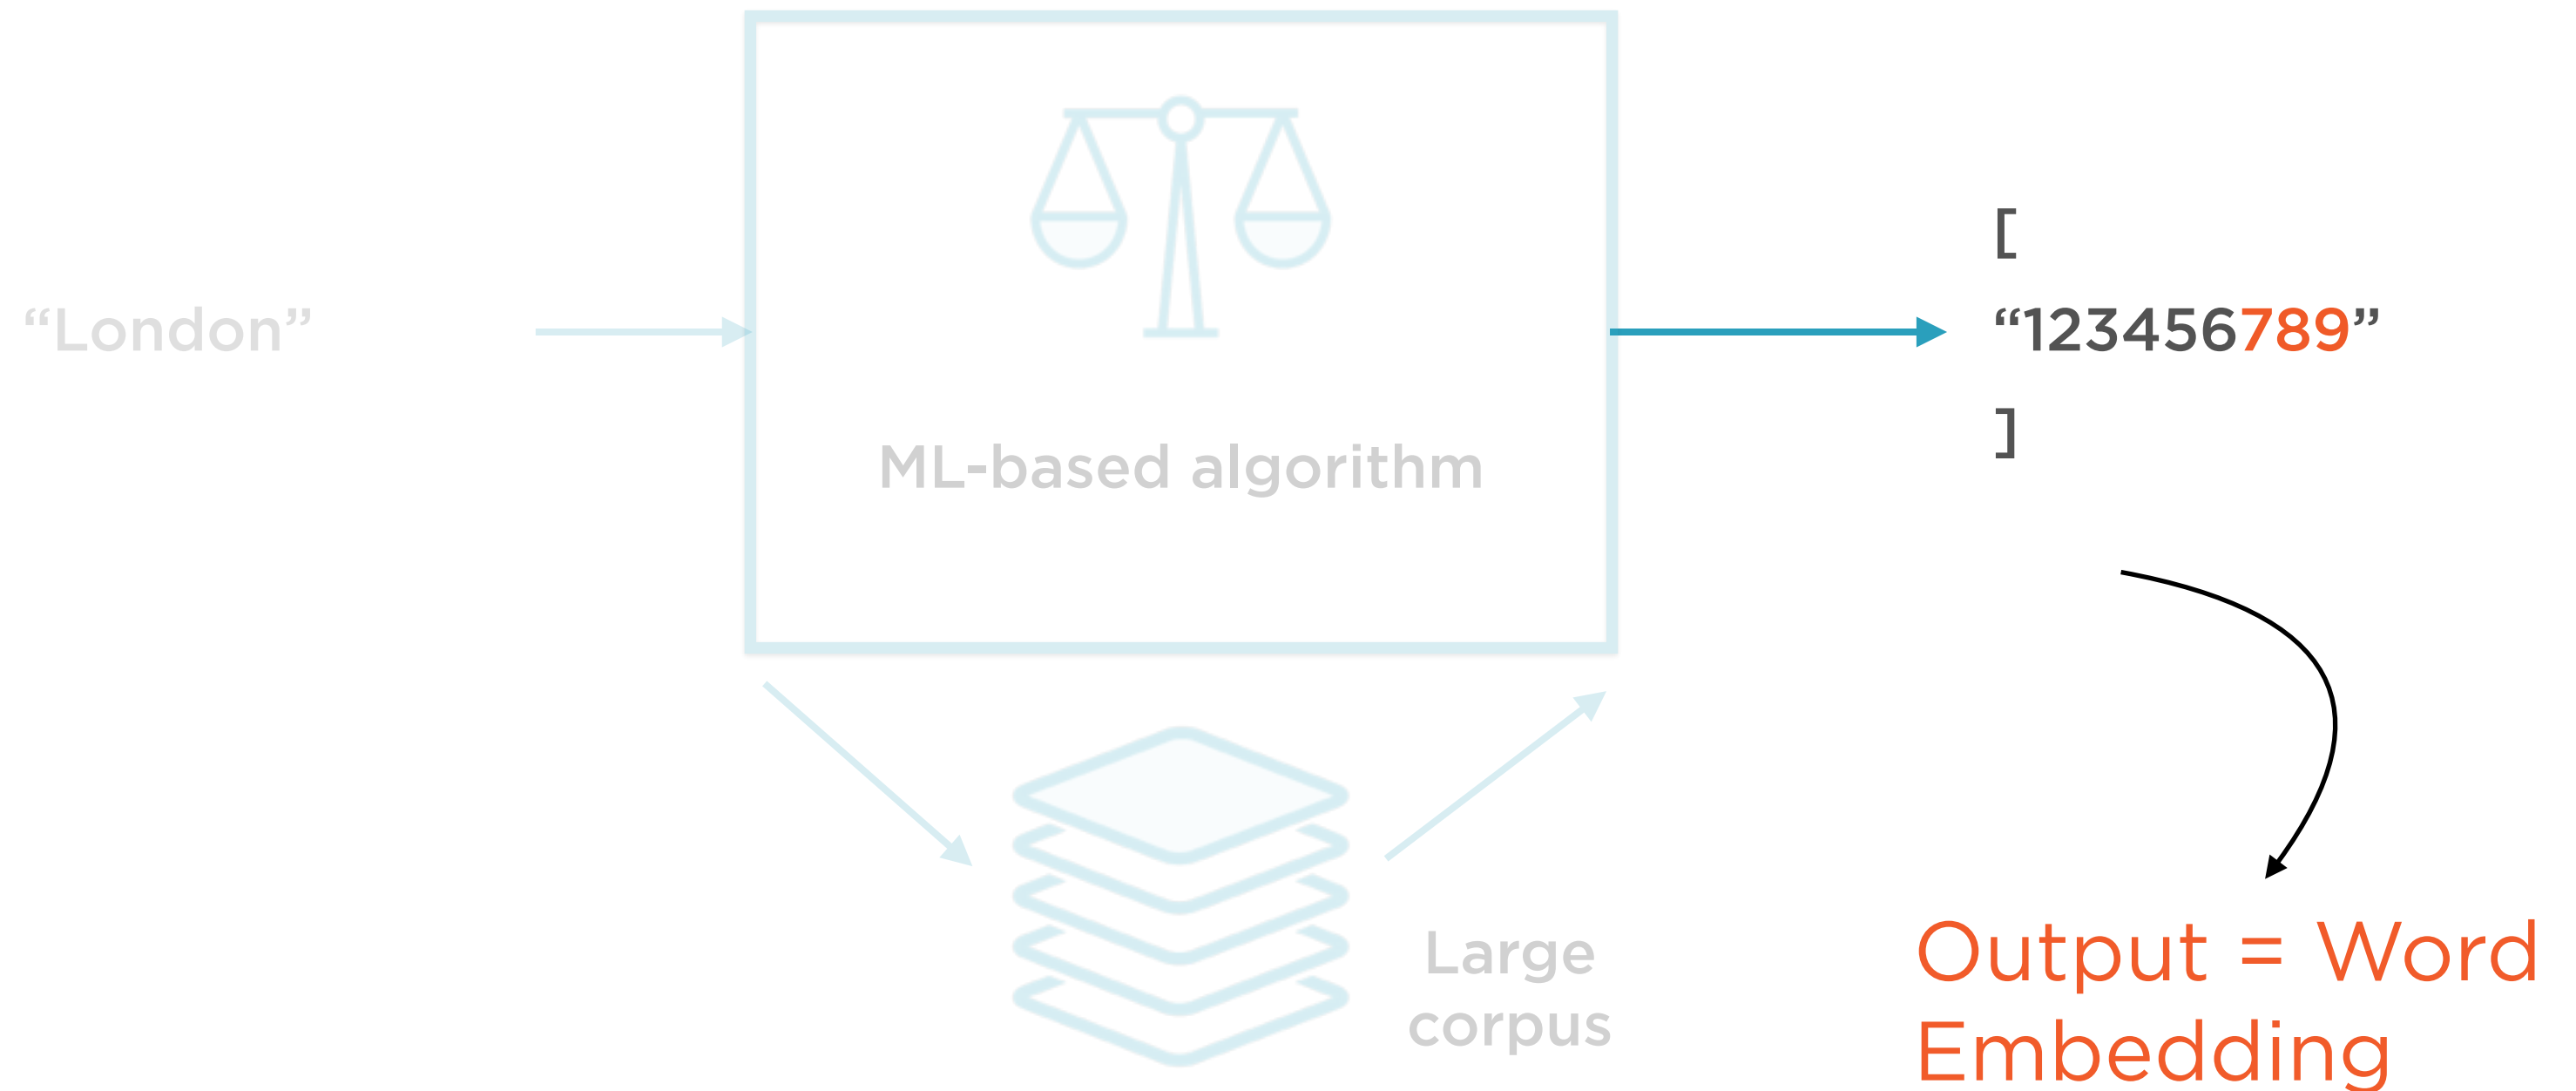
Prediction-based Word Embeddings



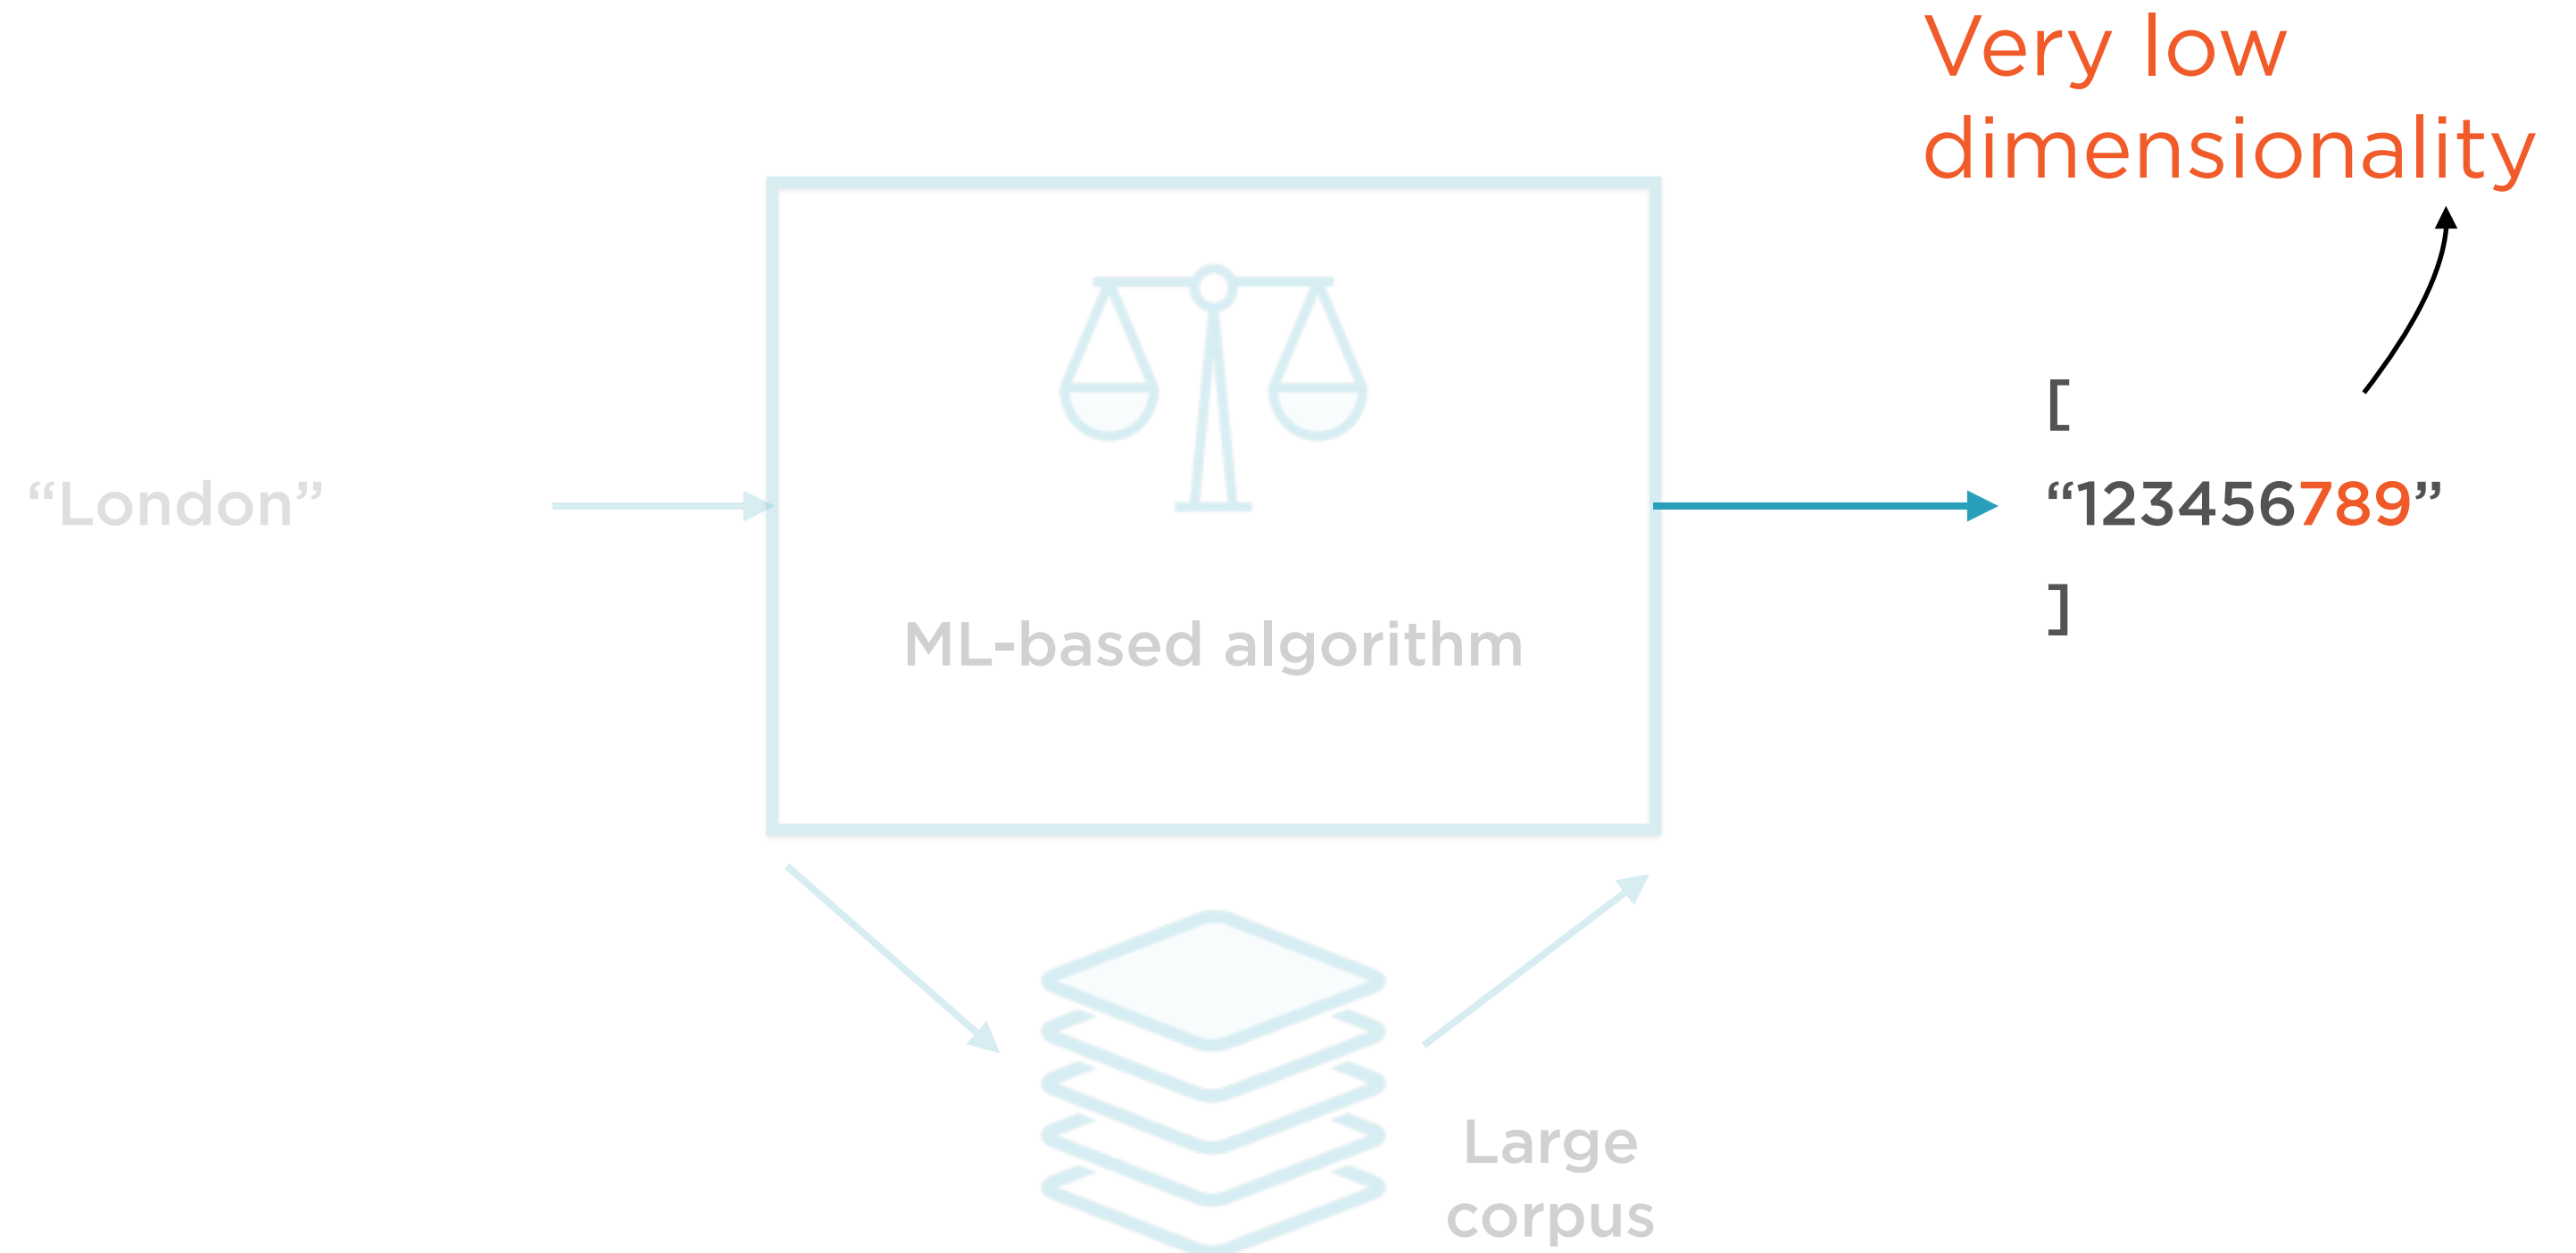
Prediction-based Word Embeddings



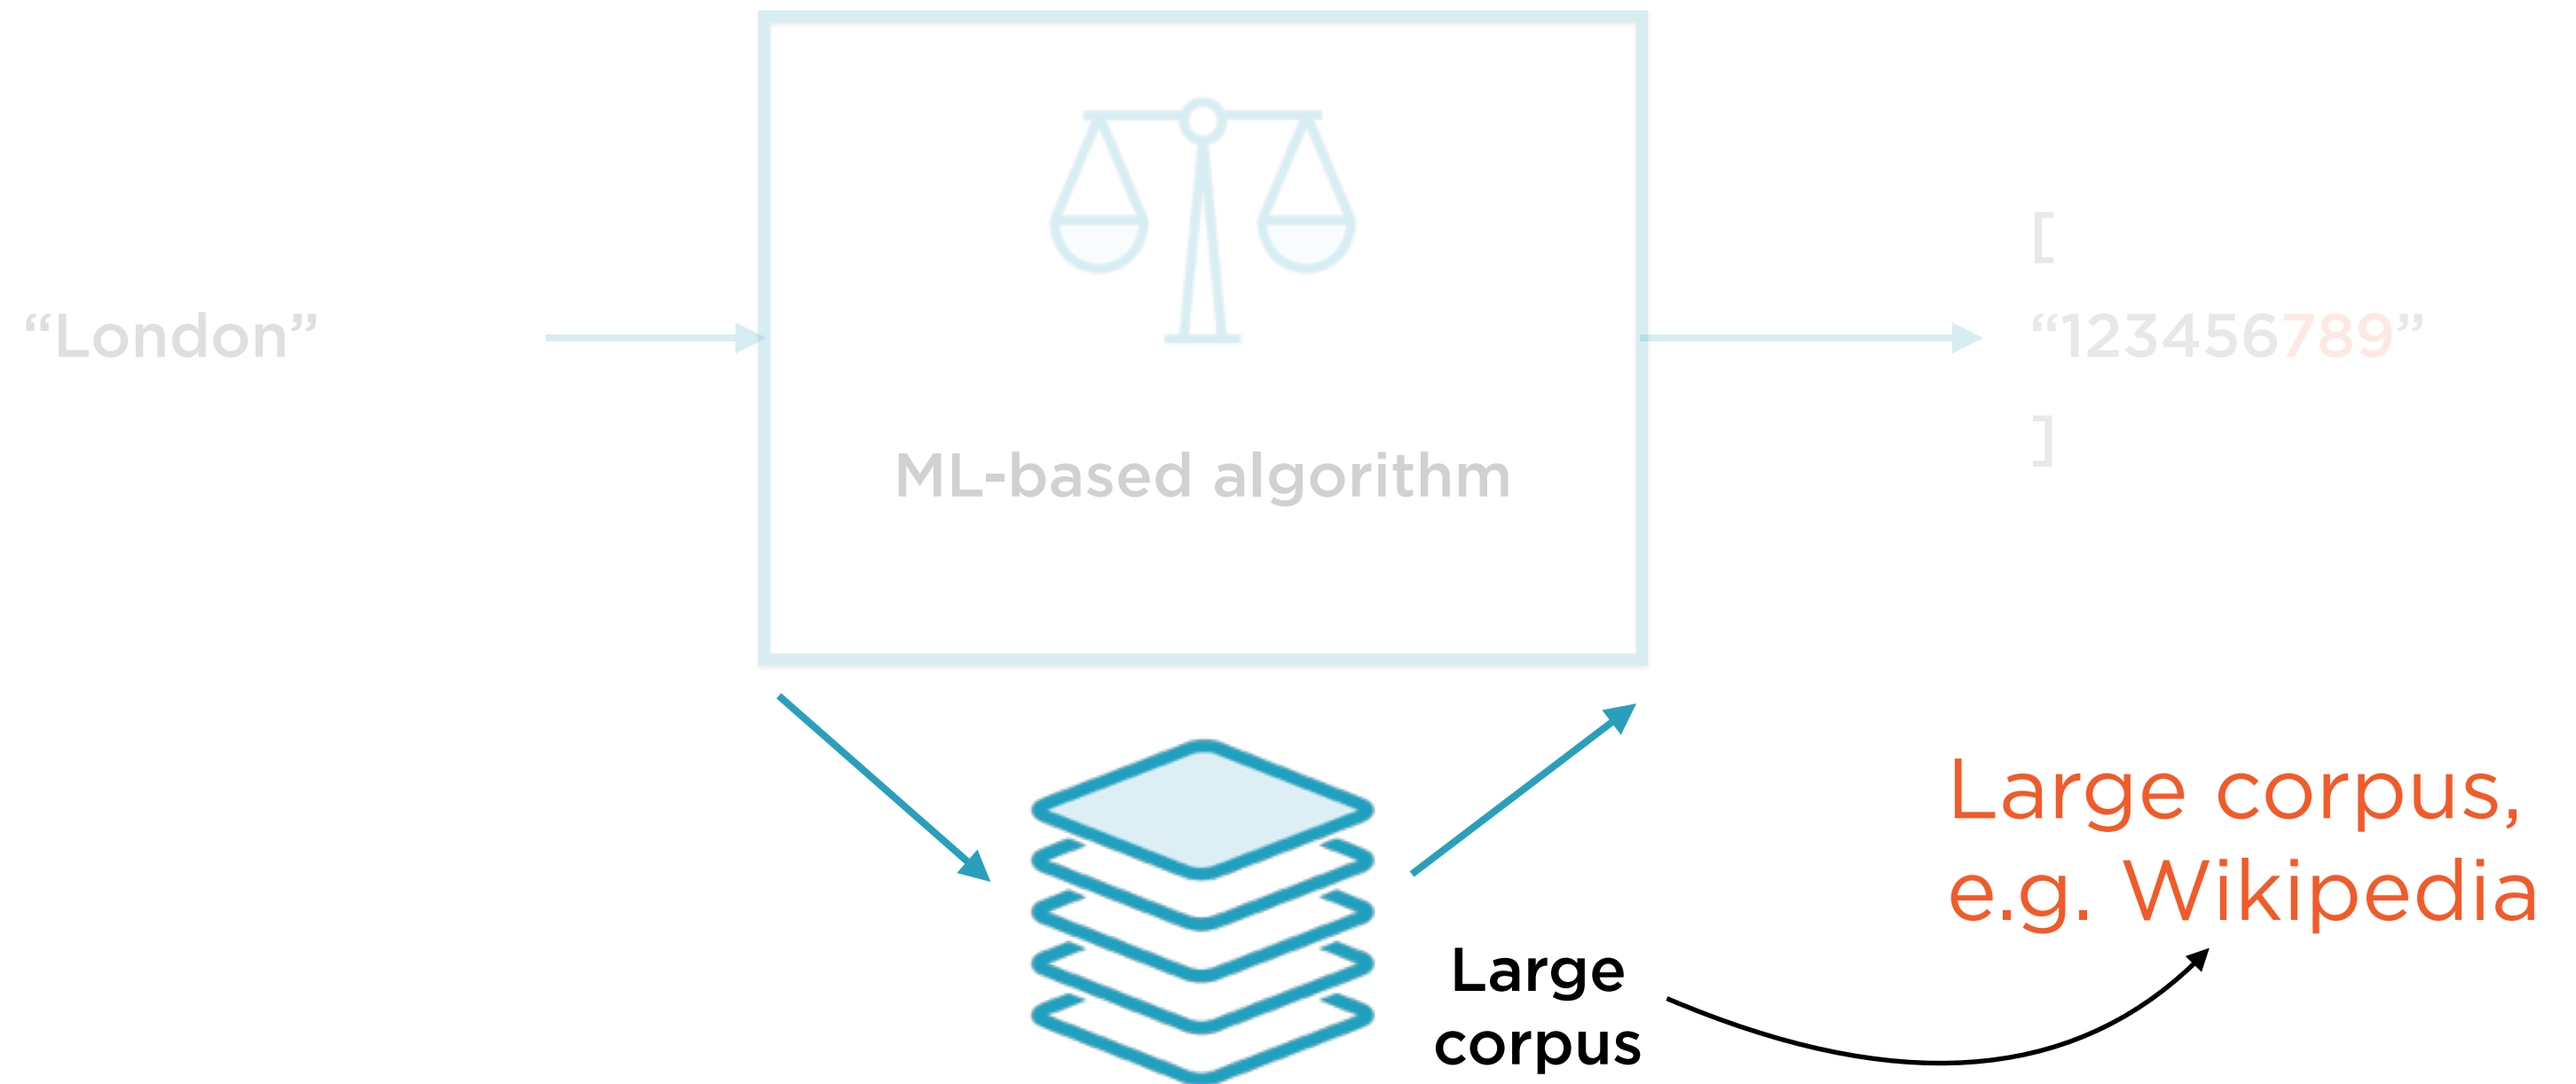
Prediction-based Word Embeddings



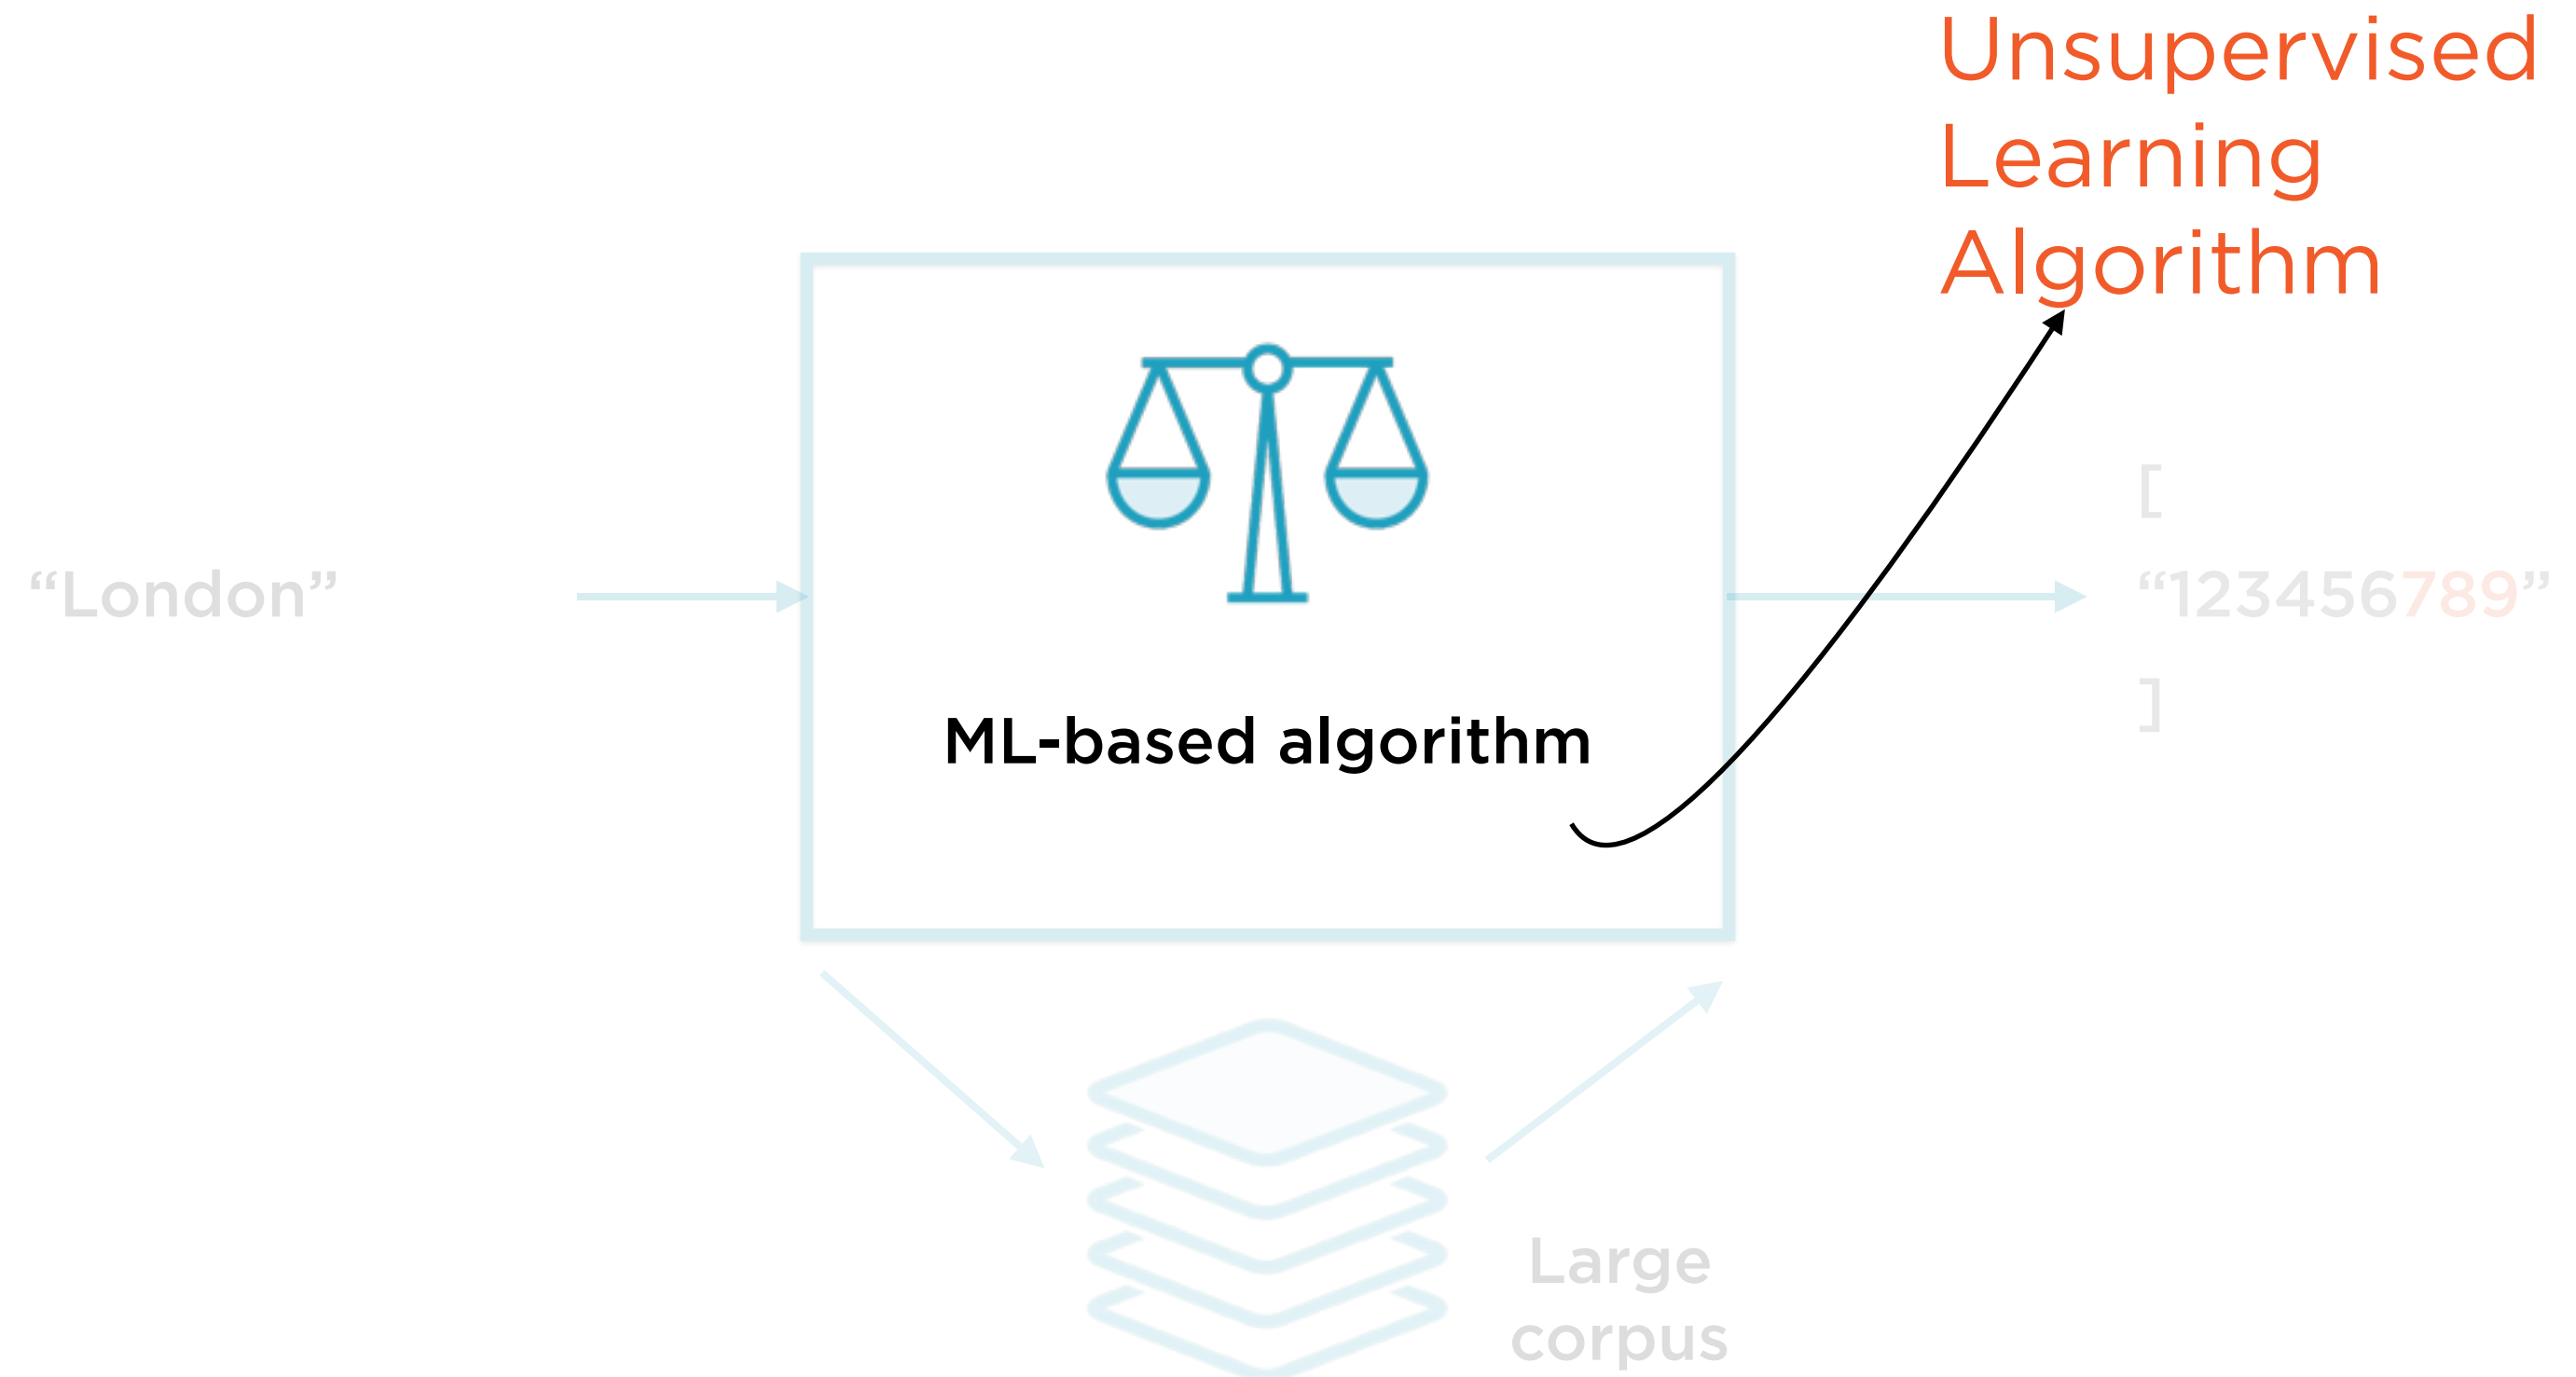
Prediction-based Word Embeddings



Prediction-based Word Embeddings



Prediction-based Word Embeddings





Magic

Word embeddings capture meaning

“Queen” ~ “King” + “Woman” - “Man”

“Paris” ~ “France” + “London” - “England”

Dramatic dimensionality reduction



Word Embeddings as Unsupervised ML

Learnt using ML, often neural networks

Unsupervised deep learning

Pre-processing step before classification

Embeddings are a way to encode words
capturing the **context** around them



Word2Vec

Most popular word embedding model

Mikolov (Google), 2013

Use simple NN (not deep) to learn embeddings



GloVe

Global Vectors for Word Representation

Jeffrey Pennington, Richard Socher, Christopher D. Manning, (Stanford) 2014

Uses word-word co-occurrence matrix, nearest-neighbors for word relationships

Two Familiar Problems

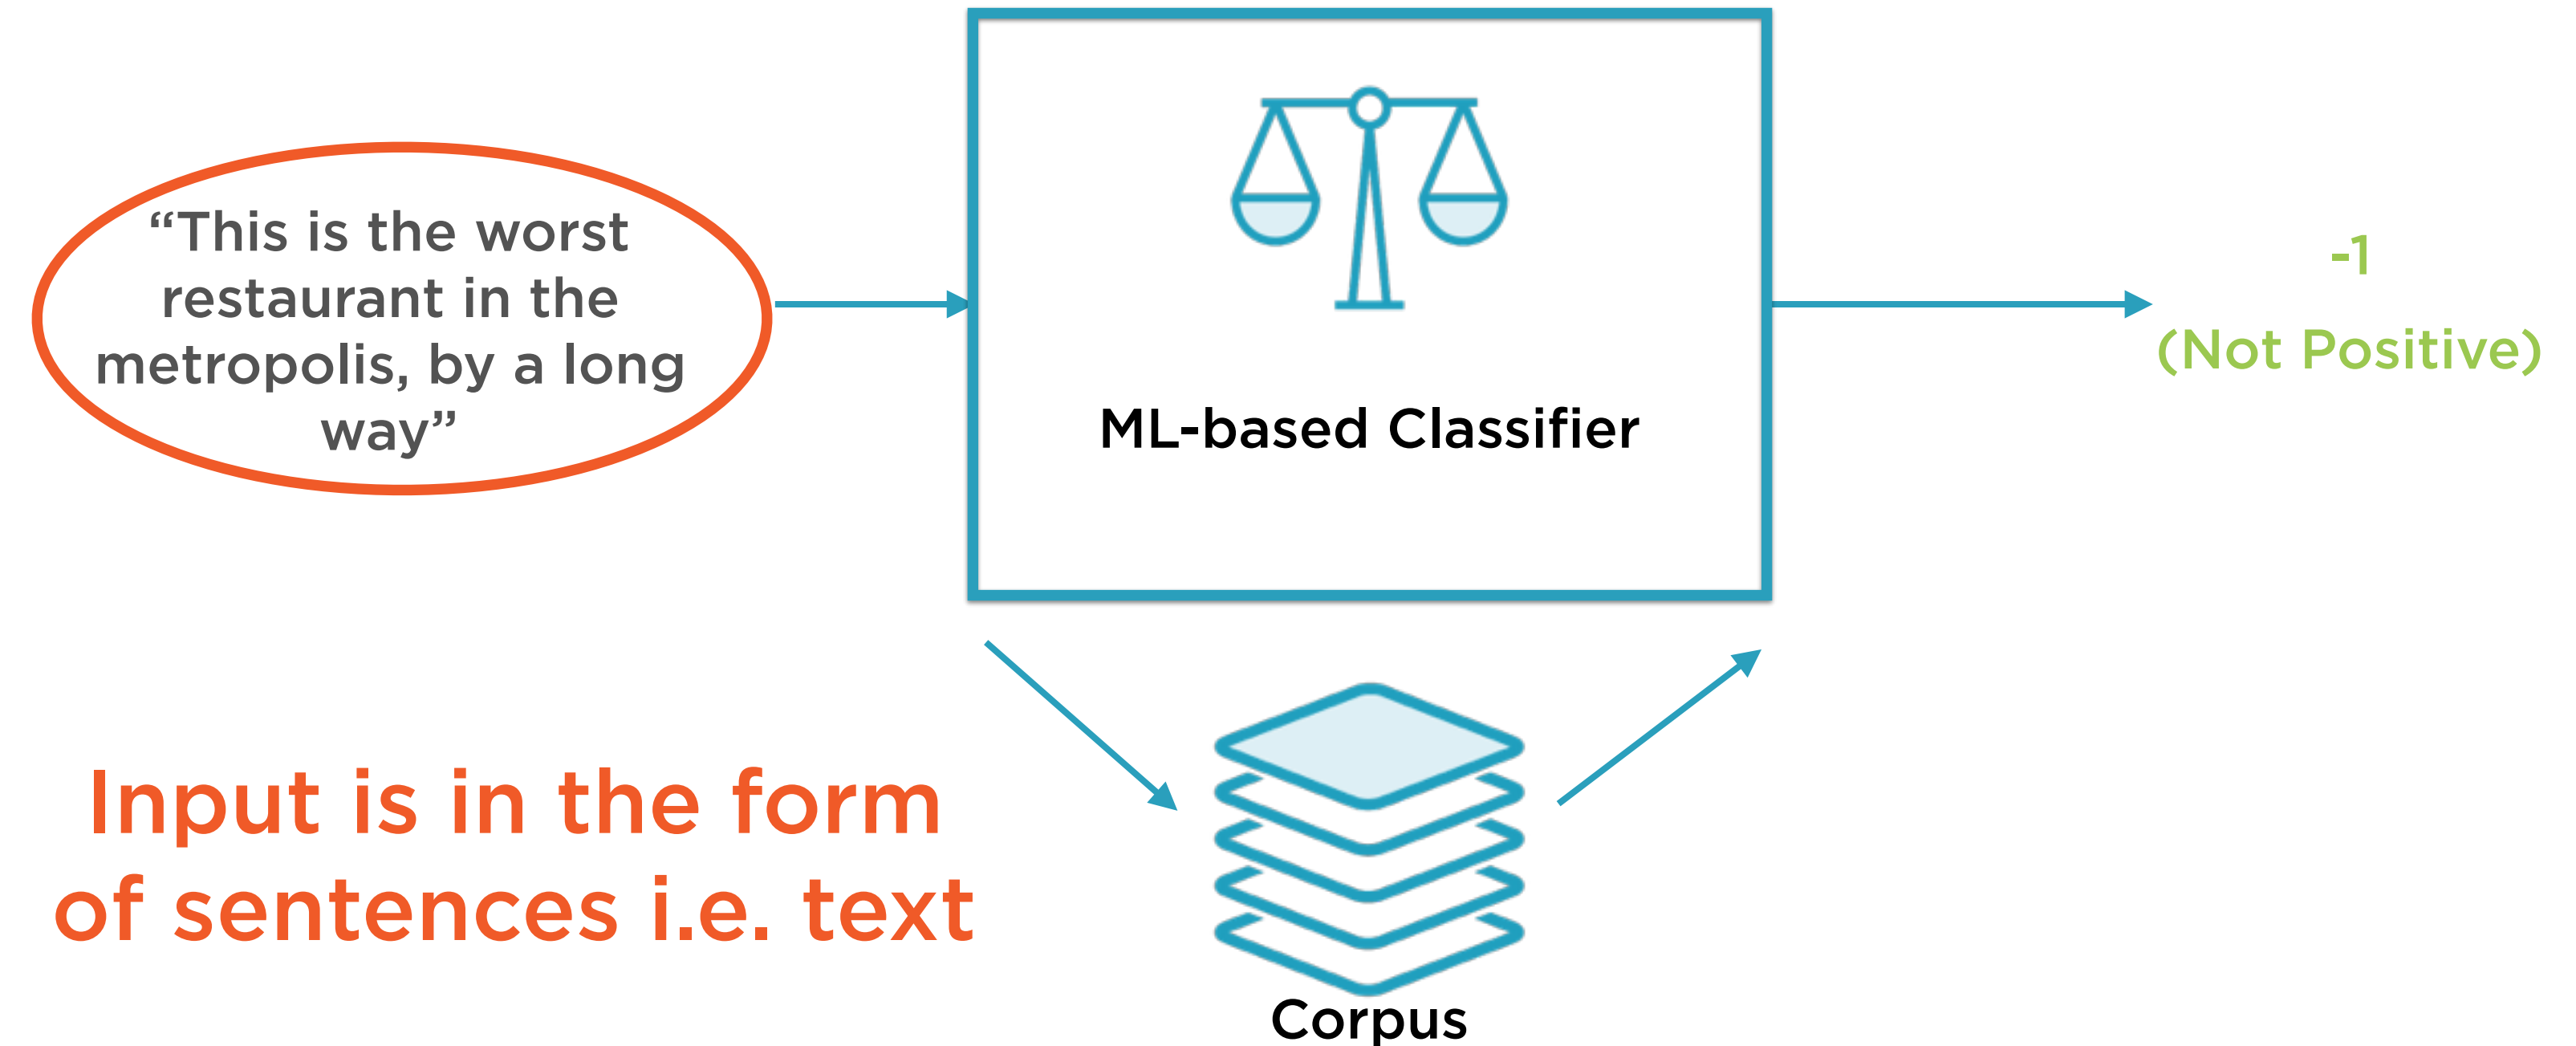
Word Embeddings

Express a word in terms of context
in numeric form

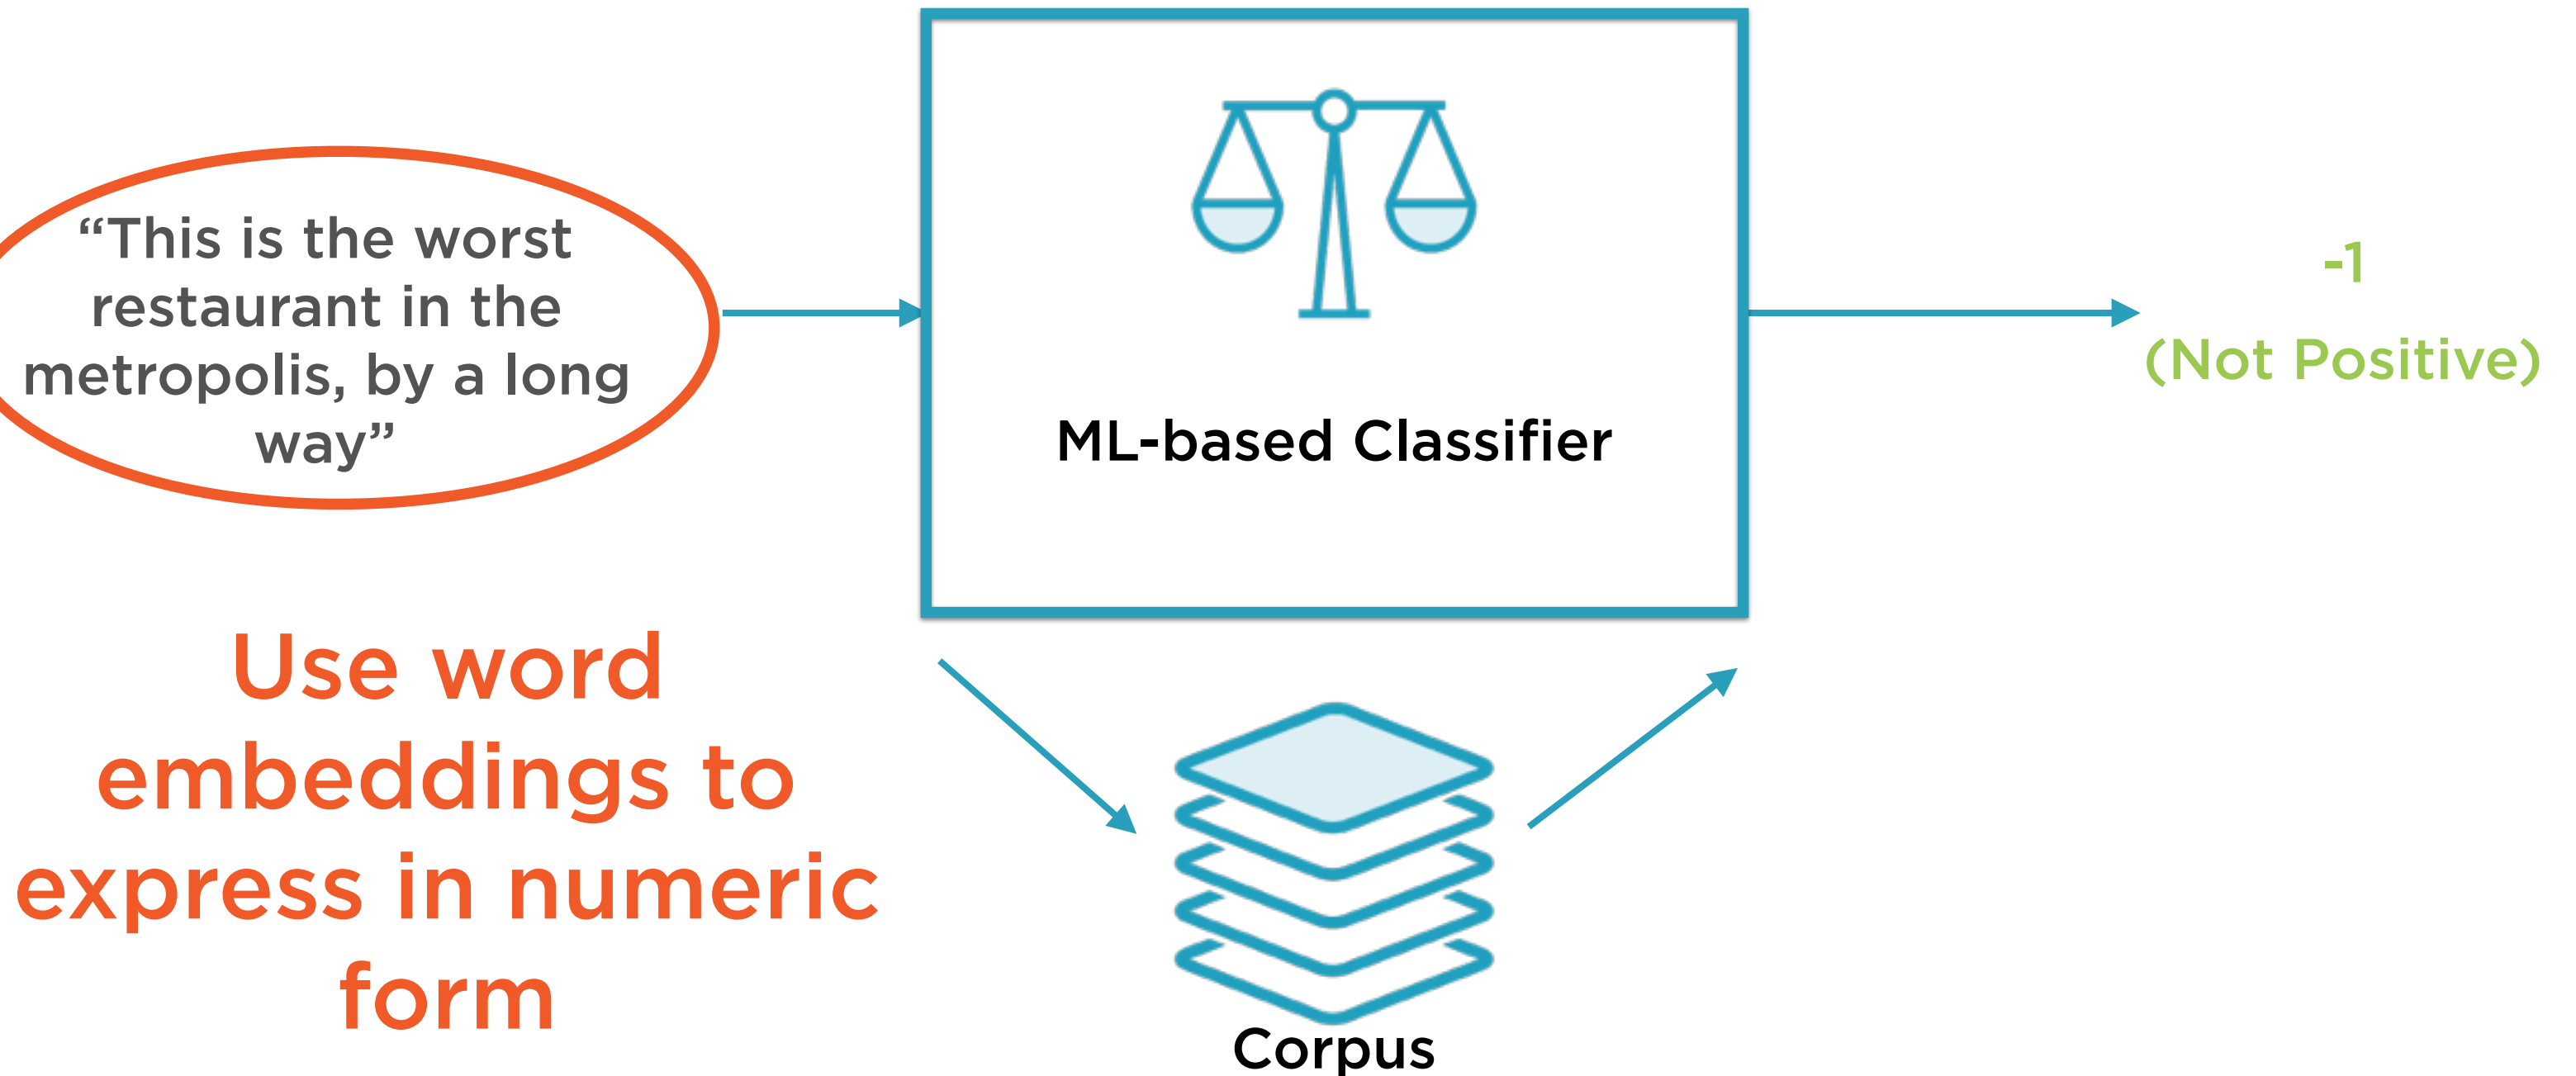
Sentiment Analysis

Classify a set of words

Sentiment Analysis Using Neural Networks



Sentiment Analysis Using Neural Networks



Neural networks are widely used in
language modeling

SemEval-2017

Each year SIGLEX publishes tasks as open challenges

**Semantic
comparison**

**Sentiment, humor
and truth**

**Parsing semantic
structures**

Semantic comparison

Semantic Textual Similarity

Given two sentences, return 0-5 score

- **0: Sentences meanings are unrelated**
- **5: Sentences have the same meaning**

Sub-tasks for different language pairs

- **Cross-lingual: Arabic-English, Spanish-English**
- **Mono-lingual: English-English, Spanish-Spanish**

**Semantic
comparison**

Community Question Answering

Input:

- Question
- Large number of user-submitted answers

Output:

- Ranking of user relevance

Additional sub-tasks: Reduce forum clutter

- Question similarity
- Relevance classification

SemEval-2017

Each year SIGLEX publishes tasks as open challenges

**Semantic
comparison**

**Sentiment, humor
and truth**

**Parsing semantic
structures**

**Sentiment, humor
and truth**

Sentiment Analysis

Twitter sentiment analysis

Fine-grained analysis on financial microblogs

**Sentiment, humor
and truth**

Detecting Humor

Given hashtag, find funniest tweet

Humor more subjective than sentiment

Binary classification approaches - simplistic for humor

Inside jokes/references: how to incorporate external knowledge?

**Sentiment, humor
and truth**

Puns

Word sense disambiguation (WSD)

Homophonic (perfect) puns

“I used to be a banker but I lost interest”

Heterophonic (imperfect) puns

**“With fronds like these, who needs
anemones?”**

SemEval-2017

Each year SIGLEX publishes tasks as open challenges

**Semantic
comparison**

**Sentiment, humor
and truth**

**Parsing semantic
structures**

**Parsing semantic
structures**

Extracting Keyphrases

Given academic publication, extract key phrases and relationships

Identifying relationships is a classic language modeling task

Cause-effect Identification: Given a sentence, tag cause and effect

Parsing semantic structures

Types of Relationships

Cause-effect

“The *cancer* was caused by *radiation*”

Instrument-agency

“The *catcher* used a *mitt*”

Product-producer

“The *craftsman* built fine *watches*”

Content-container

“Old *wine* in new *bottles*”

Entity-origin

“*Friends* from *faraway places*”

Types of Cells Used in RNNs

Types of Neurons

Simple Neuron

Affine transformation,
activation function

LSTM Cell

Maintain complex additional
state for long-memory

Multi-RNN Cell

Wrap multiple GRU cells into
single multi-layer cell

Recurrent Neuron

Feed output back as
another input

GRU Cell

Similar results as LSTM,
but simpler internals

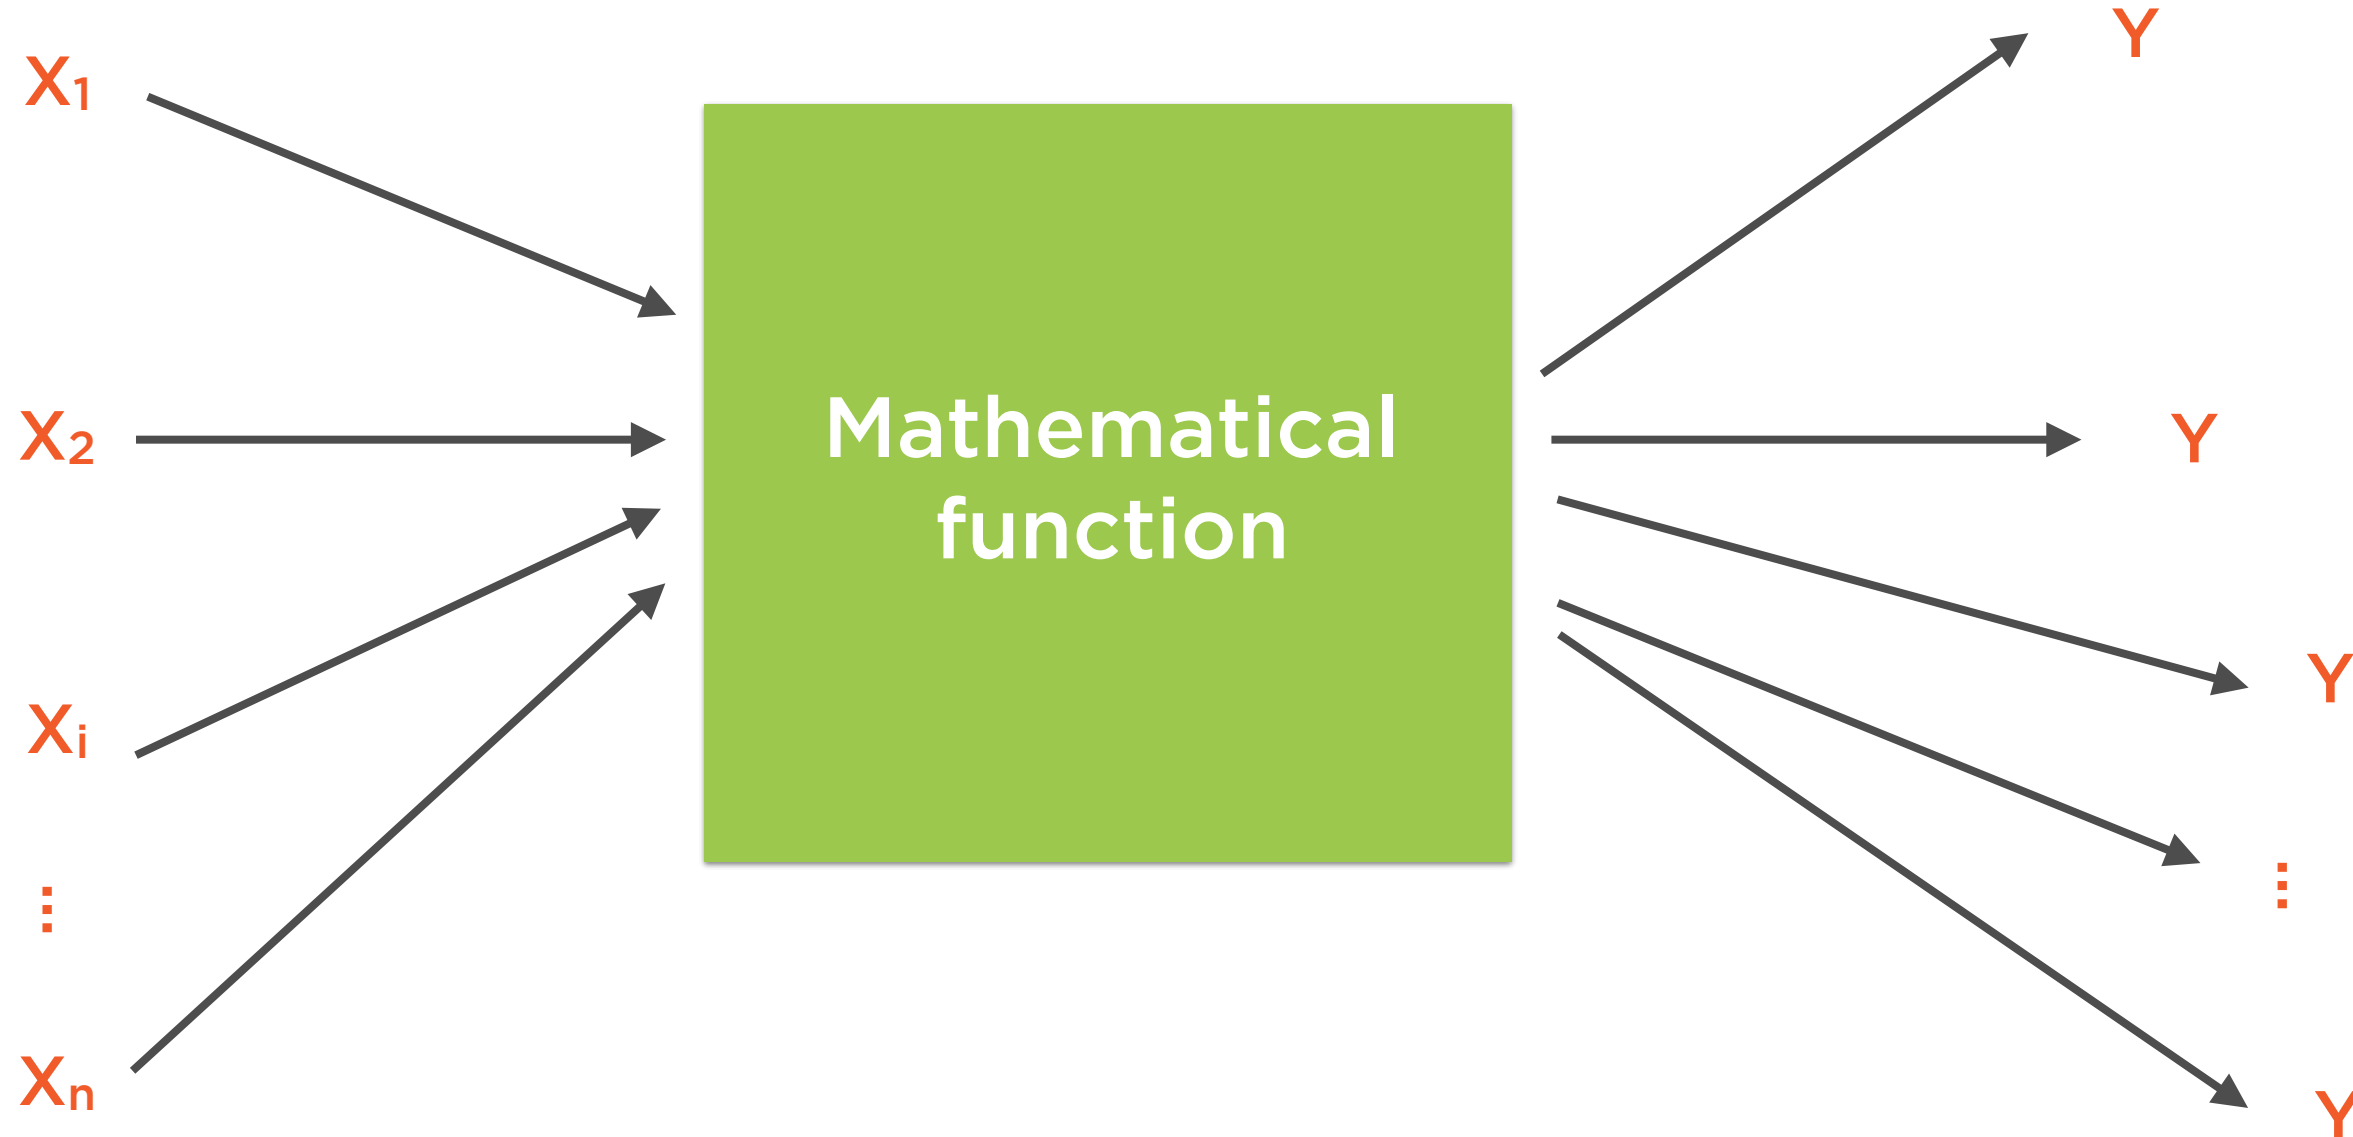
Types of Neurons

Simple Neuron

Affine transformation,
activation function




Operation of a Single Neuron




For an active neuron a change in inputs should trigger a corresponding change in the outputs

Operation of a Single Neuron




The outputs of neurons feed into the neurons from the next layer

Operation of a Single Neuron



Each connection is associated with a weight

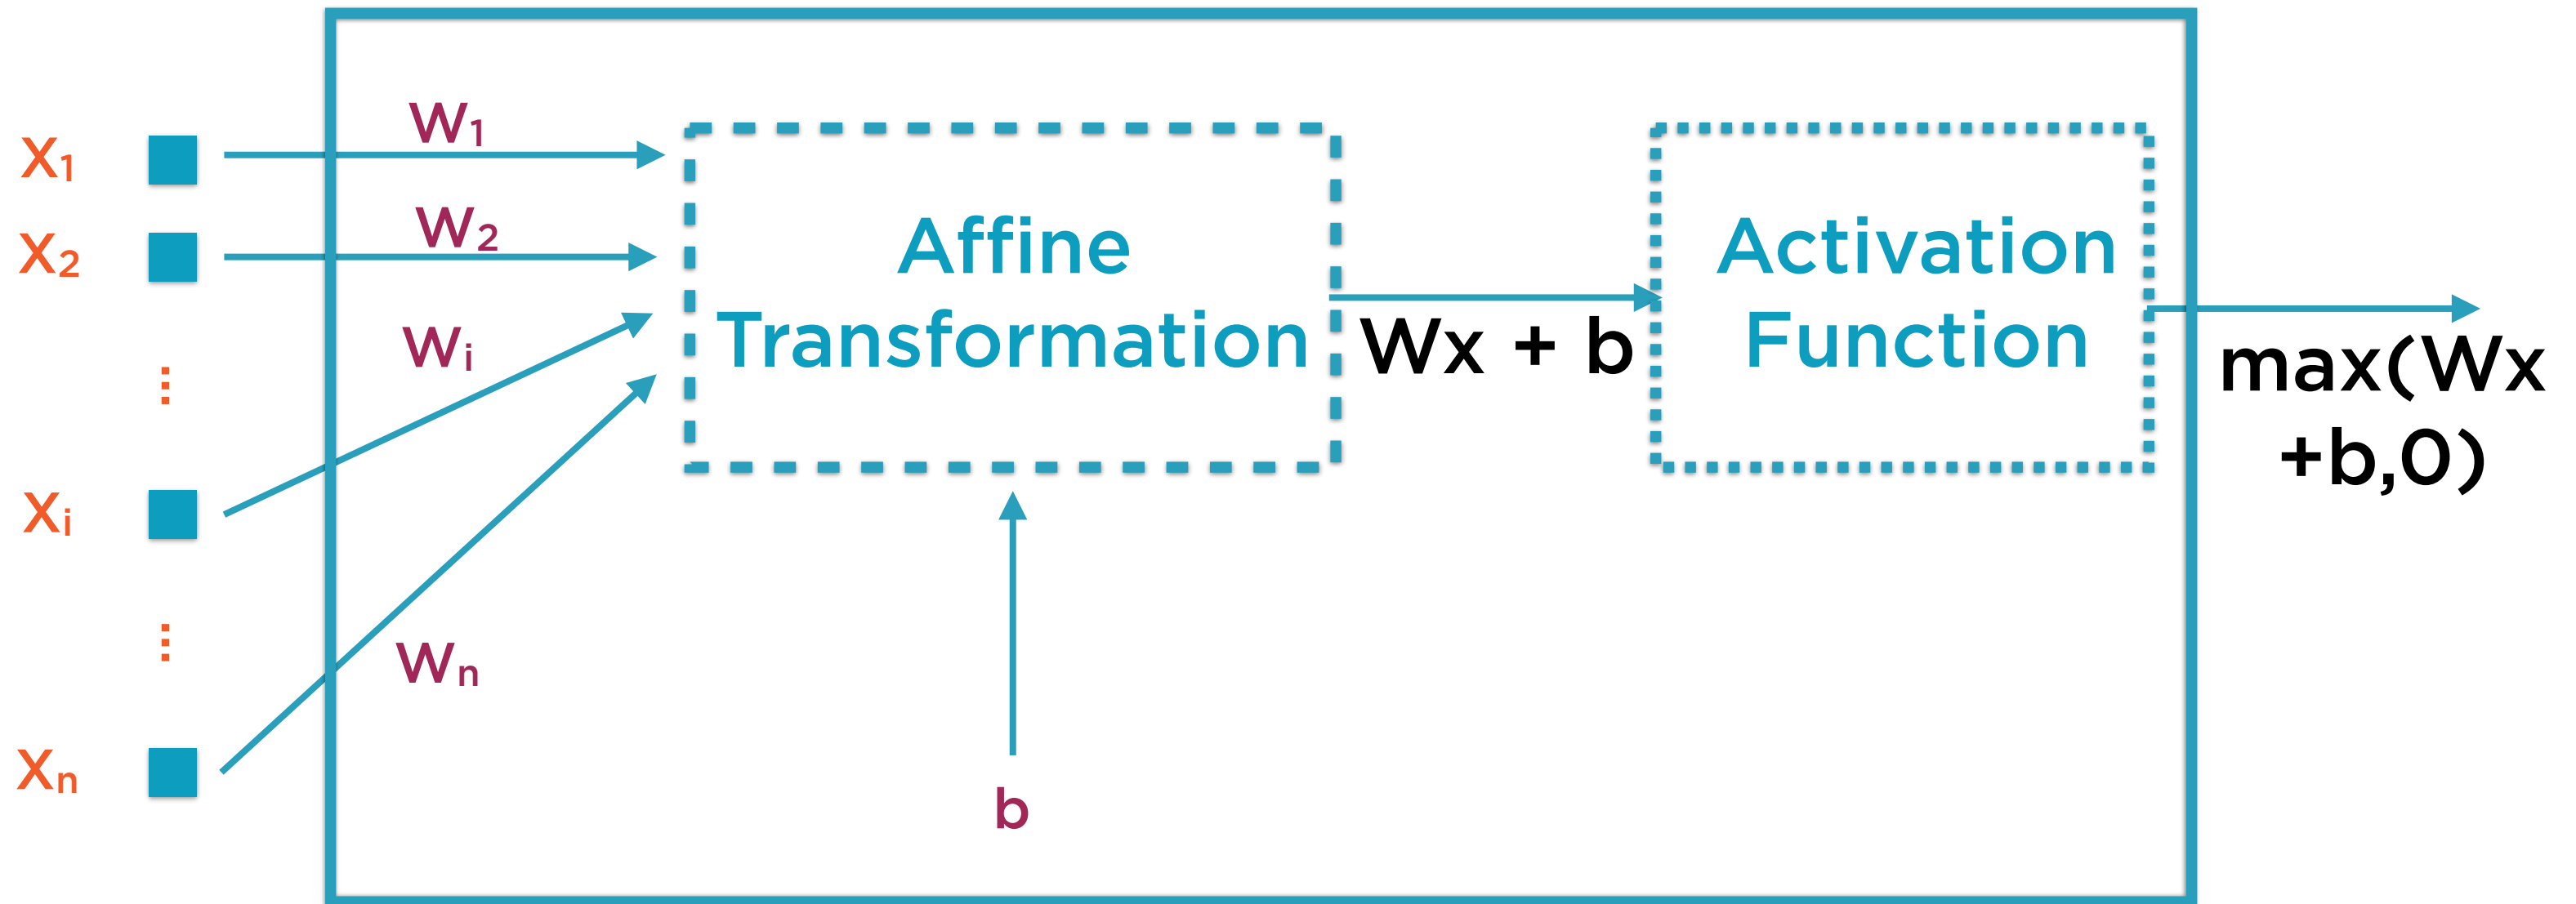
Operation of a Single Neuron



If the second neuron is sensitive to the output of the first neuron, the **connection between them gets stronger**

W increases

Neuron as a Learning Unit



The combination of the affine transformation and the activation function allows the neuron to learn any arbitrary relationship

Types of Neurons


Simple Neuron

Affine transformation,
activation function


Recurrent Neuron

Feed output back as
another input


Simplest Recurrent Neuron




Unrolling Through Time



Unrolling Through Time




Layer of Recurrent Neurons



A layer of neurons forms an RNN cell

Layer of Recurrent Neurons



The cells unrolled through time form the layers of the
neural network

Types of Neurons

Simple Neuron

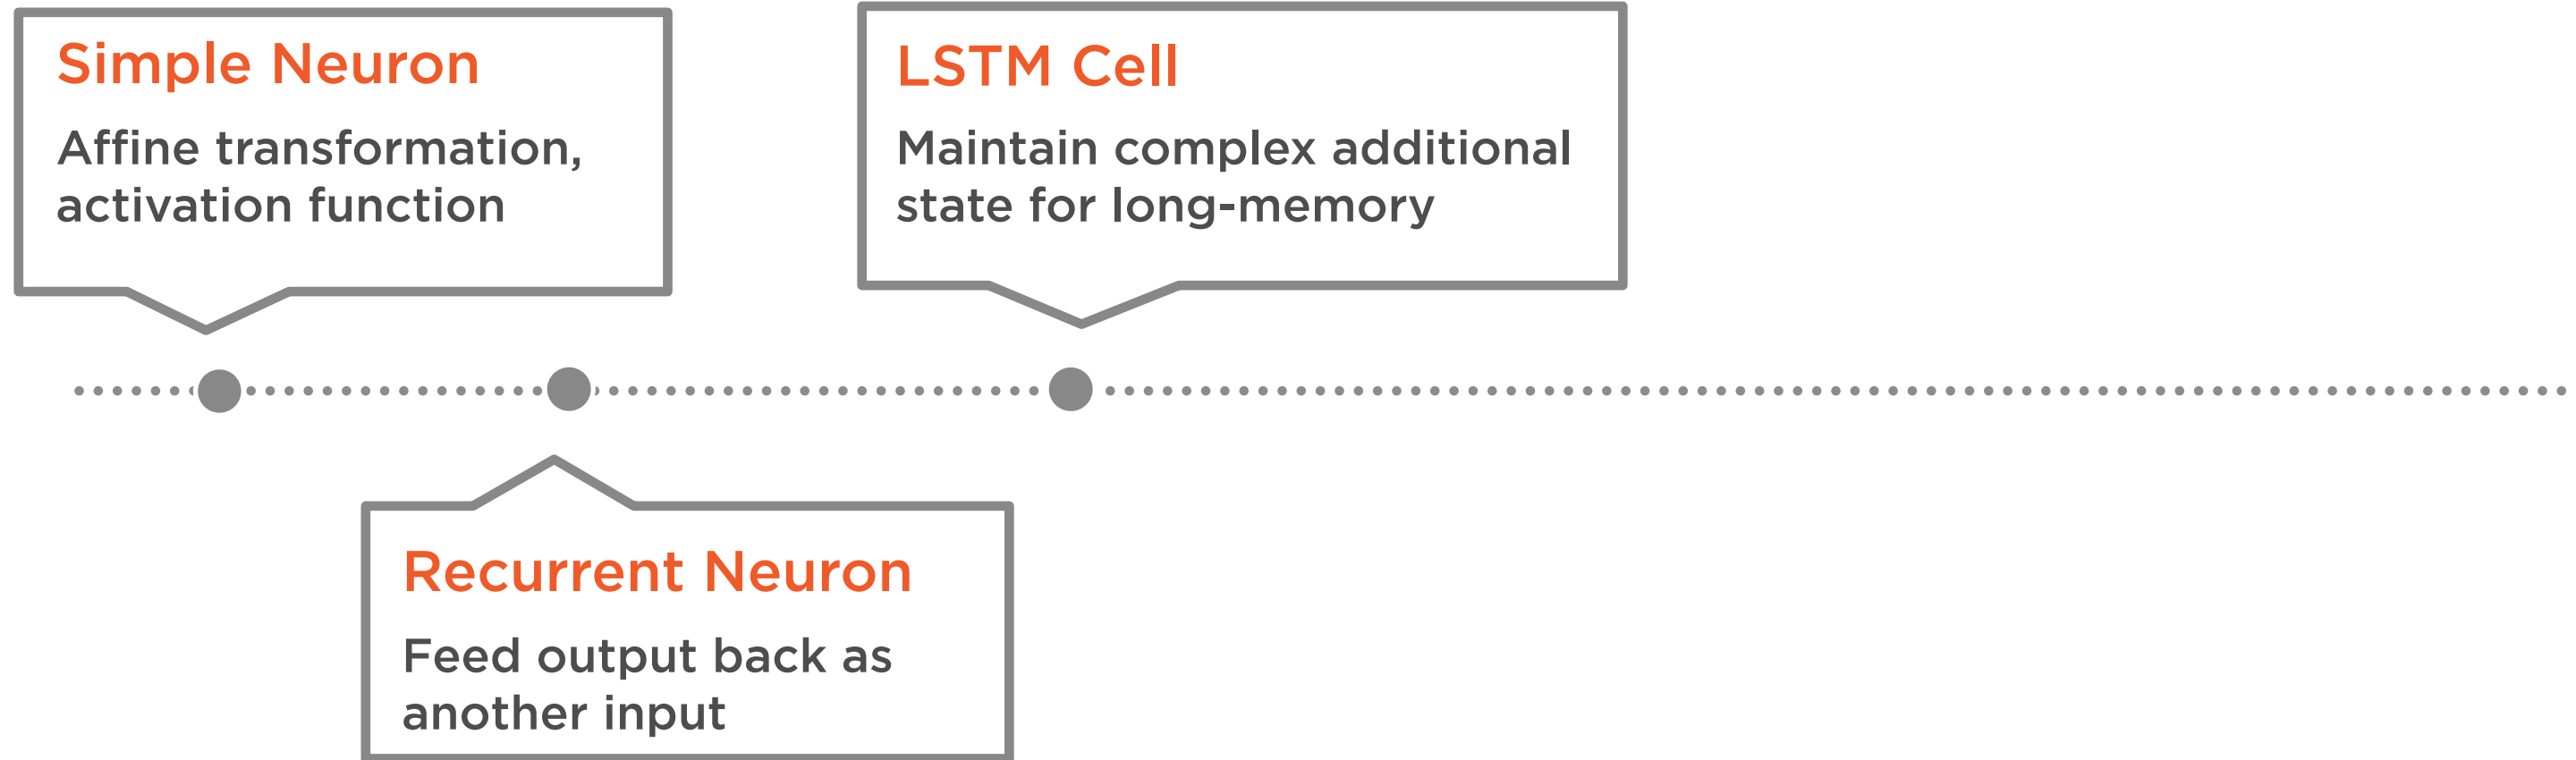
Affine transformation,
activation function

LSTM Cell


Maintain complex additional
state for long-memory

Recurrent Neuron


Feed output back as
another input




Simplest Recurrent Neuron




Simplest Recurrent Neuron



Long Memory Recurrent Neuron



Long Memory RNNs



Increase the amount of state in neuron

Effect is to increase memory of neuron

Could explicitly add:

- long-term state (c)
- short-term state (h)

Long memory neurons have several advantages over basic RNNs

Long Memory RNNs



Advantages in Training


Faster training, nicer gradients



Advantages in Prediction

No need to truncate BPTT

Long/Short-Term Memory Cell (LSTM)



Types of Neurons

Simple Neuron

Affine transformation,
activation function

LSTM Cell

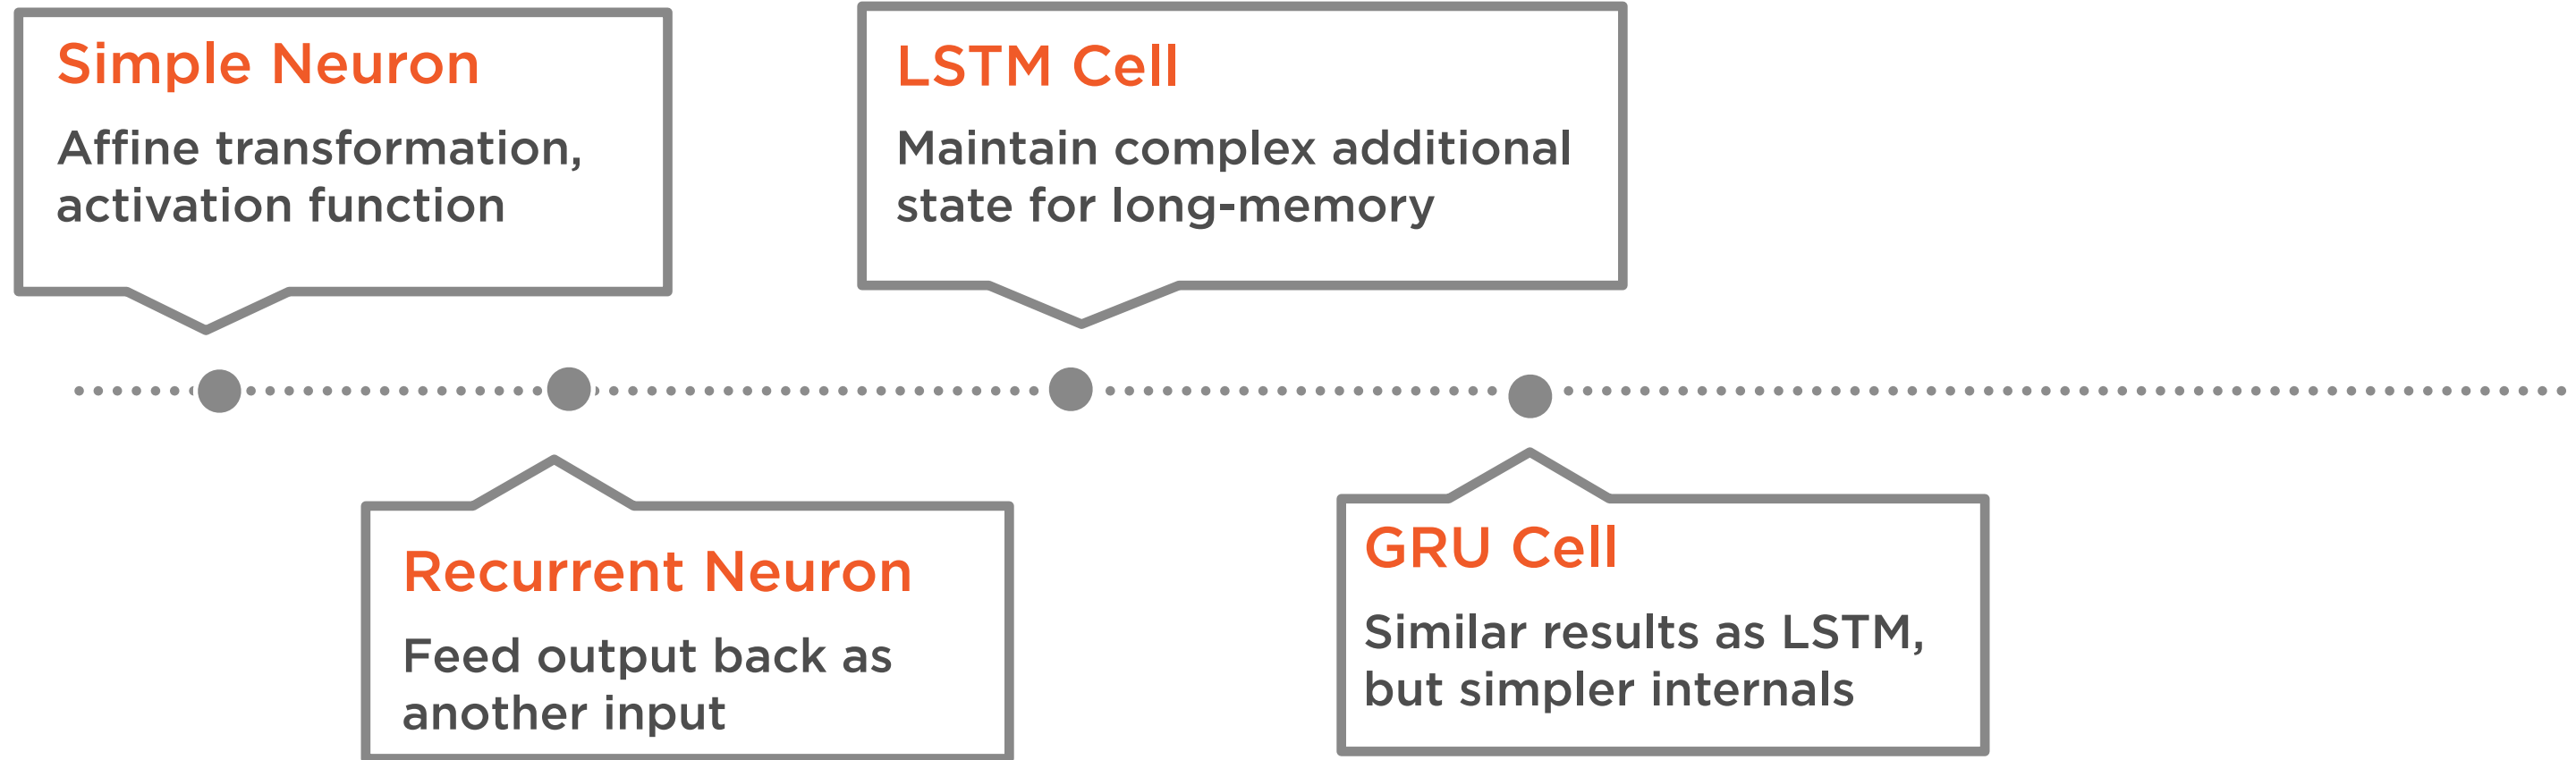
Maintain complex additional
state for long-memory

Recurrent Neuron


Feed output back as
another input

GRU Cell

Similar results as LSTM,
but simpler internals



GRU



Peephole connections: LSTM cells that store state for more than 1 period

Gated Recurrent Unit (GRU): Simplified LSTM with better performance

- Only 1 state vector
- Fewer internal gates and NNs

Multi-RNN Cell

Types of Neurons

Simple Neuron

Affine transformation,
activation function

LSTM Cell

Maintain complex additional
state for long-memory

Multi-RNN Cell


Wrap multiple GRU cells into
single 2-layer cell

Recurrent Neuron

Feed output back as
another input

GRU Cell

Similar results as LSTM,
but simpler internals




Multi-RNN Cell

Stack multiple RNN cells into “combined” RNN cell

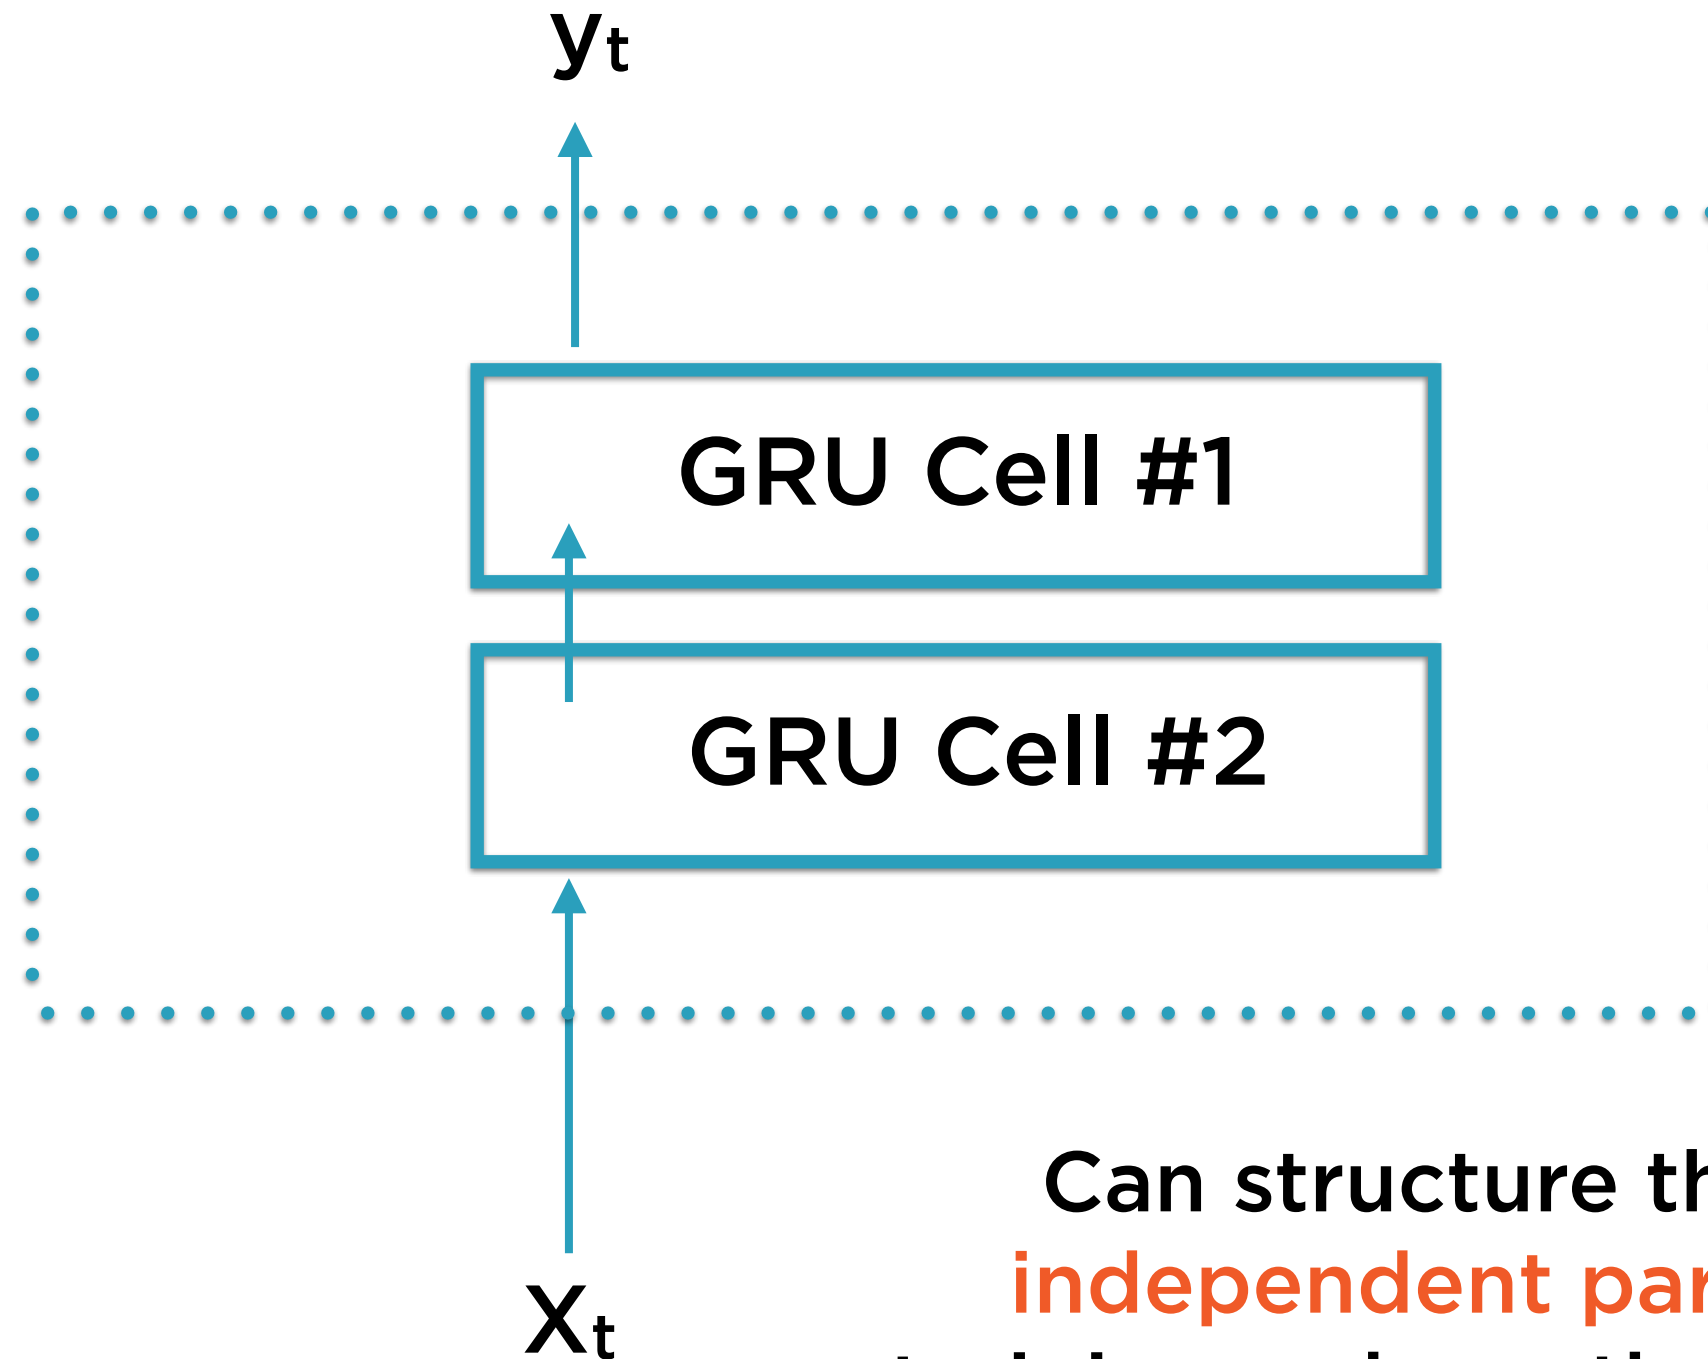
In our example, use 2 GRU cells inside each multi-RNN cell

Multi-RNN Cell



This can hold any number of cells
which are stacked one on top of
another

Multi-RNN Cell



Can structure the cells to have
independent parameters during
training or have the **same parameters**

Multi-RNN cells allow you to wrap multiple cells allowing them to **look and behave like a single cell**

$$y_t = f(x_t, y_{t-1})$$

Learning the (Recent) Past

Unrolling the RNN through time helps learn the past

$$y_t = f(x_t, y_{t-1}, y_{t-2} \dots, y_{t-1000})$$

Learning the Distant Past

The unrolled RNN will be very, very deep - many layers to train, the gradient has to be propagated a long way

$$y_t = f(x_t, y_{t-1}, y_{t-2} \dots, y_{t-1000})$$

Learning the Distant Past

Using LSTM and GRU cells helps maintain memory of the distant past with internal state rather than large number of layers

Key insight: Smart re-use of prior period state is key to prediction

$$y_t, \text{Prev_Internal_State}_t = f(x_t, y_{t-1}, \text{Prev_Internal_State}_{t-1})$$

Alternative Approach to the Distant Past

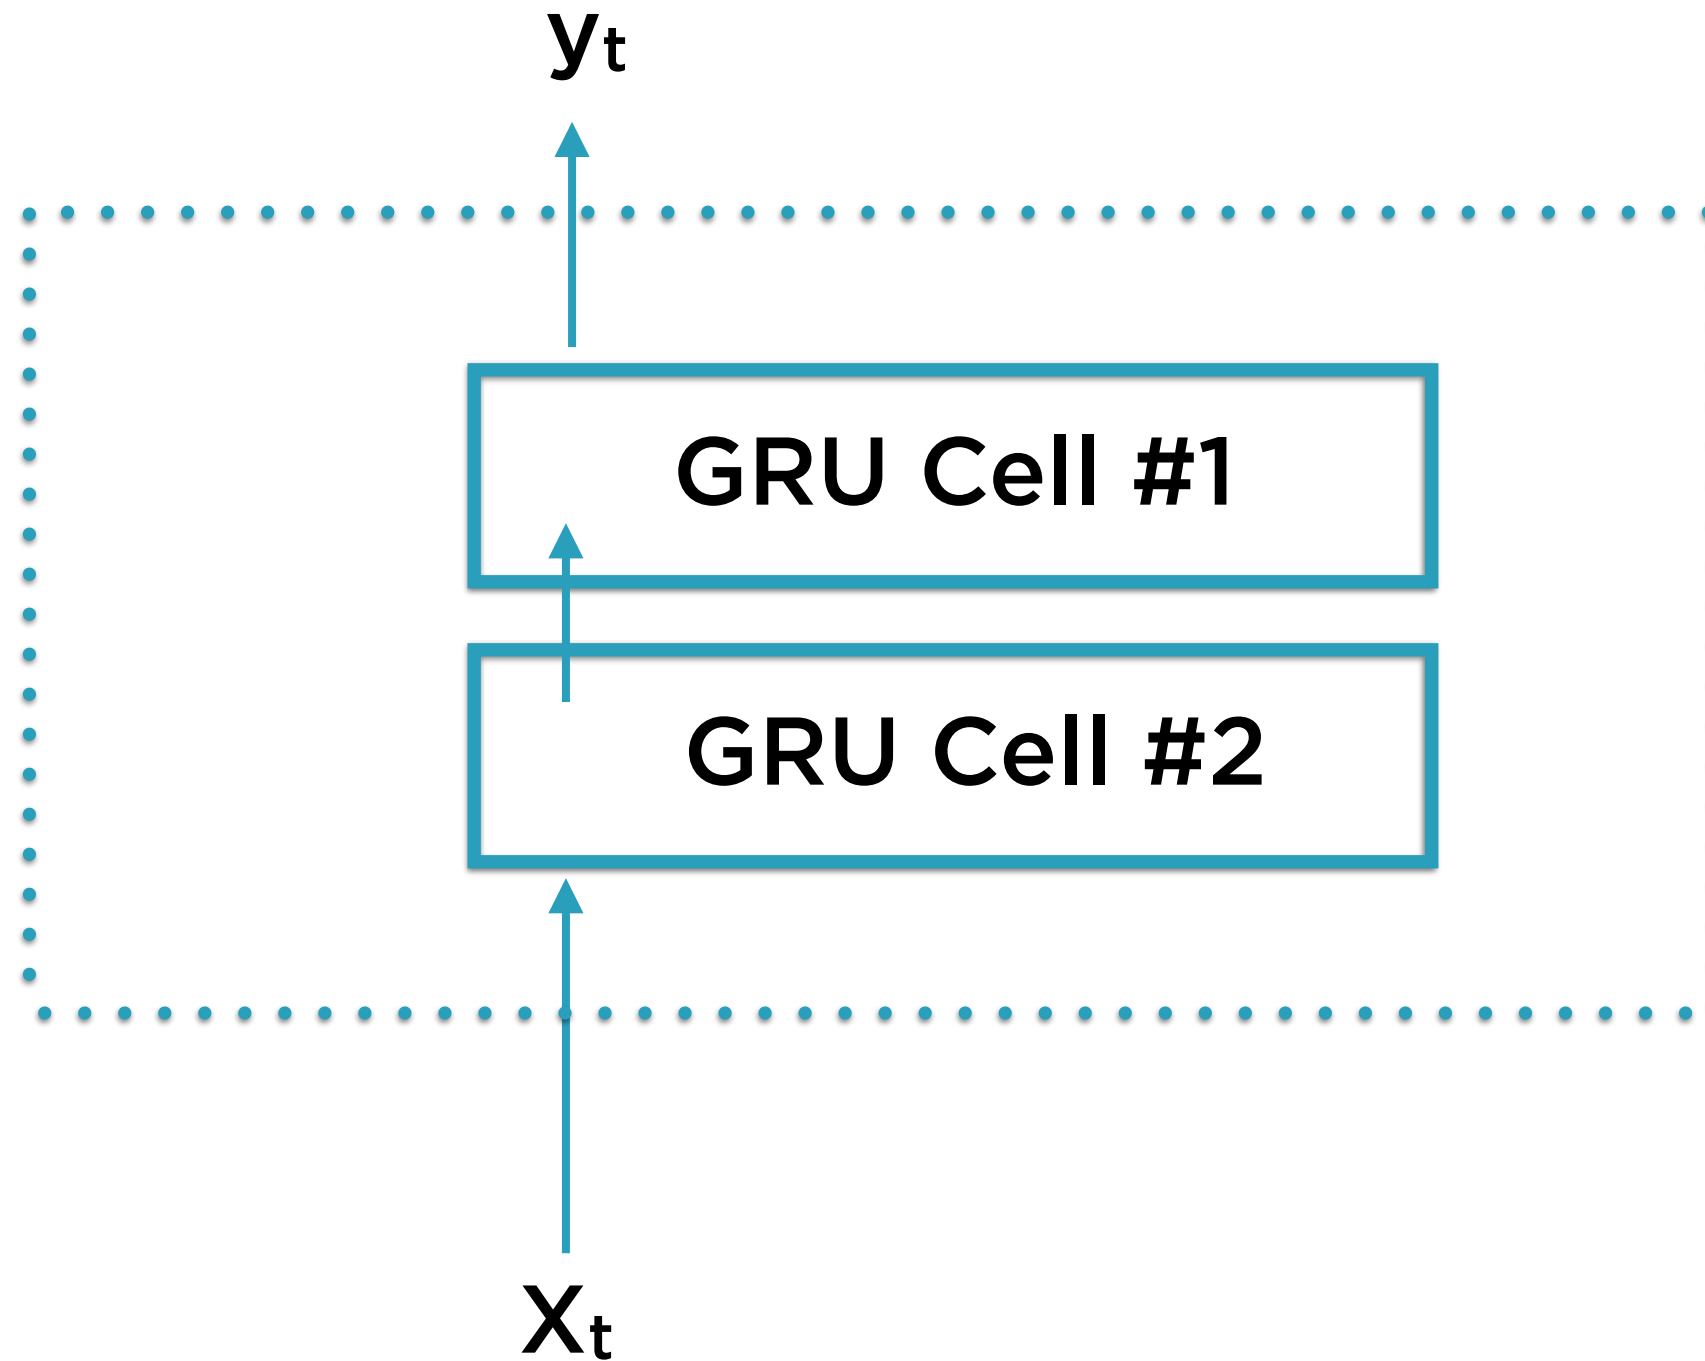
Re-using internal state in addition to using GRU gives great performance with less input

$$y_t, \text{Prev_Internal_State}_t = f(x_t, y_{t-1}, \text{Prev_Internal_State}_{t-1})$$

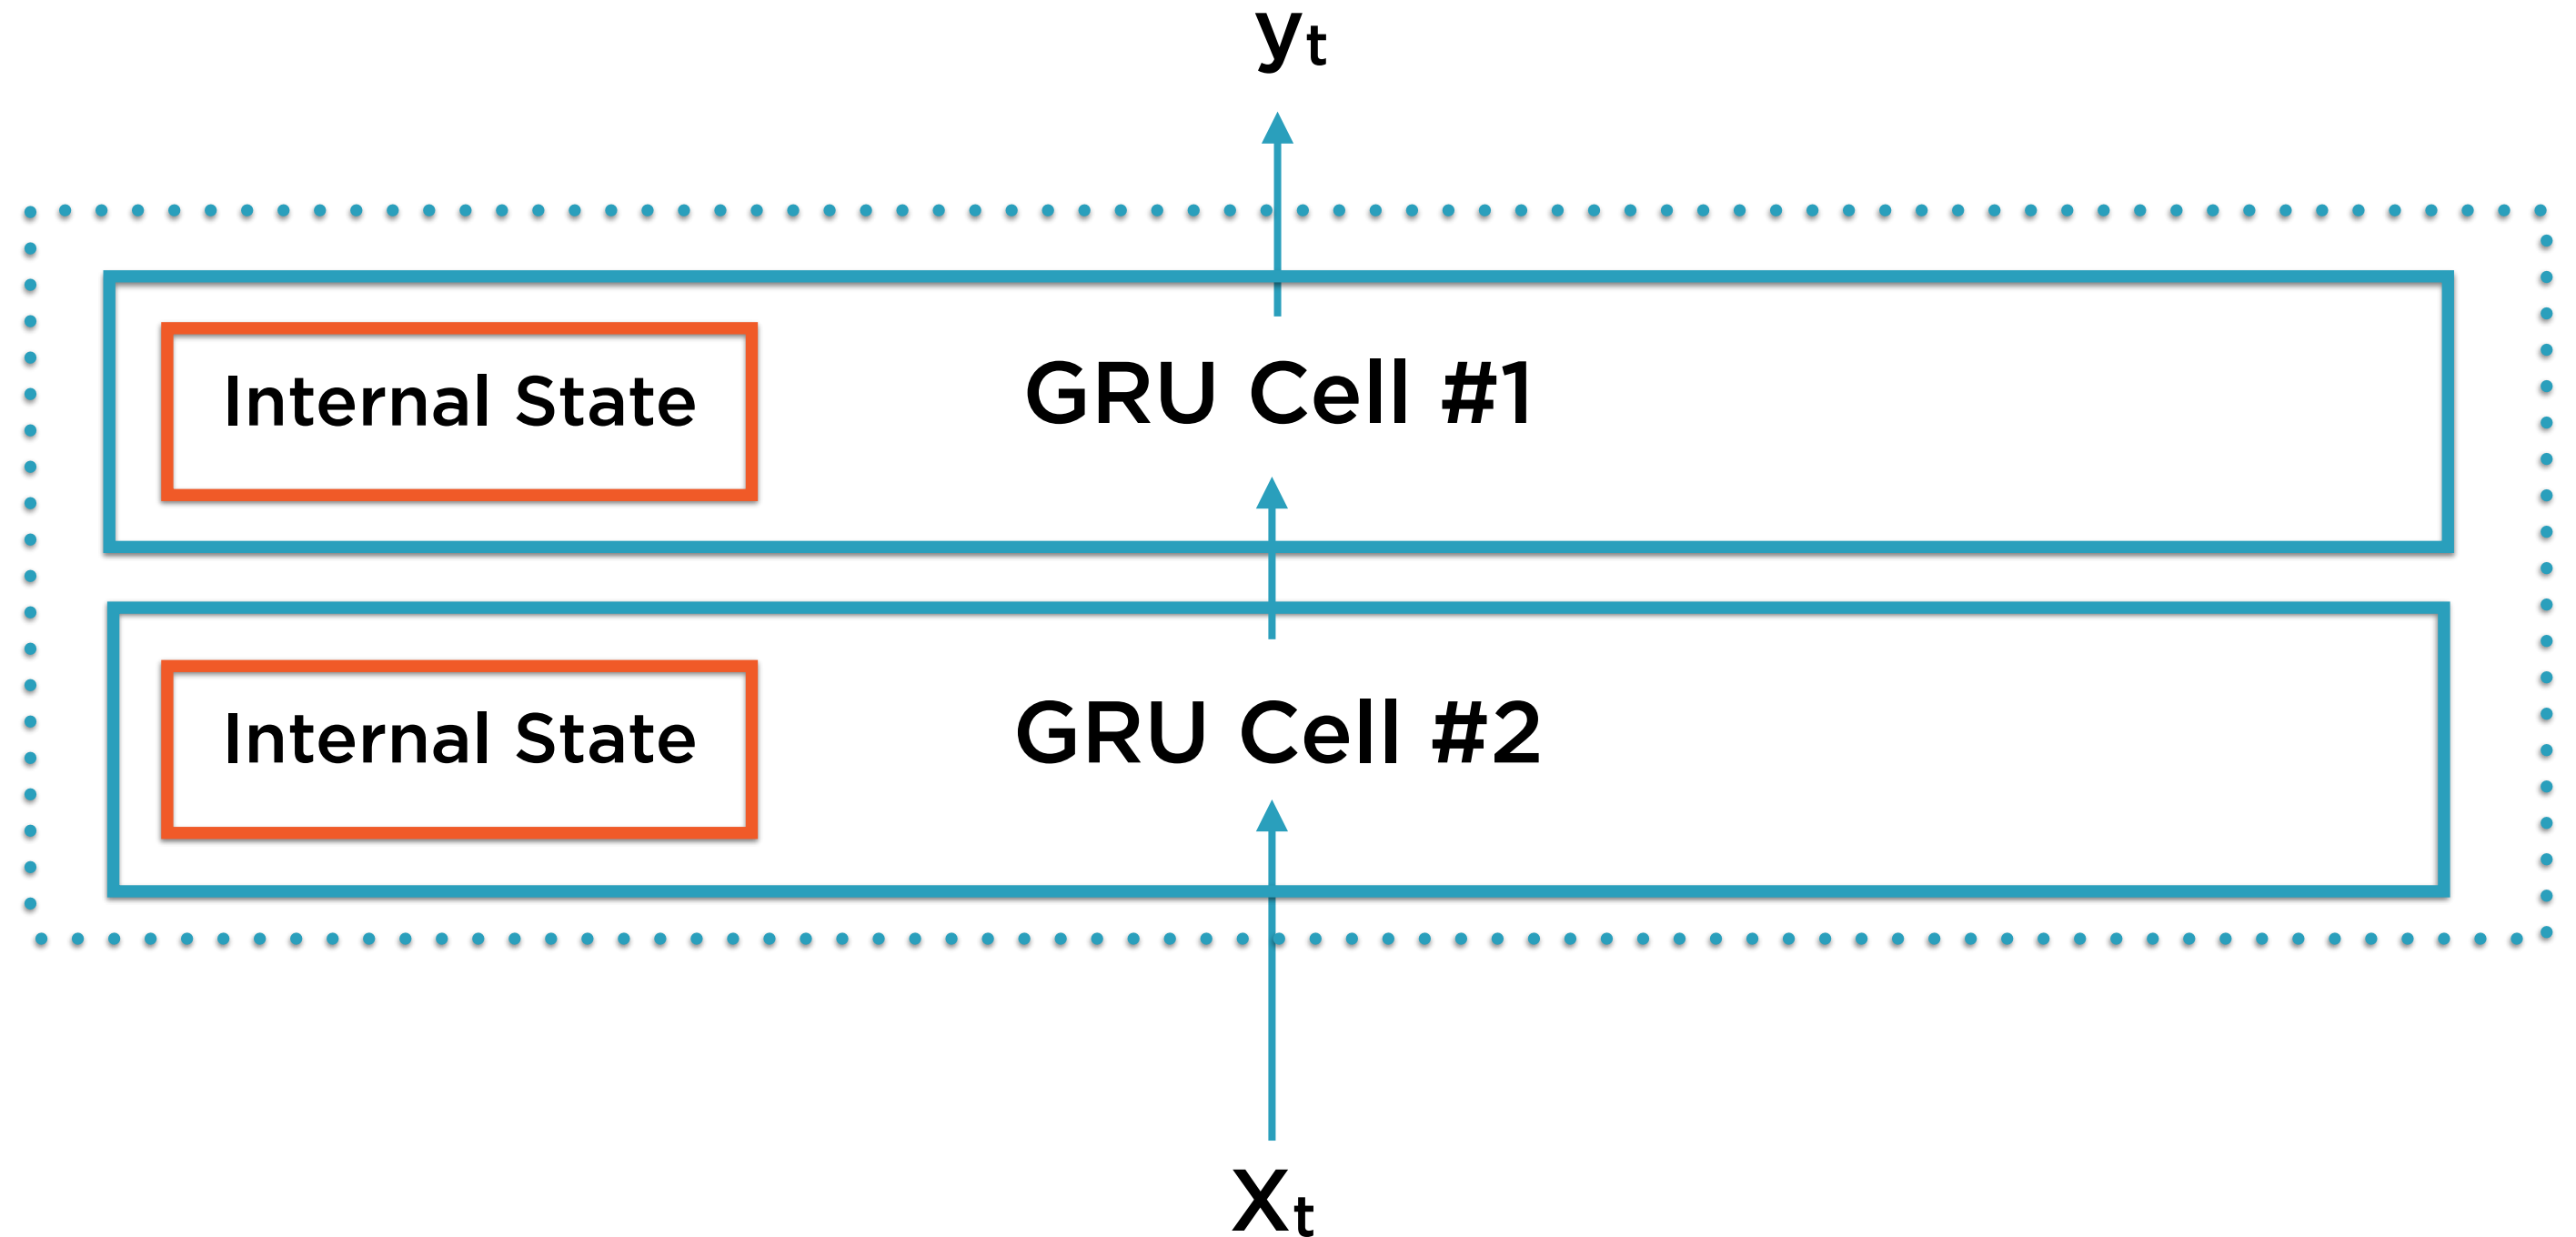
Feed the Previous State as Initial Values

Re-using internal state in addition to using GRU gives great performance with less input

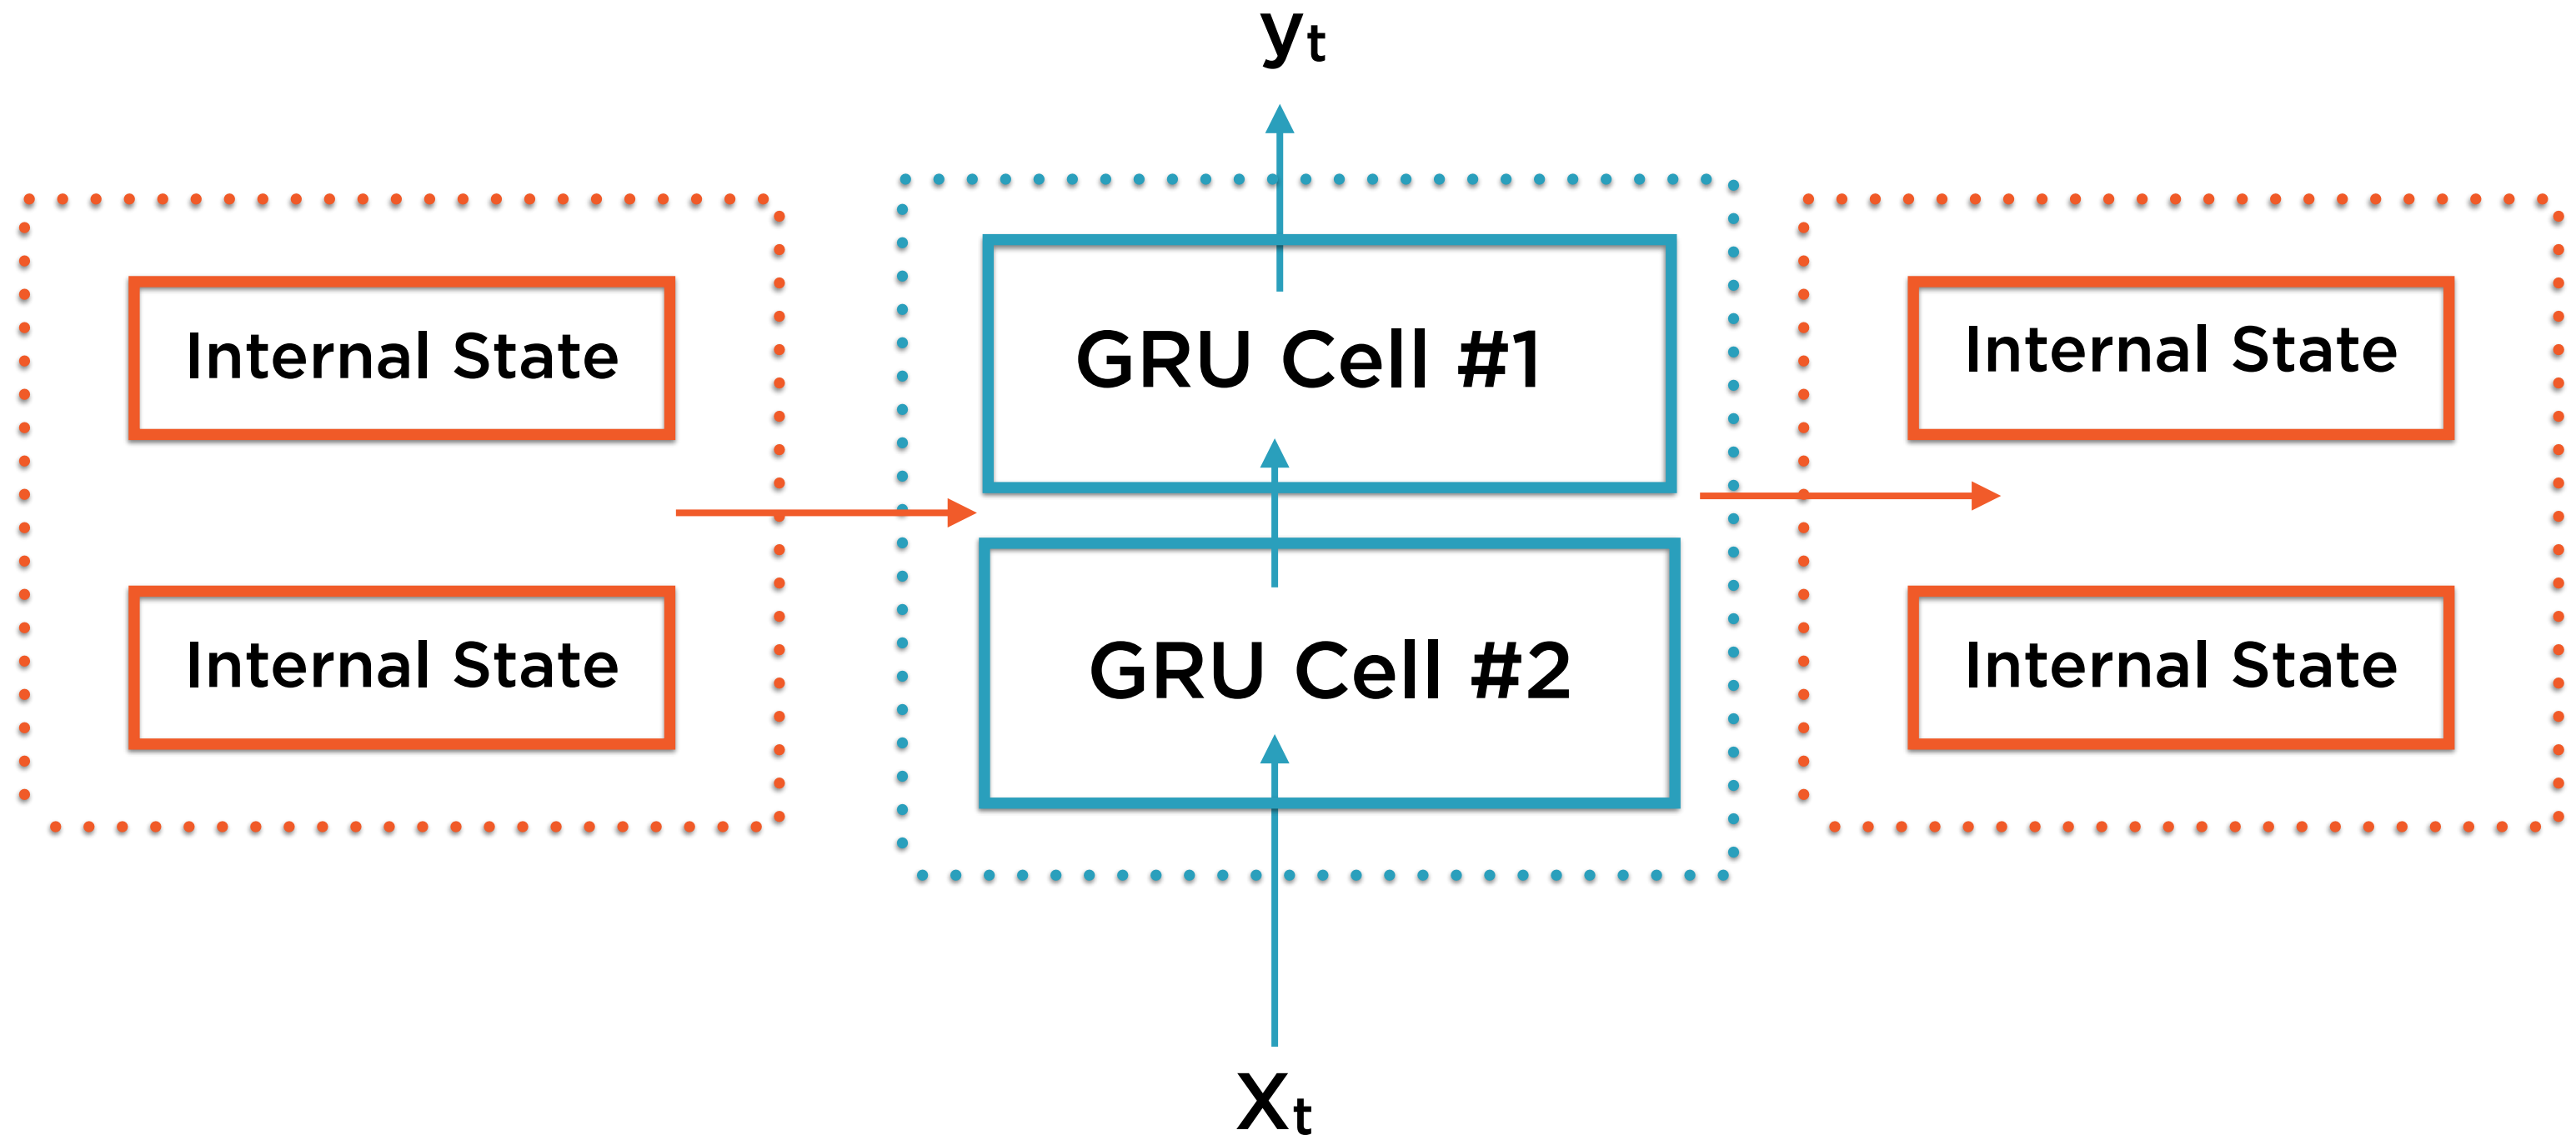
Re-using Internal State



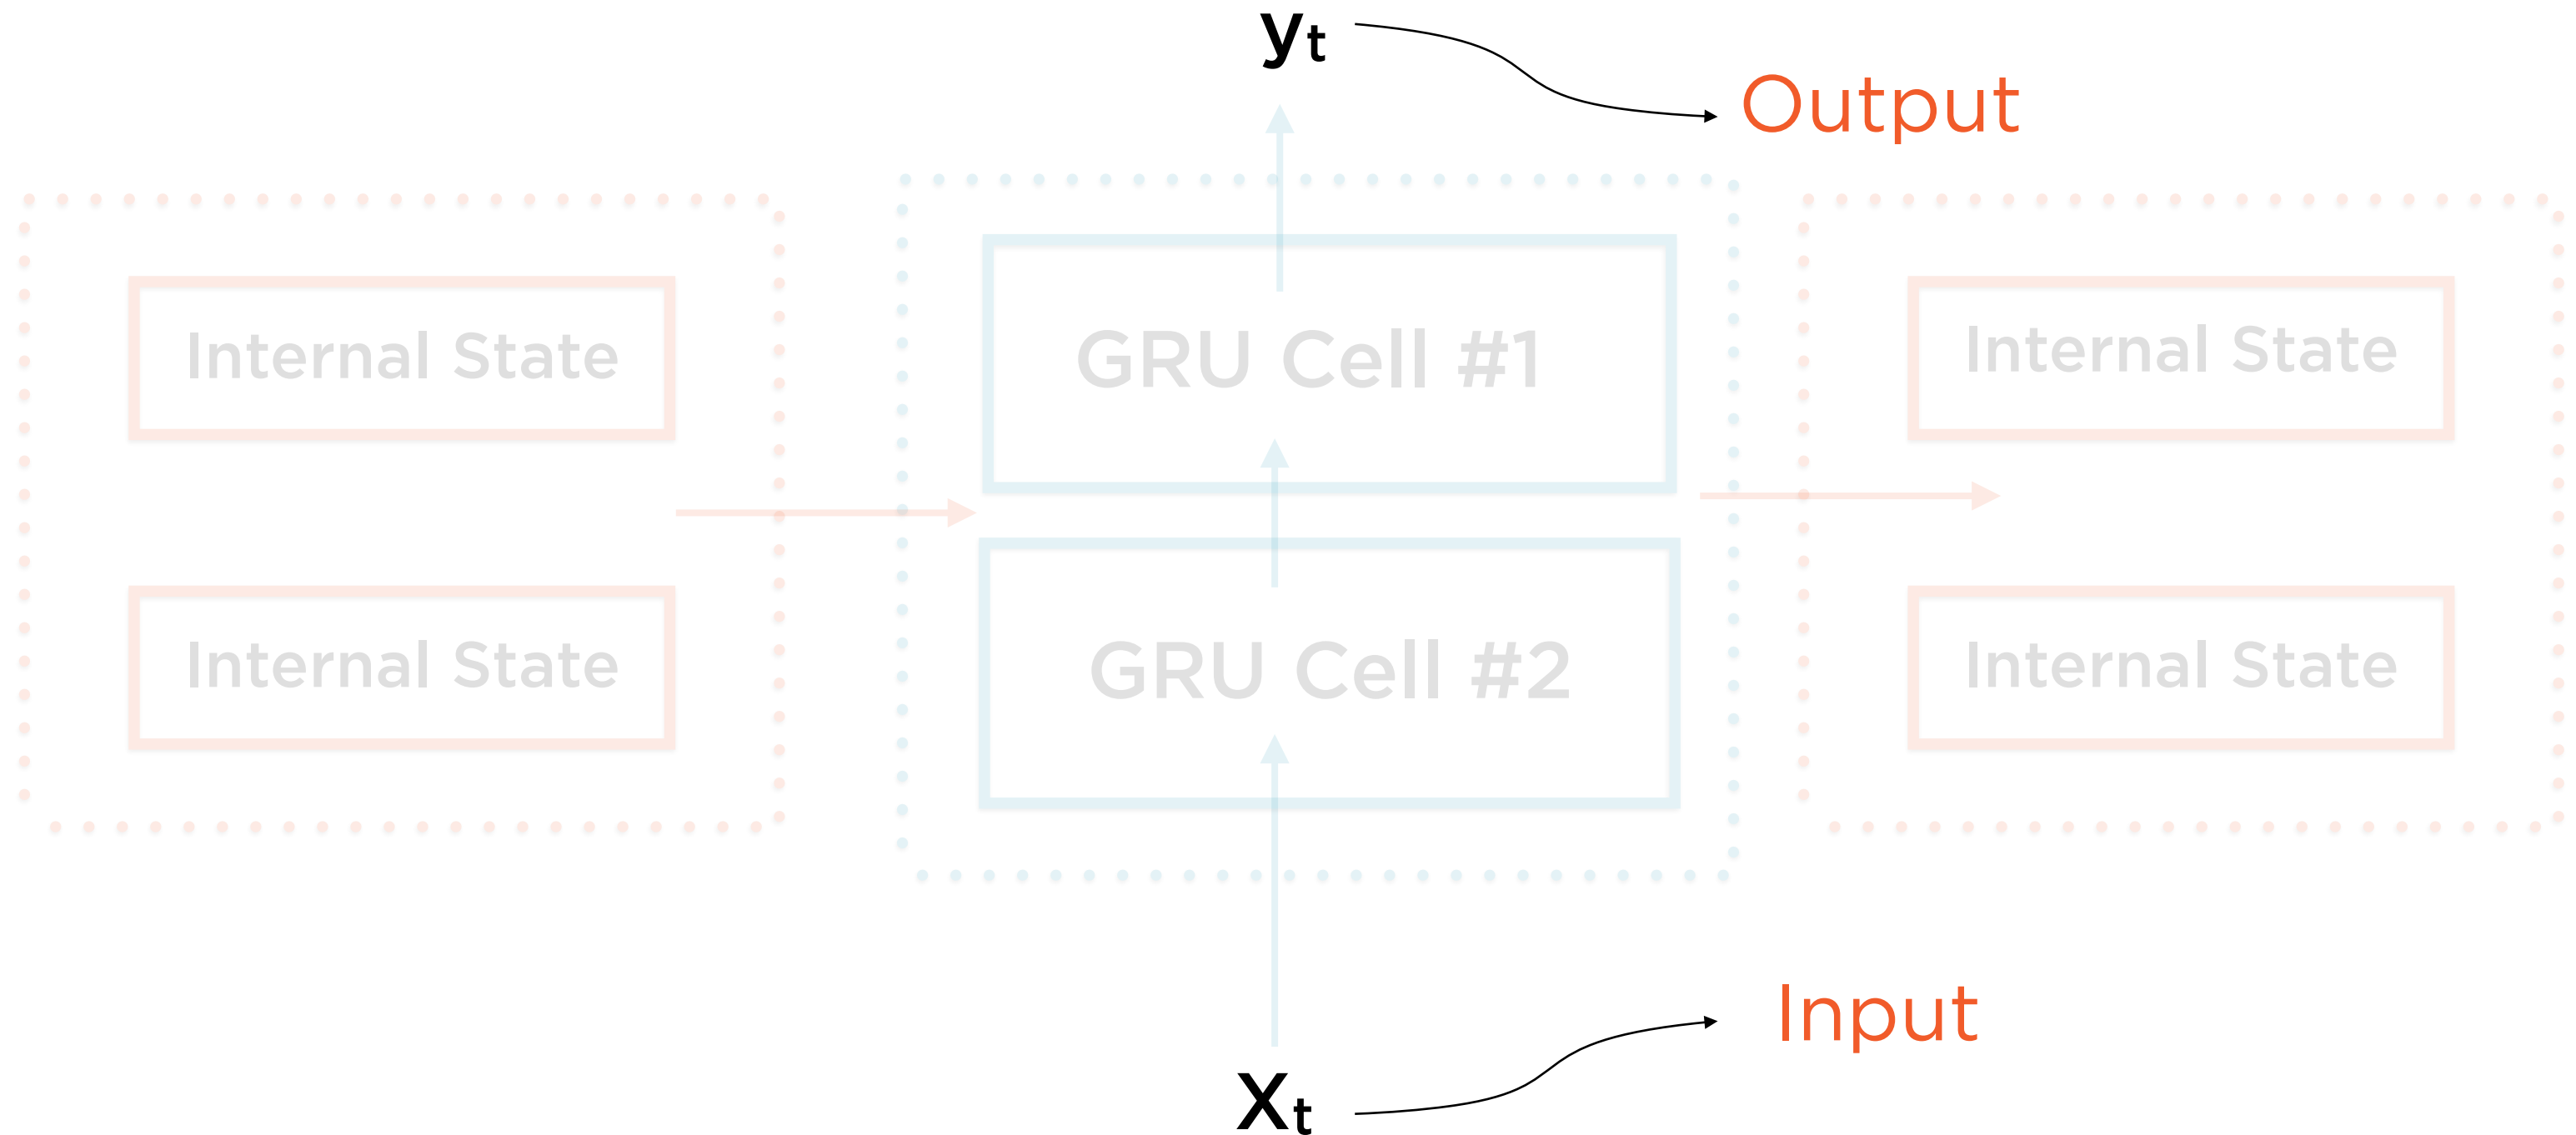
Re-using Internal State




Re-using Internal State



Re-using Internal State




Re-using Internal State




Re-using Internal State

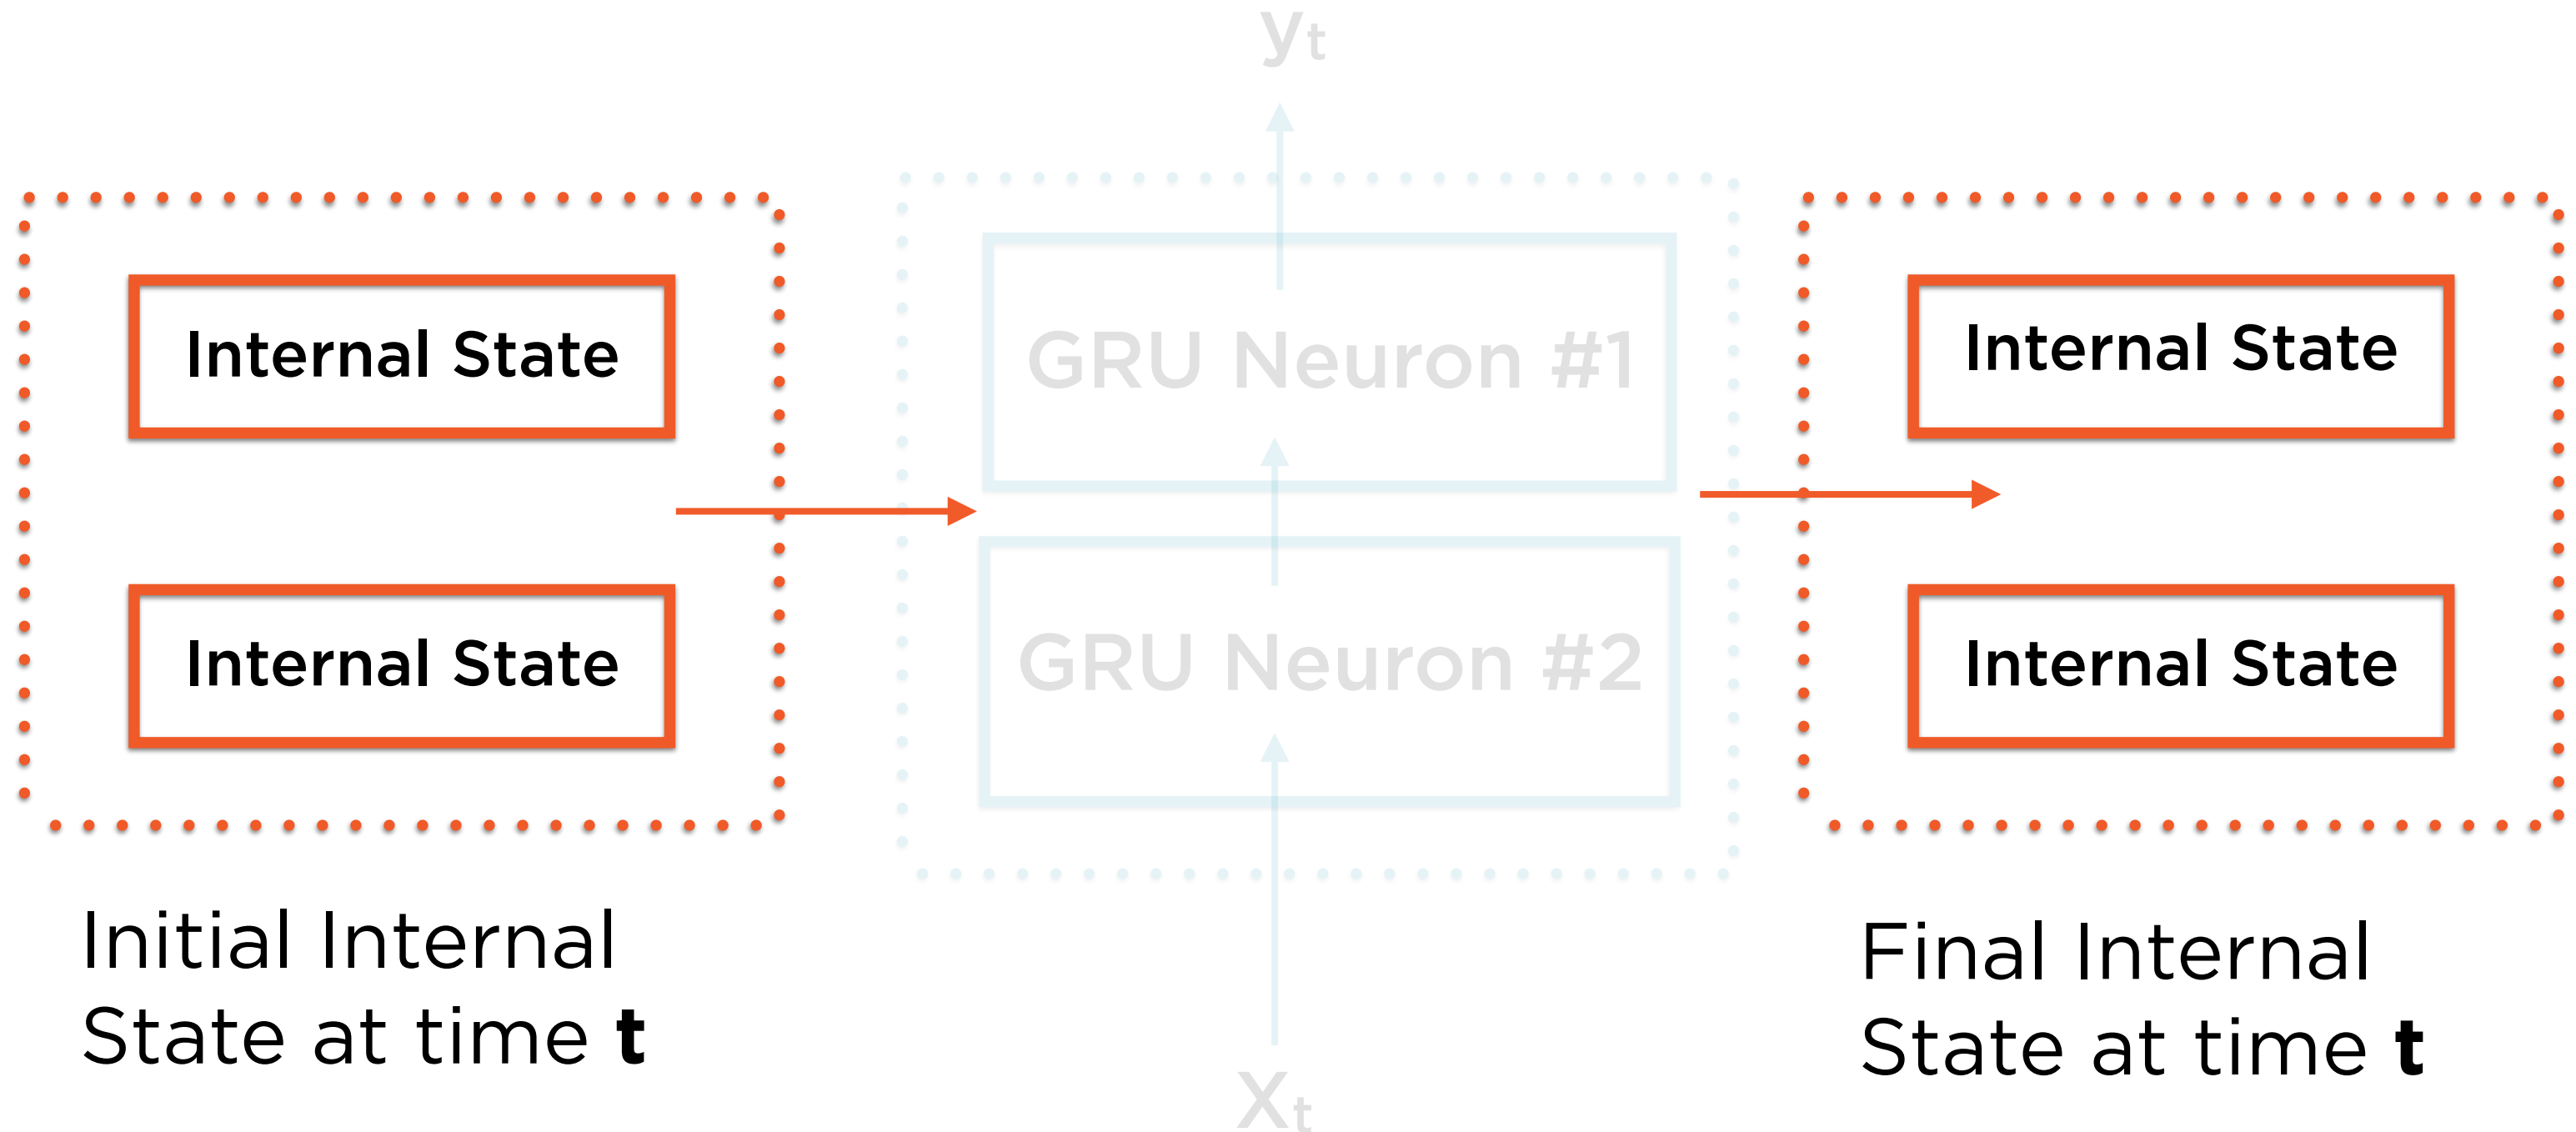
Initial Internal State



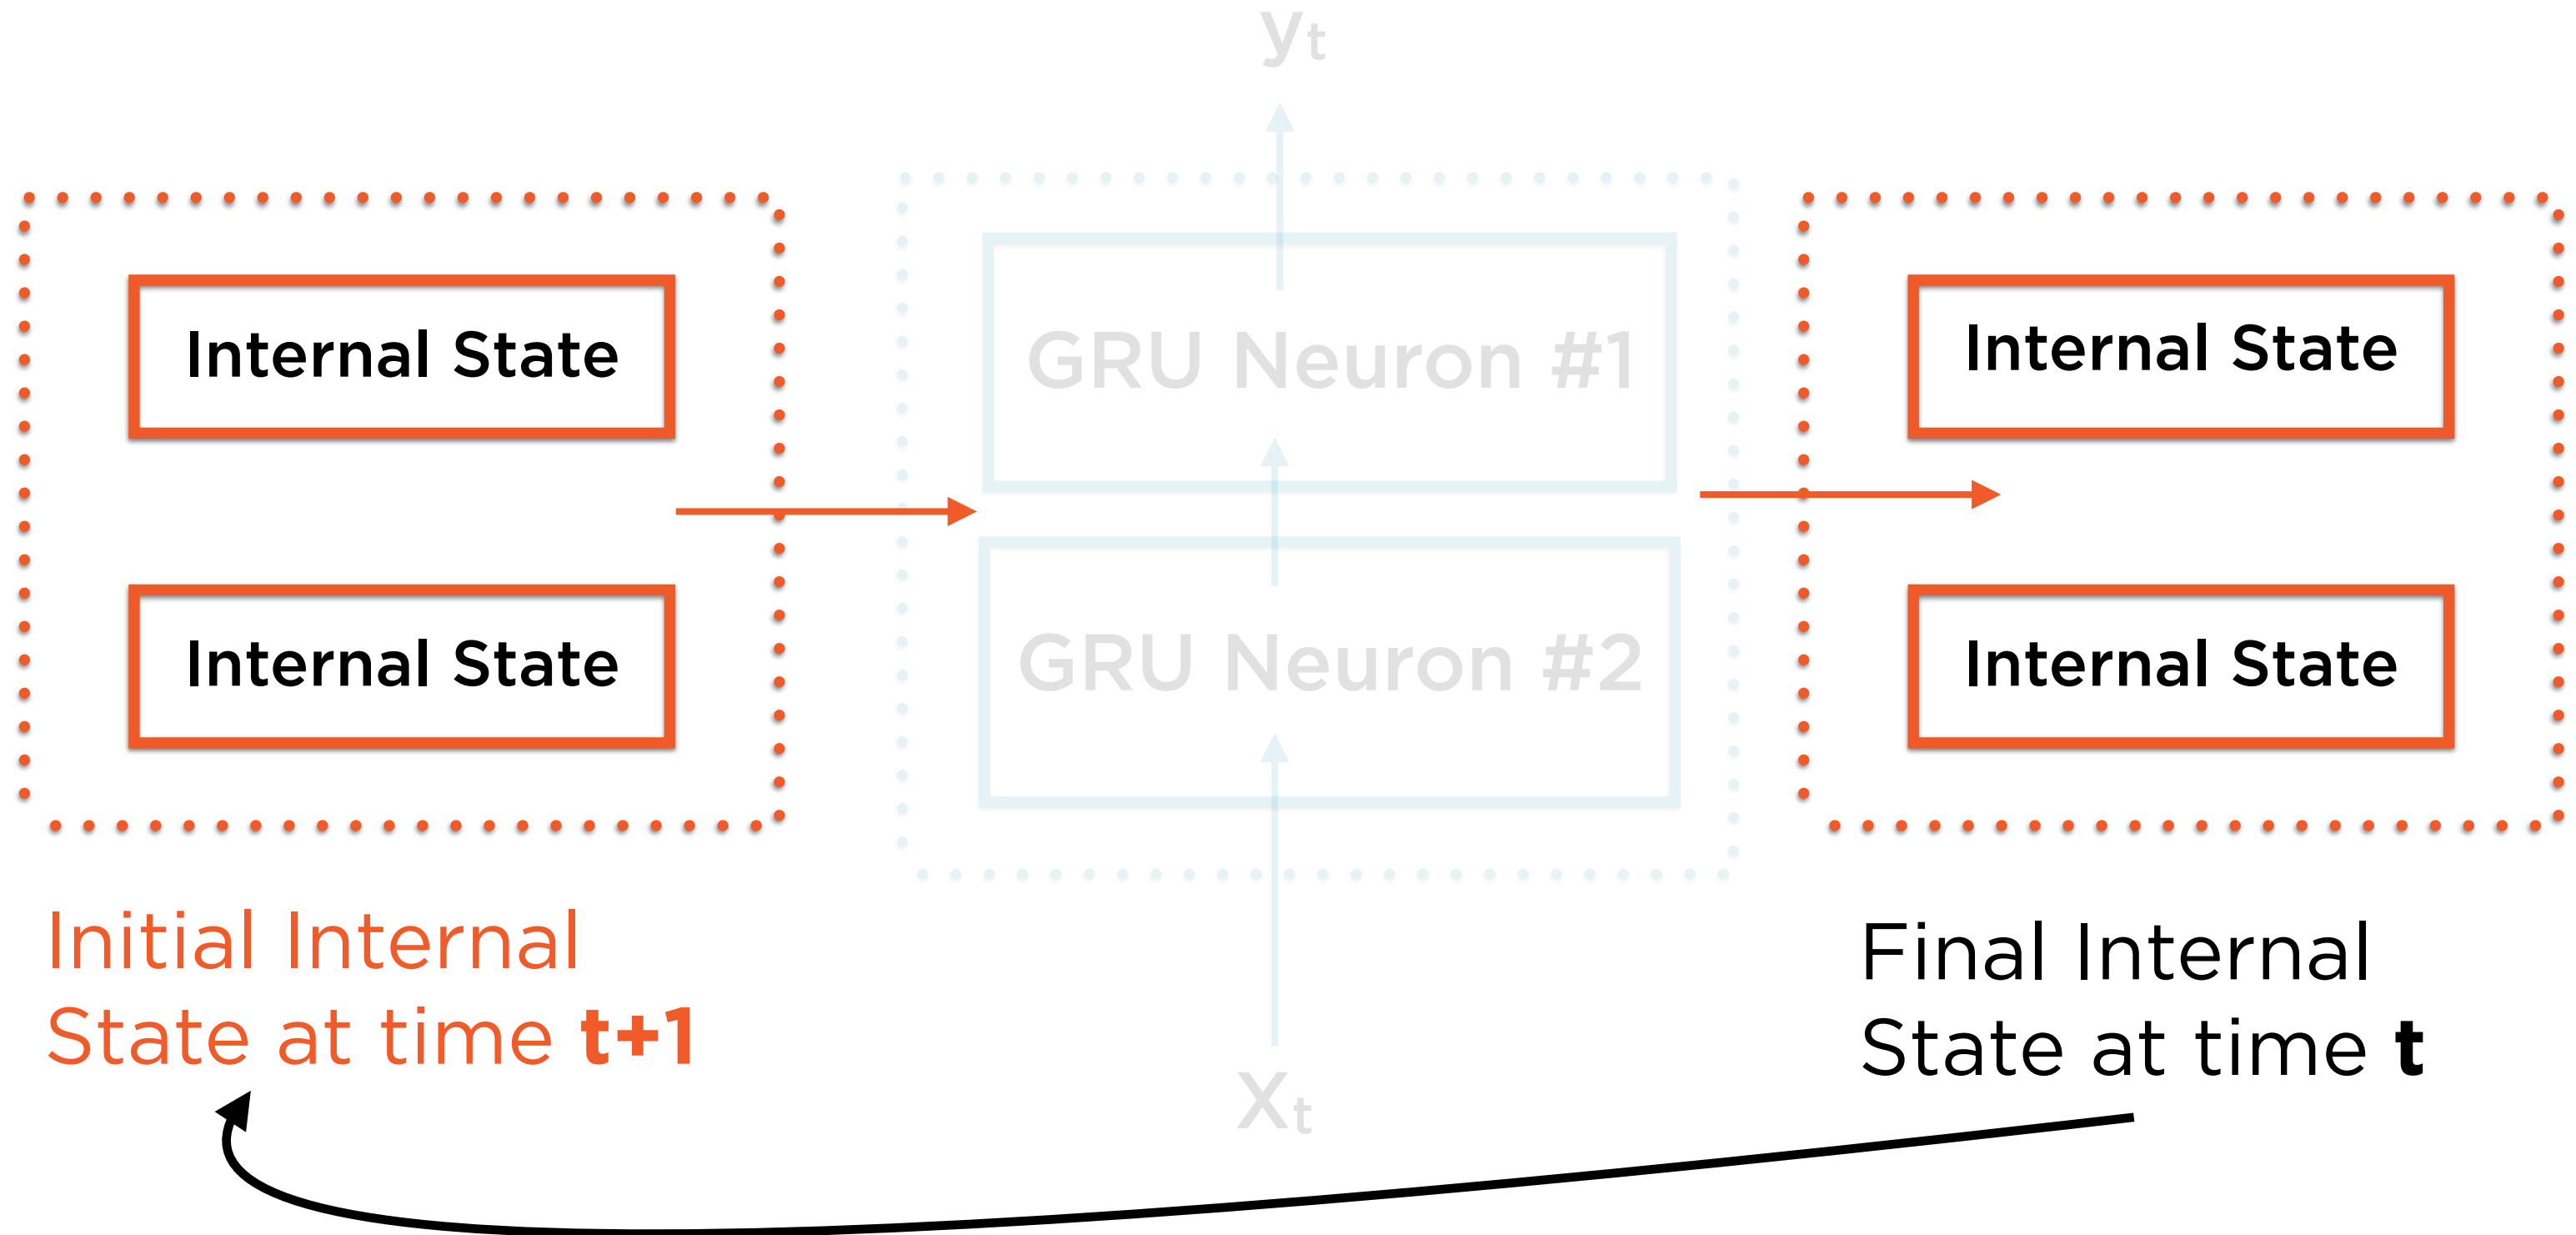
Re-using Internal State



Re-using Internal State



Re-using Internal State



Character Prediction to Generate Text

Problem statement: Given a sequence,
predict what follows

Training and Prediction



Training

Train RNN with a huge dataset of character sequences



Prediction

Use trained model to generate text by feeding back internal state

Training and Prediction



Training

Train RNN with a huge dataset of character sequences



Prediction

Use trained model to generate text by feeding back internal state

Training Dataset of Technical Papers

```
<?xml version="1.0" encoding="UTF-8"?>
<feed xmlns="http://www.w3.org/2005/Atom">
  <link href="http://arxiv.org/api/query?search_query%3Dall%3Aneural%26id_list%3D%26start%3D0%26max_results%3D10" rel="self" type="application/atom+xml"/>
  <title type="html">ArXiv Query: search_query=all:neural&id_list=&start=0&max_results=10</title>
  <id>http://arxiv.org/api/8JSYXq/Mw8td9LFNvtk3N0qiBr4</id>
  <updated>2018-01-08T00:00:00-05:00</updated>
  <opensearch:totalResults xmlns:opensearch="http://a9.com/-/spec/opensearch/1.1/">31759</opensearch:totalResults>
  <opensearch:startIndex xmlns:opensearch="http://a9.com/-/spec/opensearch/1.1/">0</opensearch:startIndex>
  <opensearch:itemsPerPage xmlns:opensearch="http://a9.com/-/spec/opensearch/1.1/">10</opensearch:itemsPerPage>
  <entry>
    <id>http://arxiv.org/abs/cs/0504056v1</id>
    <updated>2005-04-13T13:59:55Z</updated>
    <published>2005-04-13T13:59:55Z</published>
    <title>Self-Organizing Multilayered Neural Networks of Optimal Complexity</title>
    <summary> The principles of self-organizing the neural networks of optimal complexity
is considered under the unrepresentative learning set. The method of
self-organizing the multi-layered neural networks is offered and used to train
the logical neural networks which were applied to the medical diagnostics.
</summary>
    <author>
      <name>V. Schetinin</name>
    </author>
    <link href="http://arxiv.org/abs/cs/0504056v1" rel="alternate" type="text/html"/>
    <link title="pdf" href="http://arxiv.org/pdf/cs/0504056v1" rel="related" type="application/pdf"/>
    <arxiv:primary_category xmlns:arxiv="http://arxiv.org/schemas/atom" term="cs.NE" scheme="http://arxiv.org/schemas/atom"/>
    <category term="cs.NE" scheme="http://arxiv.org/schemas/atom"/>
    <category term="cs.AI" scheme="http://arxiv.org/schemas/atom"/>
  </entry>
  <entry>
    <id>http://arxiv.org/abs/cs/0608073v1</id>
    <updated>2006-08-18T08:28:23Z</updated>
    <published>2006-08-18T08:28:23Z</published>
    <title>Parametrical Neural Networks and Some Other Similar Architectures</title>
    <summary> A review of works on associative neural networks accomplished during last
four years in the Institute of Optical Neural Technologies RAS is given. The
presentation is based on description of parametrical neural networks (PNN). For
today PNN have record recognizing characteristics (storage capacity, noise
immunity and speed of operation). Presentation of basic ideas and principles is
accentuated.
```



```
<opensearch:itemsPerPage xmlns:opensearch="http://a9.com/-/spec/opensearch/1.1/"
```

```
<entry>
```

```
<id>http://arxiv.org/abs/cs/0504056v1</id>
```

```
<updated>2005-04-13T13:59:55Z</updated>
```

```
<published>2005-04-13T13:59:55Z</published>
```

```
<title>Self-Organizing Multilayered Neural Networks of Optimal Complexity</tit
```

```
<summary> The principles of self-organizing the neural networks of optimal co
```

```
is considered under the unrepresentative learning set. The method of
```

```
self-organizing the multi-layered neural networks is offered and used to train
```

```
the logical neural networks which were applied to the medical diagnostics.
```

```
</summary>
```

```
<author>
```

```
<name>V. Schetinin</name>
```

```
</author>
```

```
<link href="http://arxiv.org/abs/cs/0504056v1" rel="alternate" type="text/html
```

```
<link title="pdf" href="http://arxiv.org/pdf/cs/0504056v1" rel="related" type=
```

```
<arxiv:primary_category xmlns:arxiv="http://arxiv.org/schemas/atom" term="cs.N
```

```
<category term="cs.NE" scheme="http://arxiv.org/schemas/atom"/>
```

```
<category term="cs.AI" scheme="http://arxiv.org/schemas/atom"/>
```

```
</entry>
```

```
<entry>
```

```
<id>http://arxiv.org/abs/cs/0608073v1</id>
```

```
<updated>2006-08-18T08:28:23Z</updated>
```

```
<published>2006-08-18T08:28:23Z</published>
```

Sliding Window in Training

Create window of 50 characters



The quick brown fox jumps over the lazy dog

Sliding Window in Training

Create window of 50 characters



The quick brown fox jumps o

Sliding Window in Training

This window is our sequence length



The quick brown fox jumps o

Sliding Window in Training


Use 0:48 as feature

1




The quick brown fox jumps o

Sliding Window in Training



Training Phase


The quick brown fox jumps



he quick brown fox jumps o

Training Phase


The quick brown fox jumps



he quick brown fox jumps o

Training Phase


The quick brown fox jumps



he quick brown fox jumps o

Training Phase

The quick brown fox jumps



he quick brown fox jumps o

Training Phase

The quick brown fox jumps




**Multi-RNN
Predictor**

he quick brown fox jumps o

Training Phase


The **q**uick brown fox jumps



he **u**ick brown fox jumps o

Training Phase

The **u**ck brown fox jumps




he **i**ck brown fox jumps o

Training Phase

The quick brown fox jumps

s



he quick brown fox jumps

o

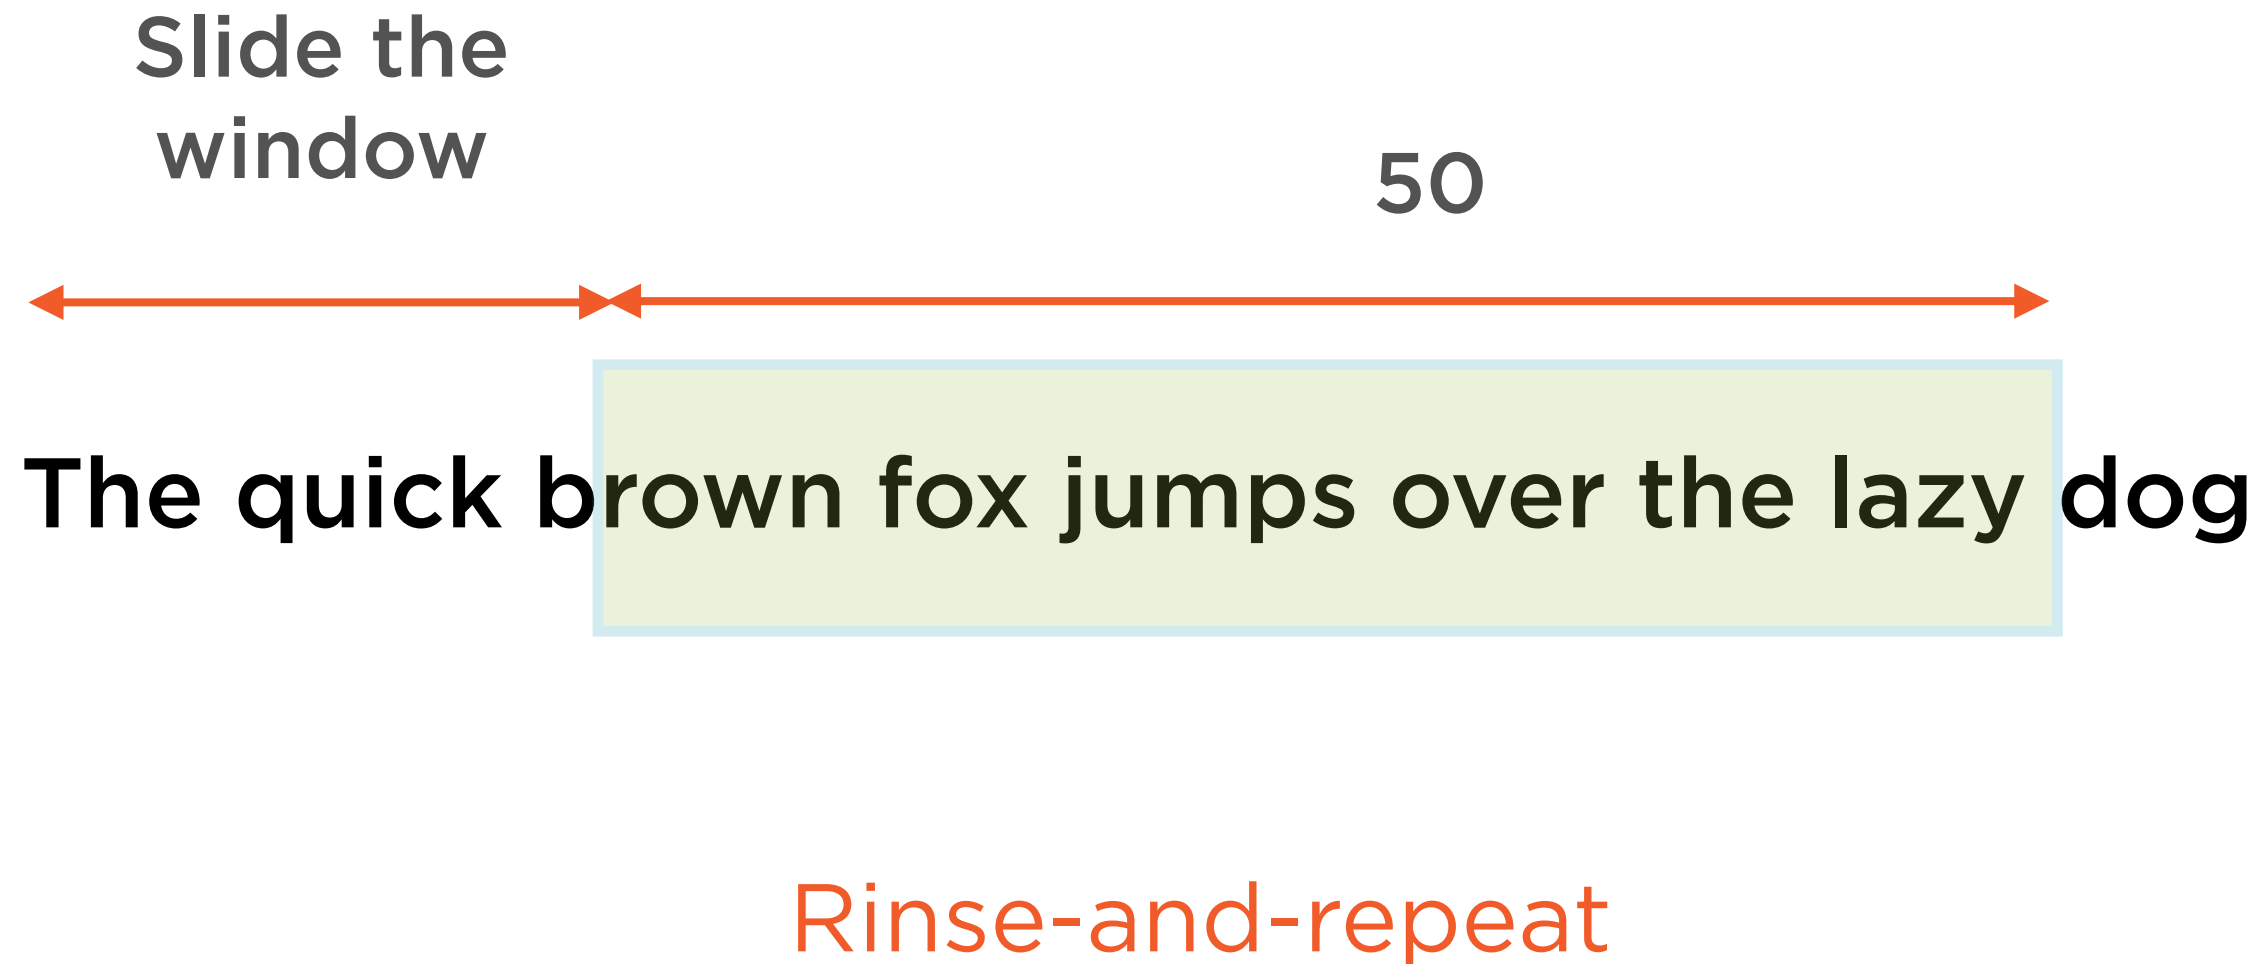
Sliding Window in Training

Create window of 50 characters




The quick brown fox jumps over the lazy dog

Sliding Window in Training



Sliding Window in Training



Rinse-and-repeat

Training and Prediction



Training

Train RNN with a huge dataset of character sequences



Prediction

Use trained model to generate text by feeding back internal state

Training and Prediction



Training

Train RNN with a huge dataset of character sequences




Prediction

Use trained model to generate text by feeding back internal state

Text Prediction

**“The quick
brown fo”**



**“x jumps over
the lazy dog”**

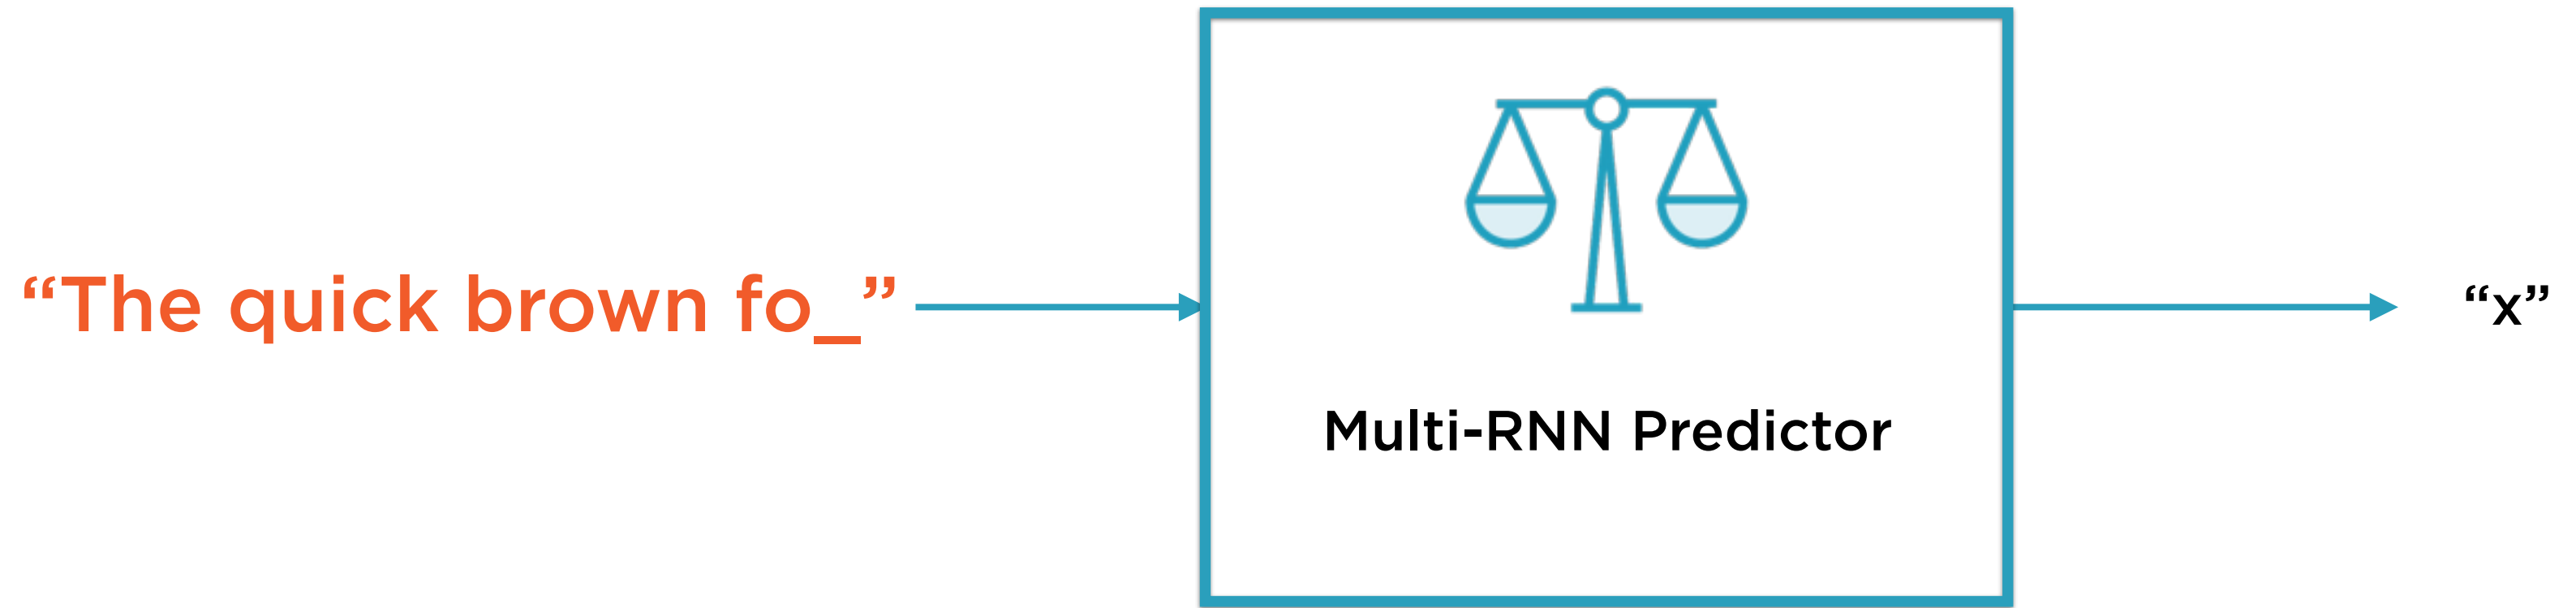
Predict One Character at a Time

“The quick brown fox”

?

A diagram illustrating a sequence prediction task. It features the sentence "The quick brown fox" in bold black text. At the end of the sentence, there is a small, empty rectangular box with an orange border. Below this box, there is a larger rectangular box with an orange border containing a bold black question mark. An orange arrow points from the question mark box up and to the right, ending at the small box at the end of the sentence.

Predict One Character at a Time



Predict One Character at a Time

“The quick brown fox_”

x



The diagram illustrates a sequence prediction task. At the top, the text "The quick brown fox_" is shown in a bold black font. The underscore character is enclosed in a small orange rectangular box. An orange arrow points from this box down to a larger orange rectangular box below it. Inside this larger box is the letter "x" in a bold black font, representing the predicted character.

Predict One Character at a Time

“The quick brown foxx**”**

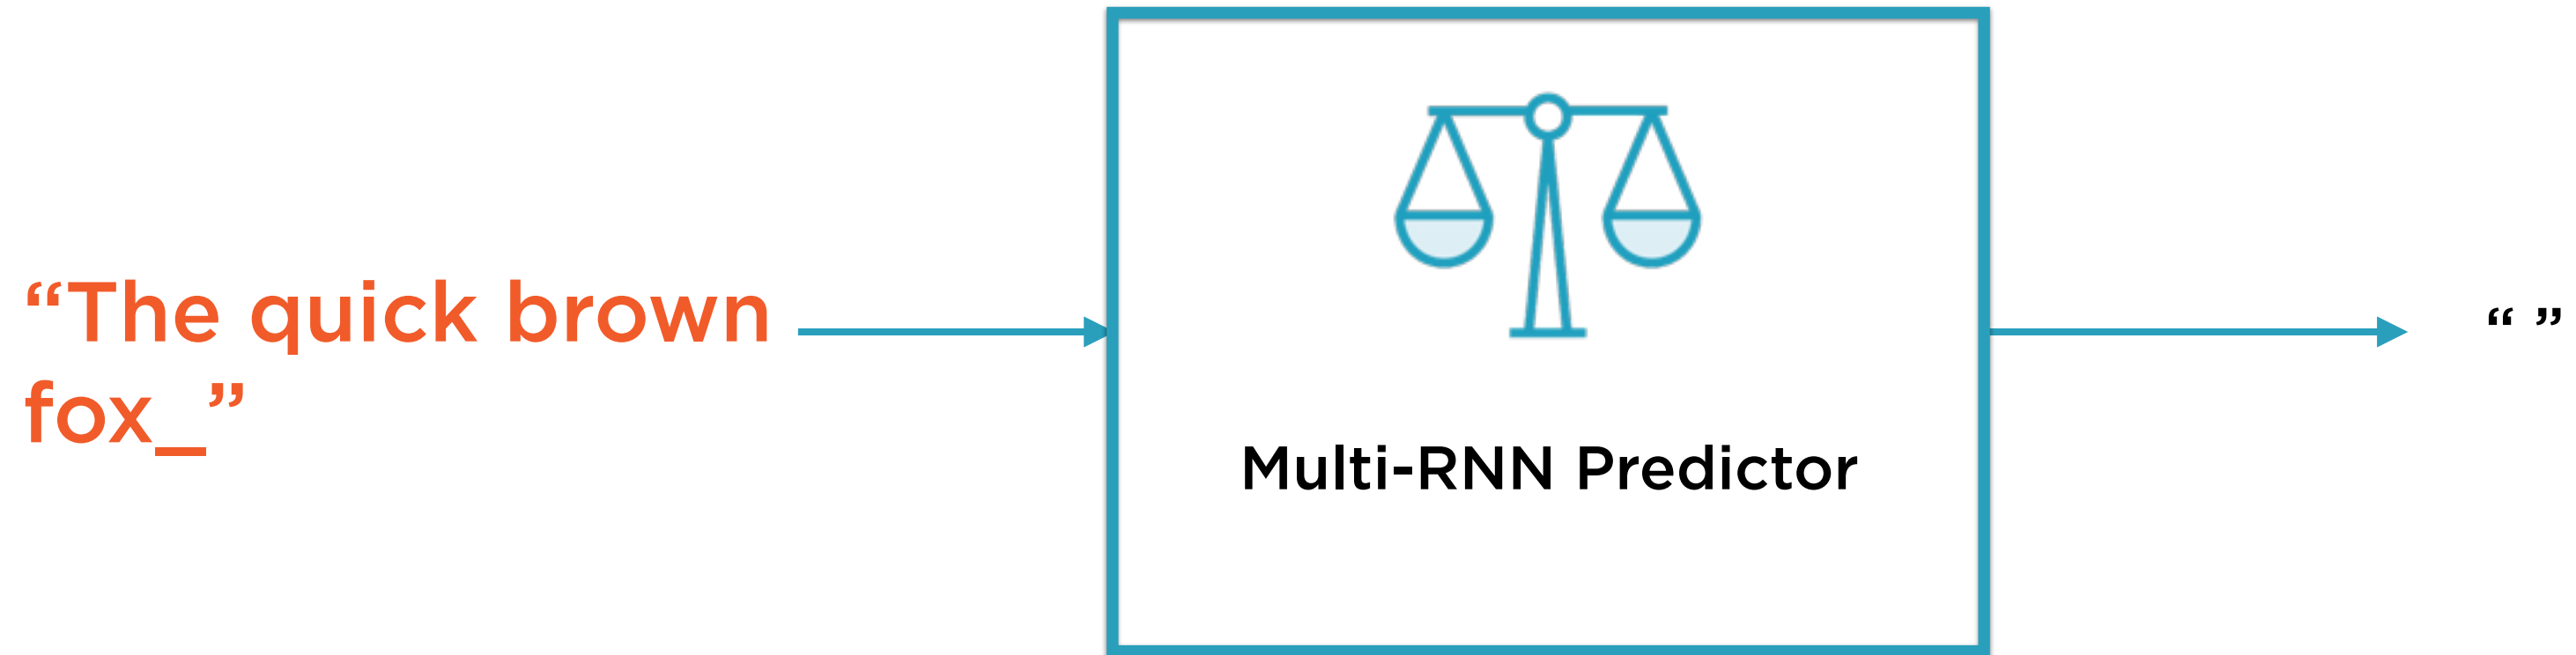
Predict One Character at a Time

“The quick brown fox_”

?



Predict One Character at a Time



Predict One Character at a Time

“The quick brown fox_”

“ ”


A diagram illustrating a sequence prediction task. At the bottom, a rectangular box with an orange border contains the opening and closing quotation marks " ". An orange arrow points from this box diagonally upwards and to the right, ending at a smaller rectangular box with an orange border. This smaller box is positioned at the end of the text "The quick brown fox_" in the line above, which is enclosed in larger quotation marks.

Predict One Character at a Time

“The quick brown fox”

Text Prediction

**“The quick brown fox jumps over the lazy
do_”**



g

Text Prediction

**“The quick brown fox jumps over the lazy
dog”**

Problem statement: Given a sequence,
predict what follows

Solution Outline: Use a RNN to predict
words, character-by-character

Contrasting Architectures

OCR Classification

Classification

Bi-directional RNN

Input 128px image, output character


Text Prediction

Prediction


Multi-RNN Cell

Input character, output next character

OCR: RNN Architecture



Text Prediction: RNN Architecture



Contrasting Architectures

OCR Classification

Input tensor

[batch_size, 14, 128]

Text Prediction


Input tensor

[batch_size, 49, 83]

Larger vocabulary because we include special characters

Text Prediction: Input Tensor for Training

[batch_size, 49, 83]



Contrasting Architectures

OCR Classification

Output tensor

[batch_size, 14, 26]

One-hot with 26 characters


Text Prediction

Output tensor


[batch_size, 49, 83]

One-hot with 83 character

OCR: Output Tensor for Predicted Values



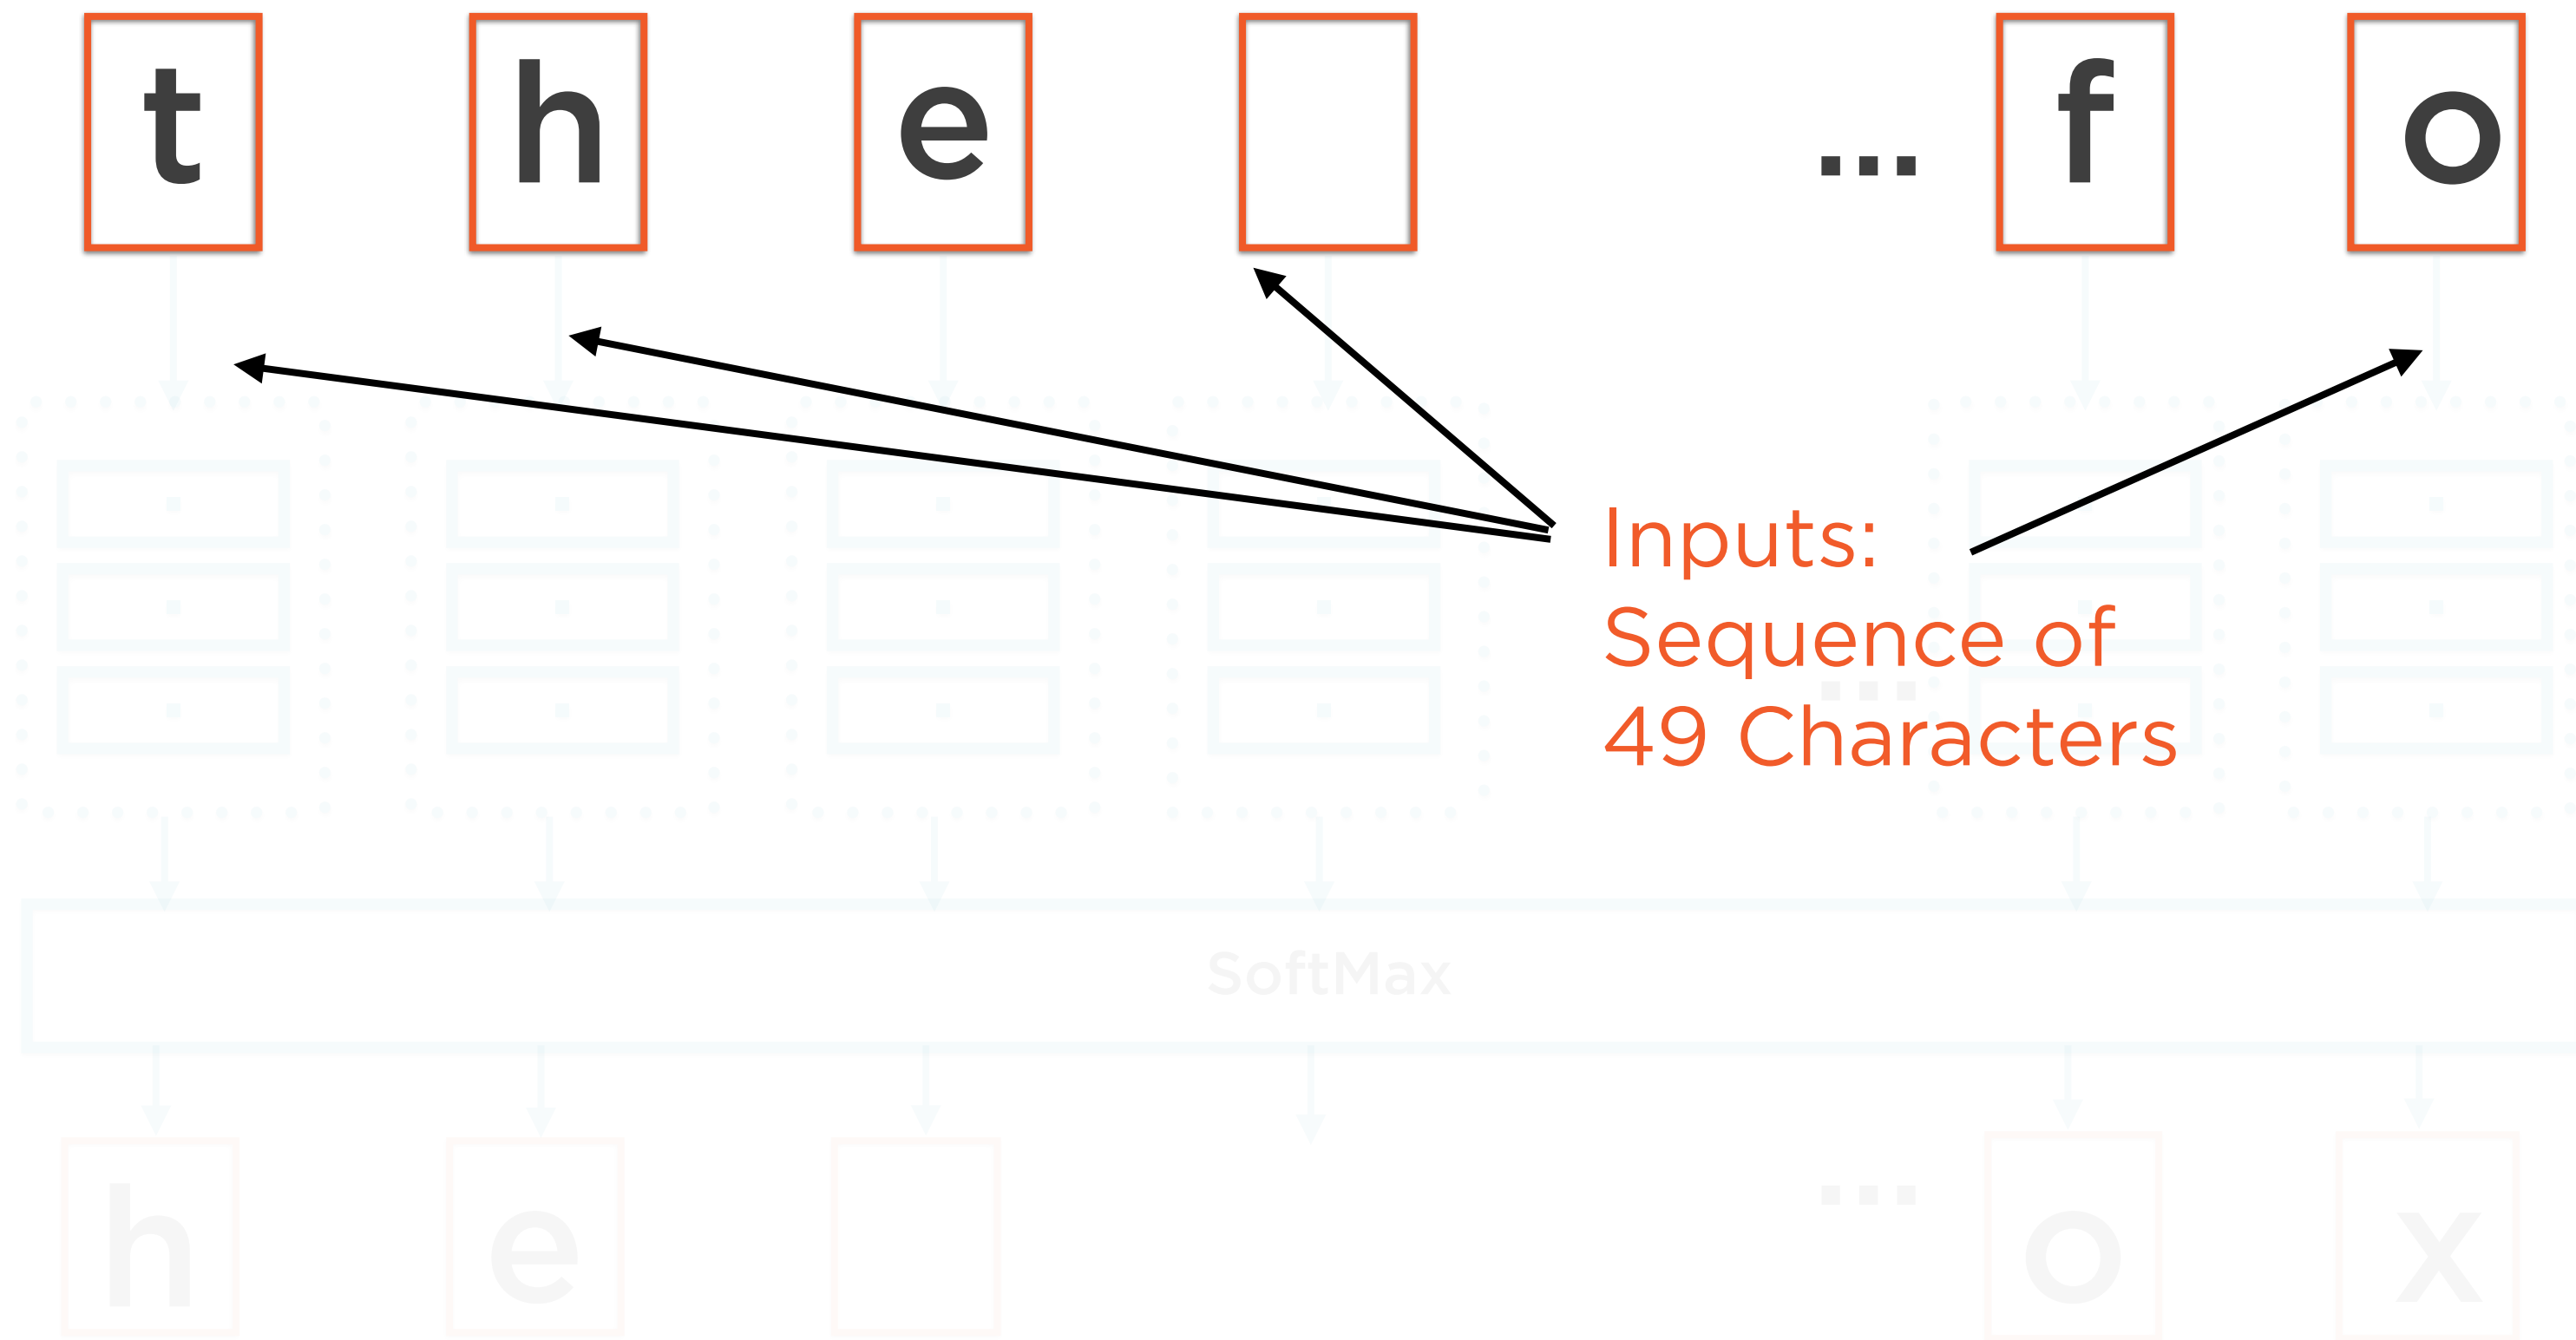
Text Prediction: RNN Architecture



OCR Input Tensors



Text Prediction: RNN Architecture



Contrasting Architectures

OCR Classification

One RNN layer per input image

14 RNN layers

Every bidirectional cell had 300 neurons

Text Prediction


One RNN layer per input character

49 RNN layers

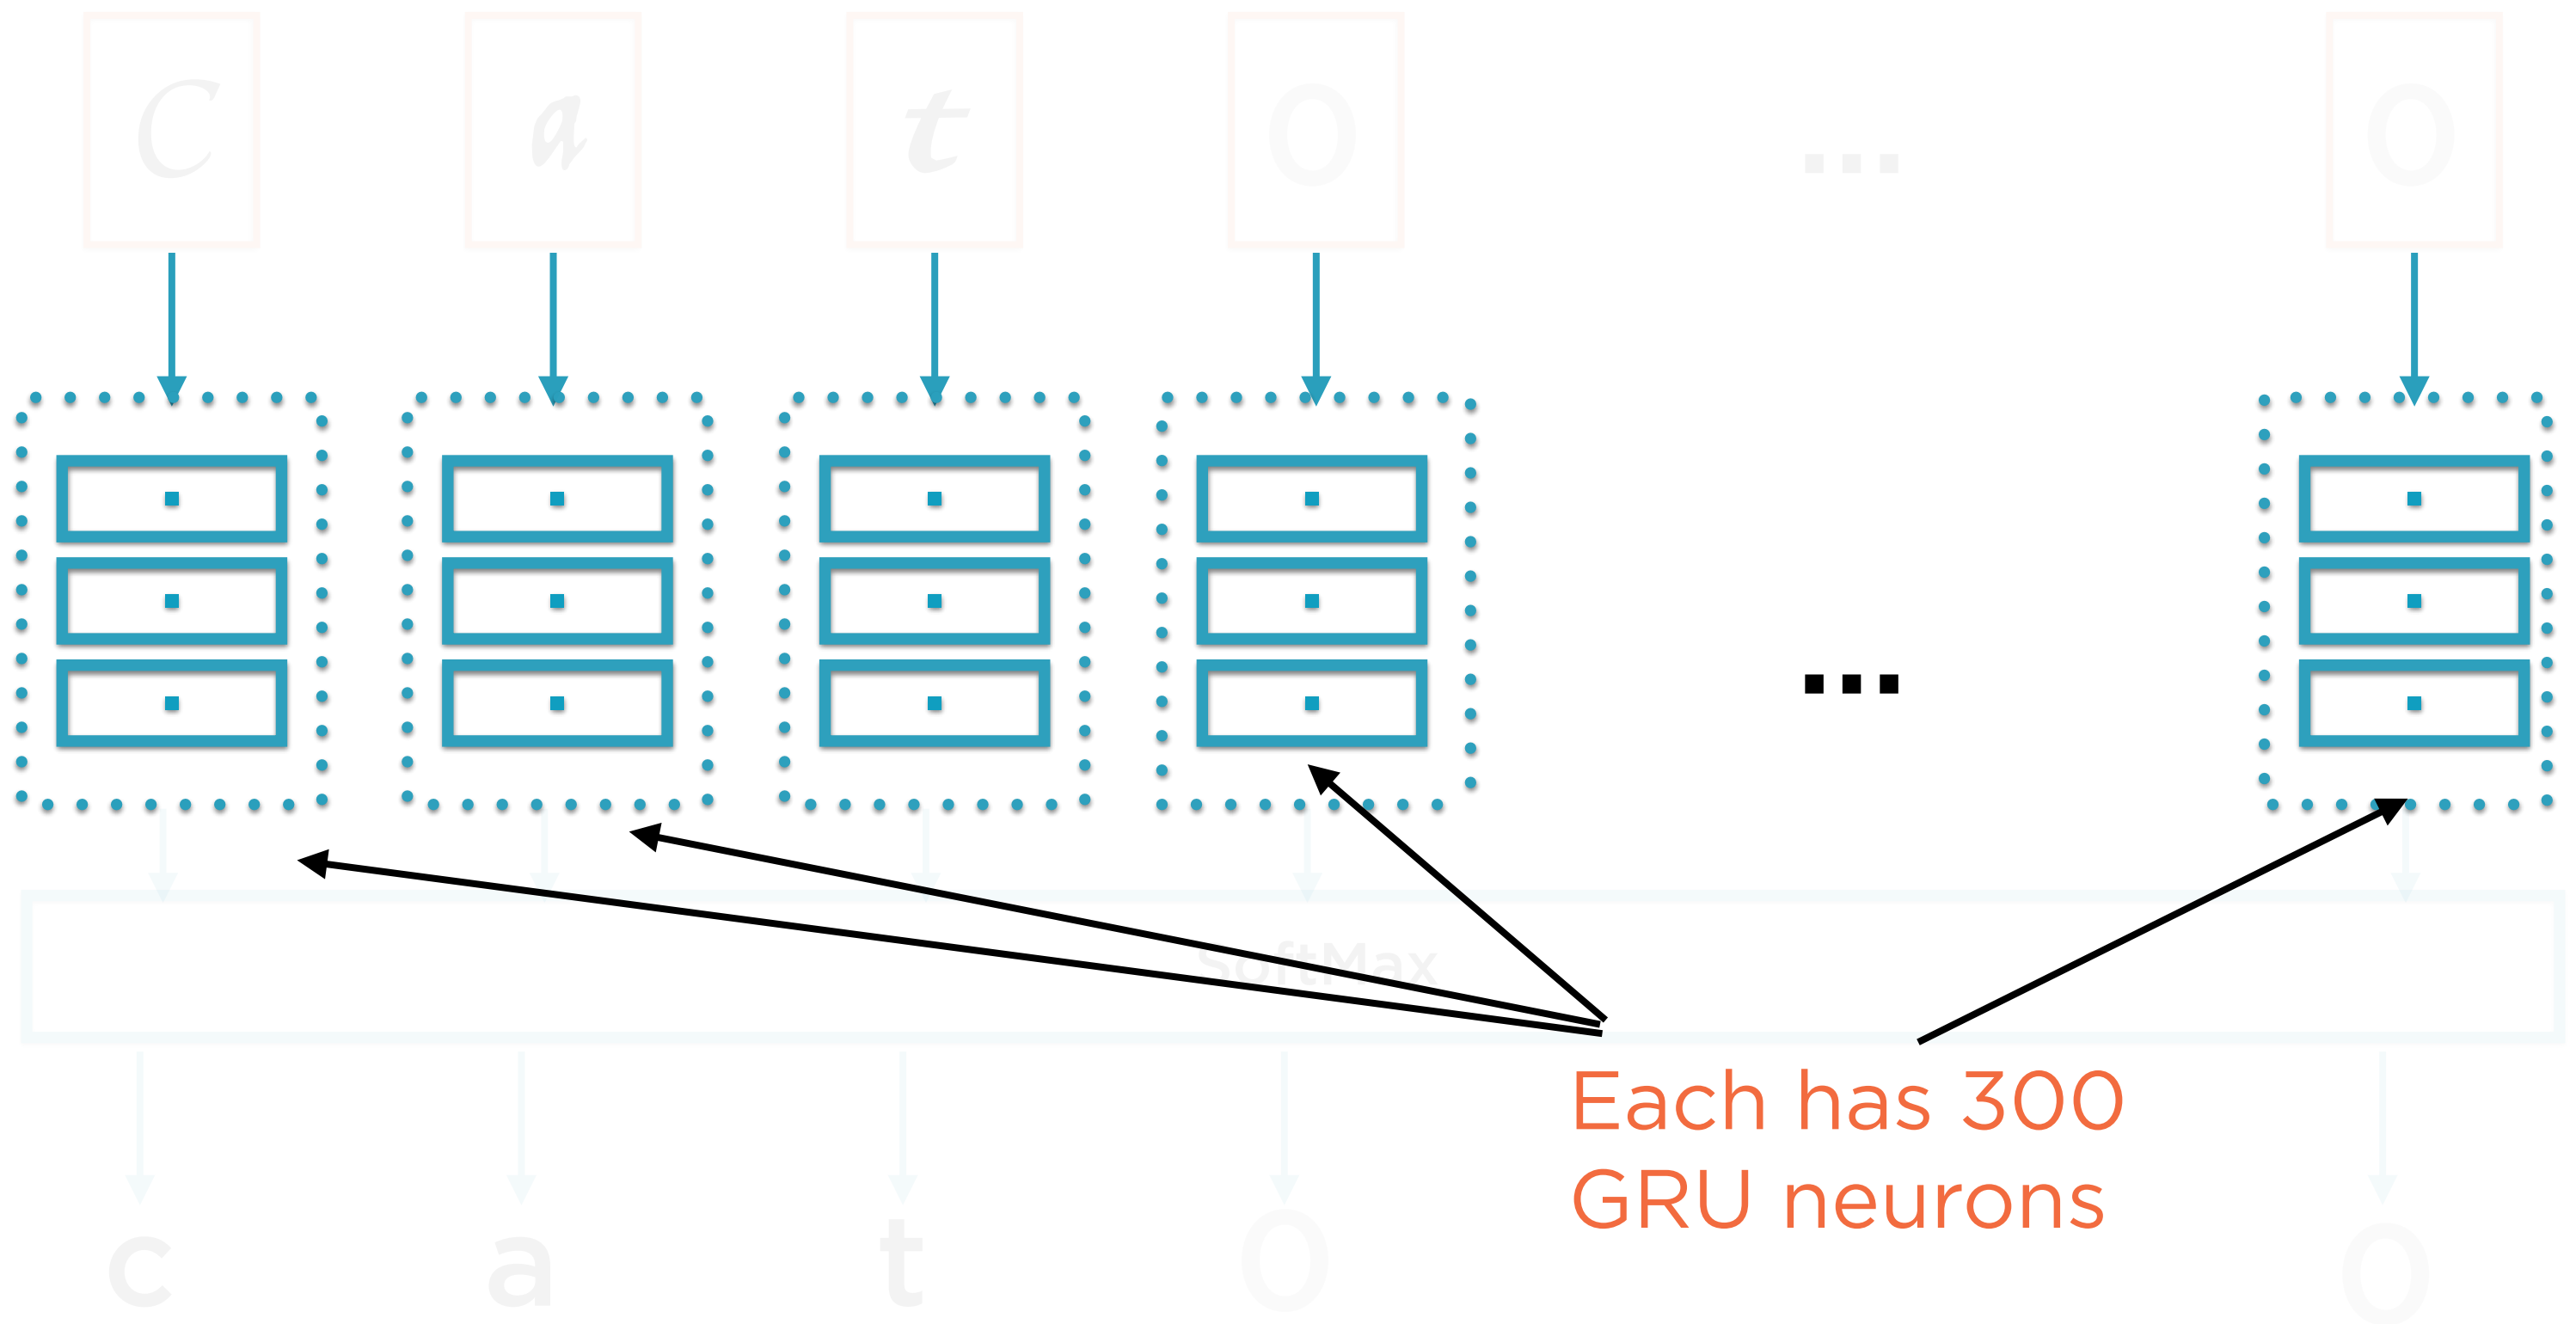
Each multi-RNN cell has 2 GRU cells

Each GRU cell has 200 neurons

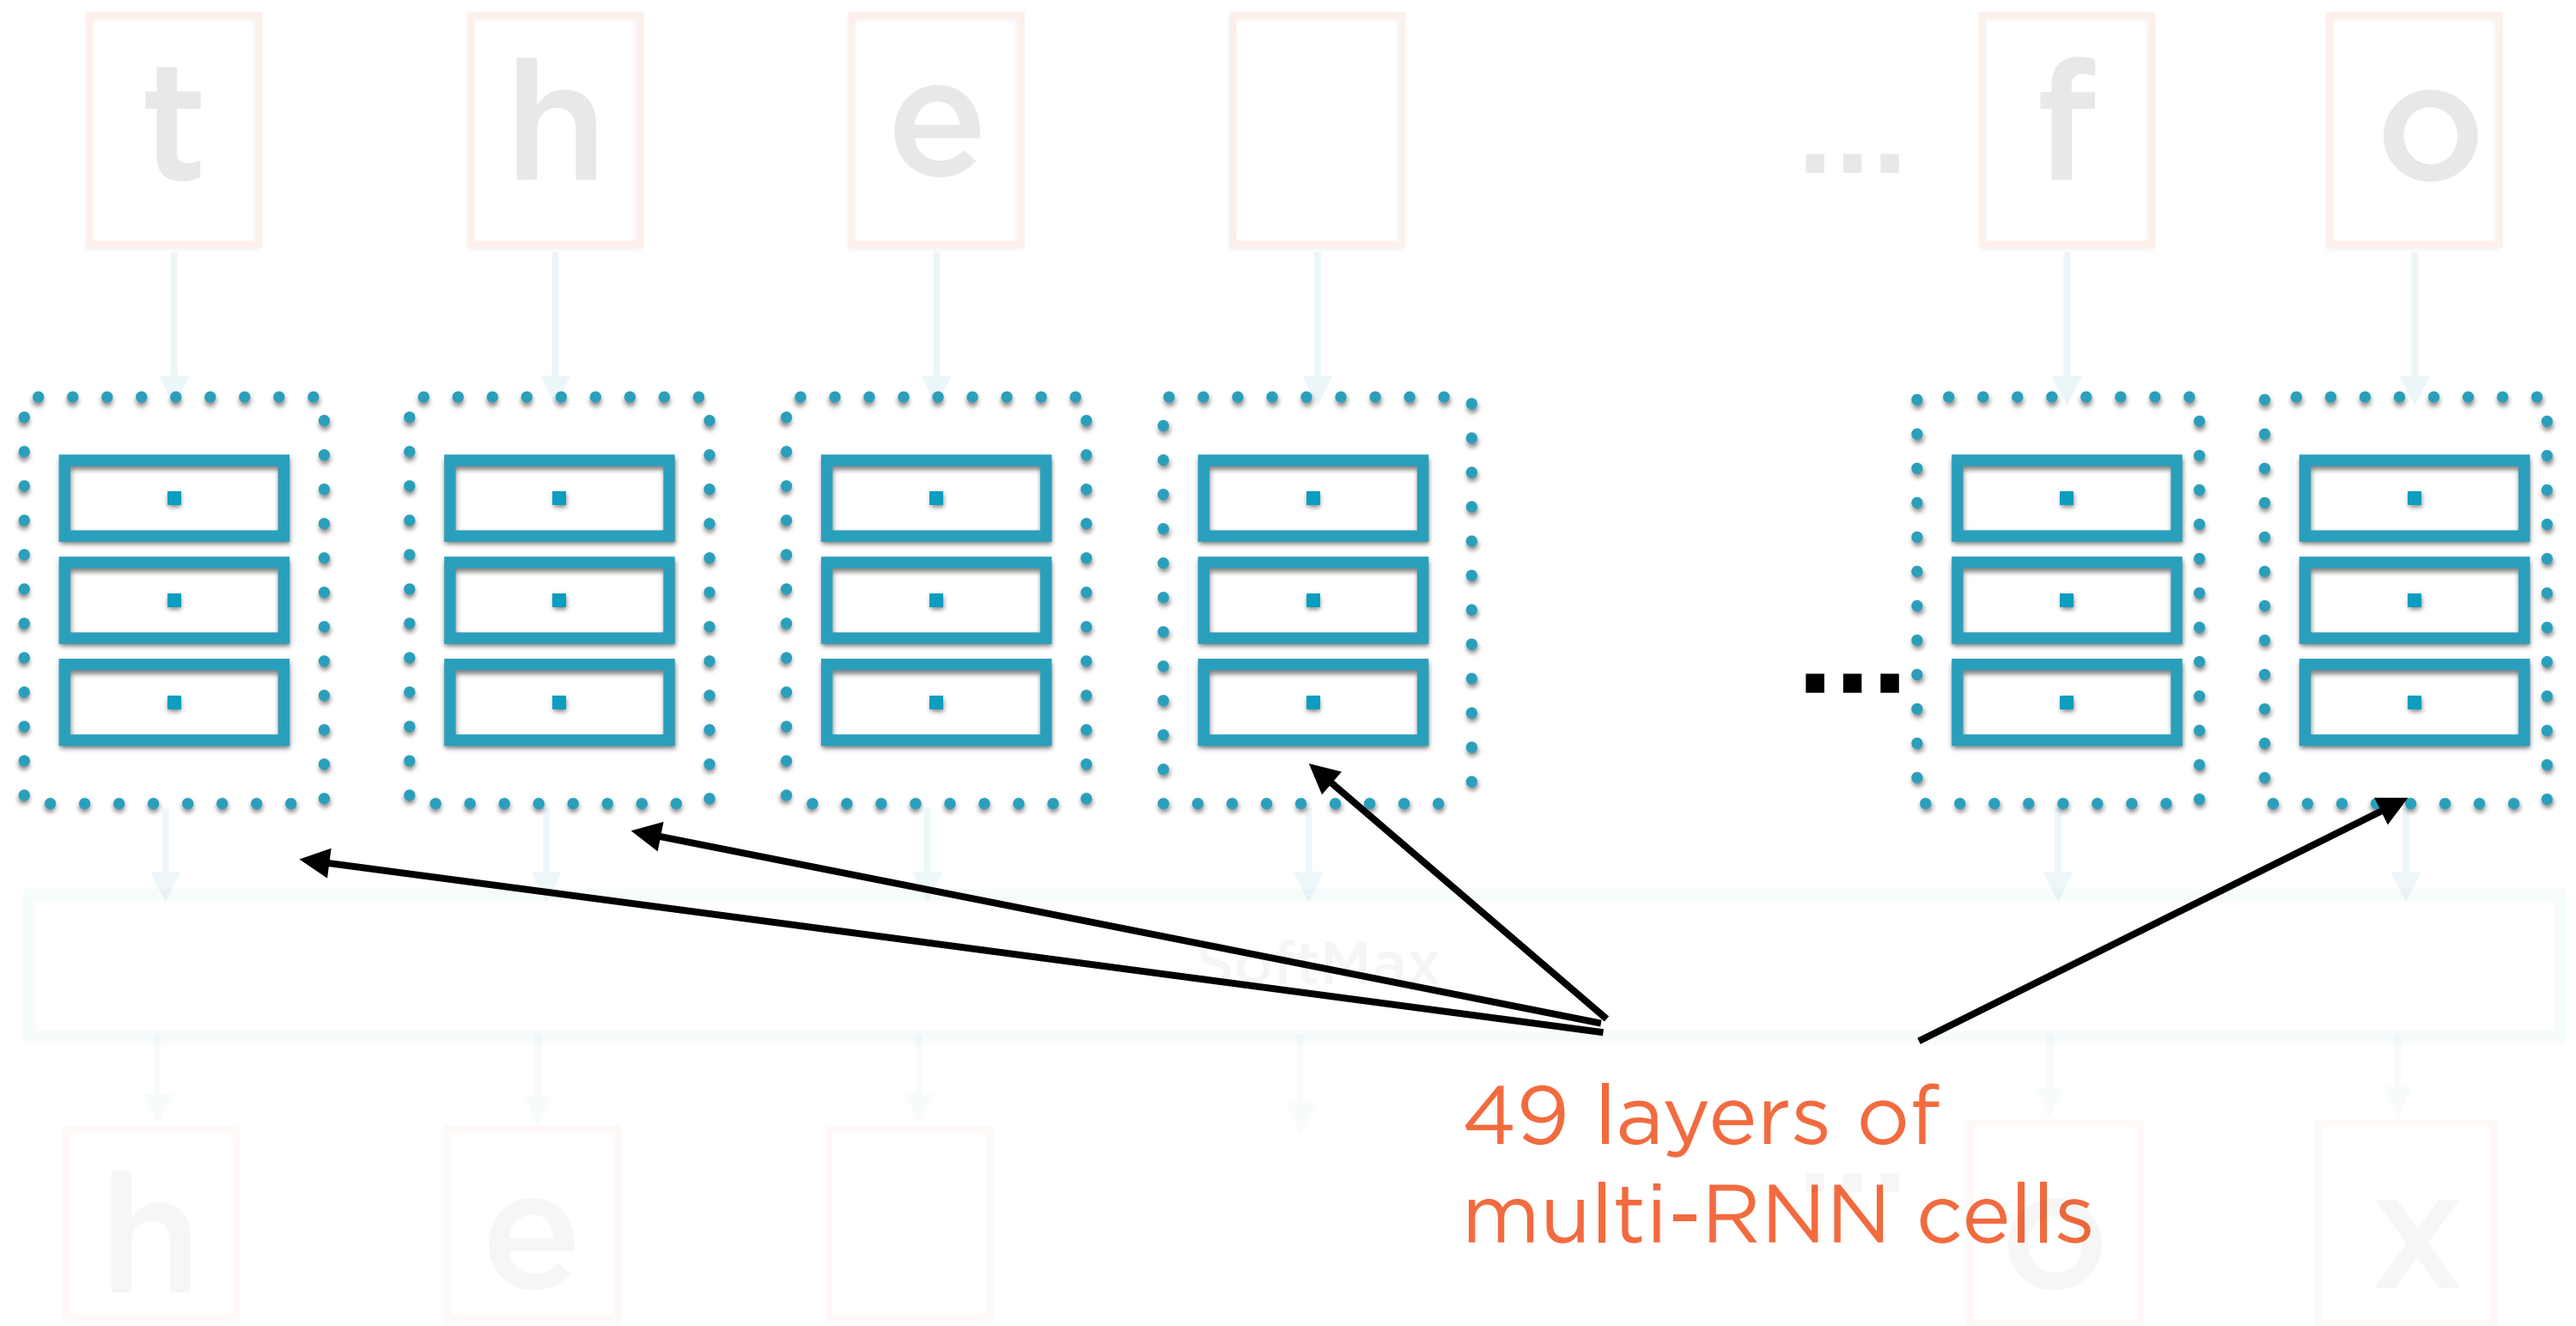
OCR: RNN Architecture



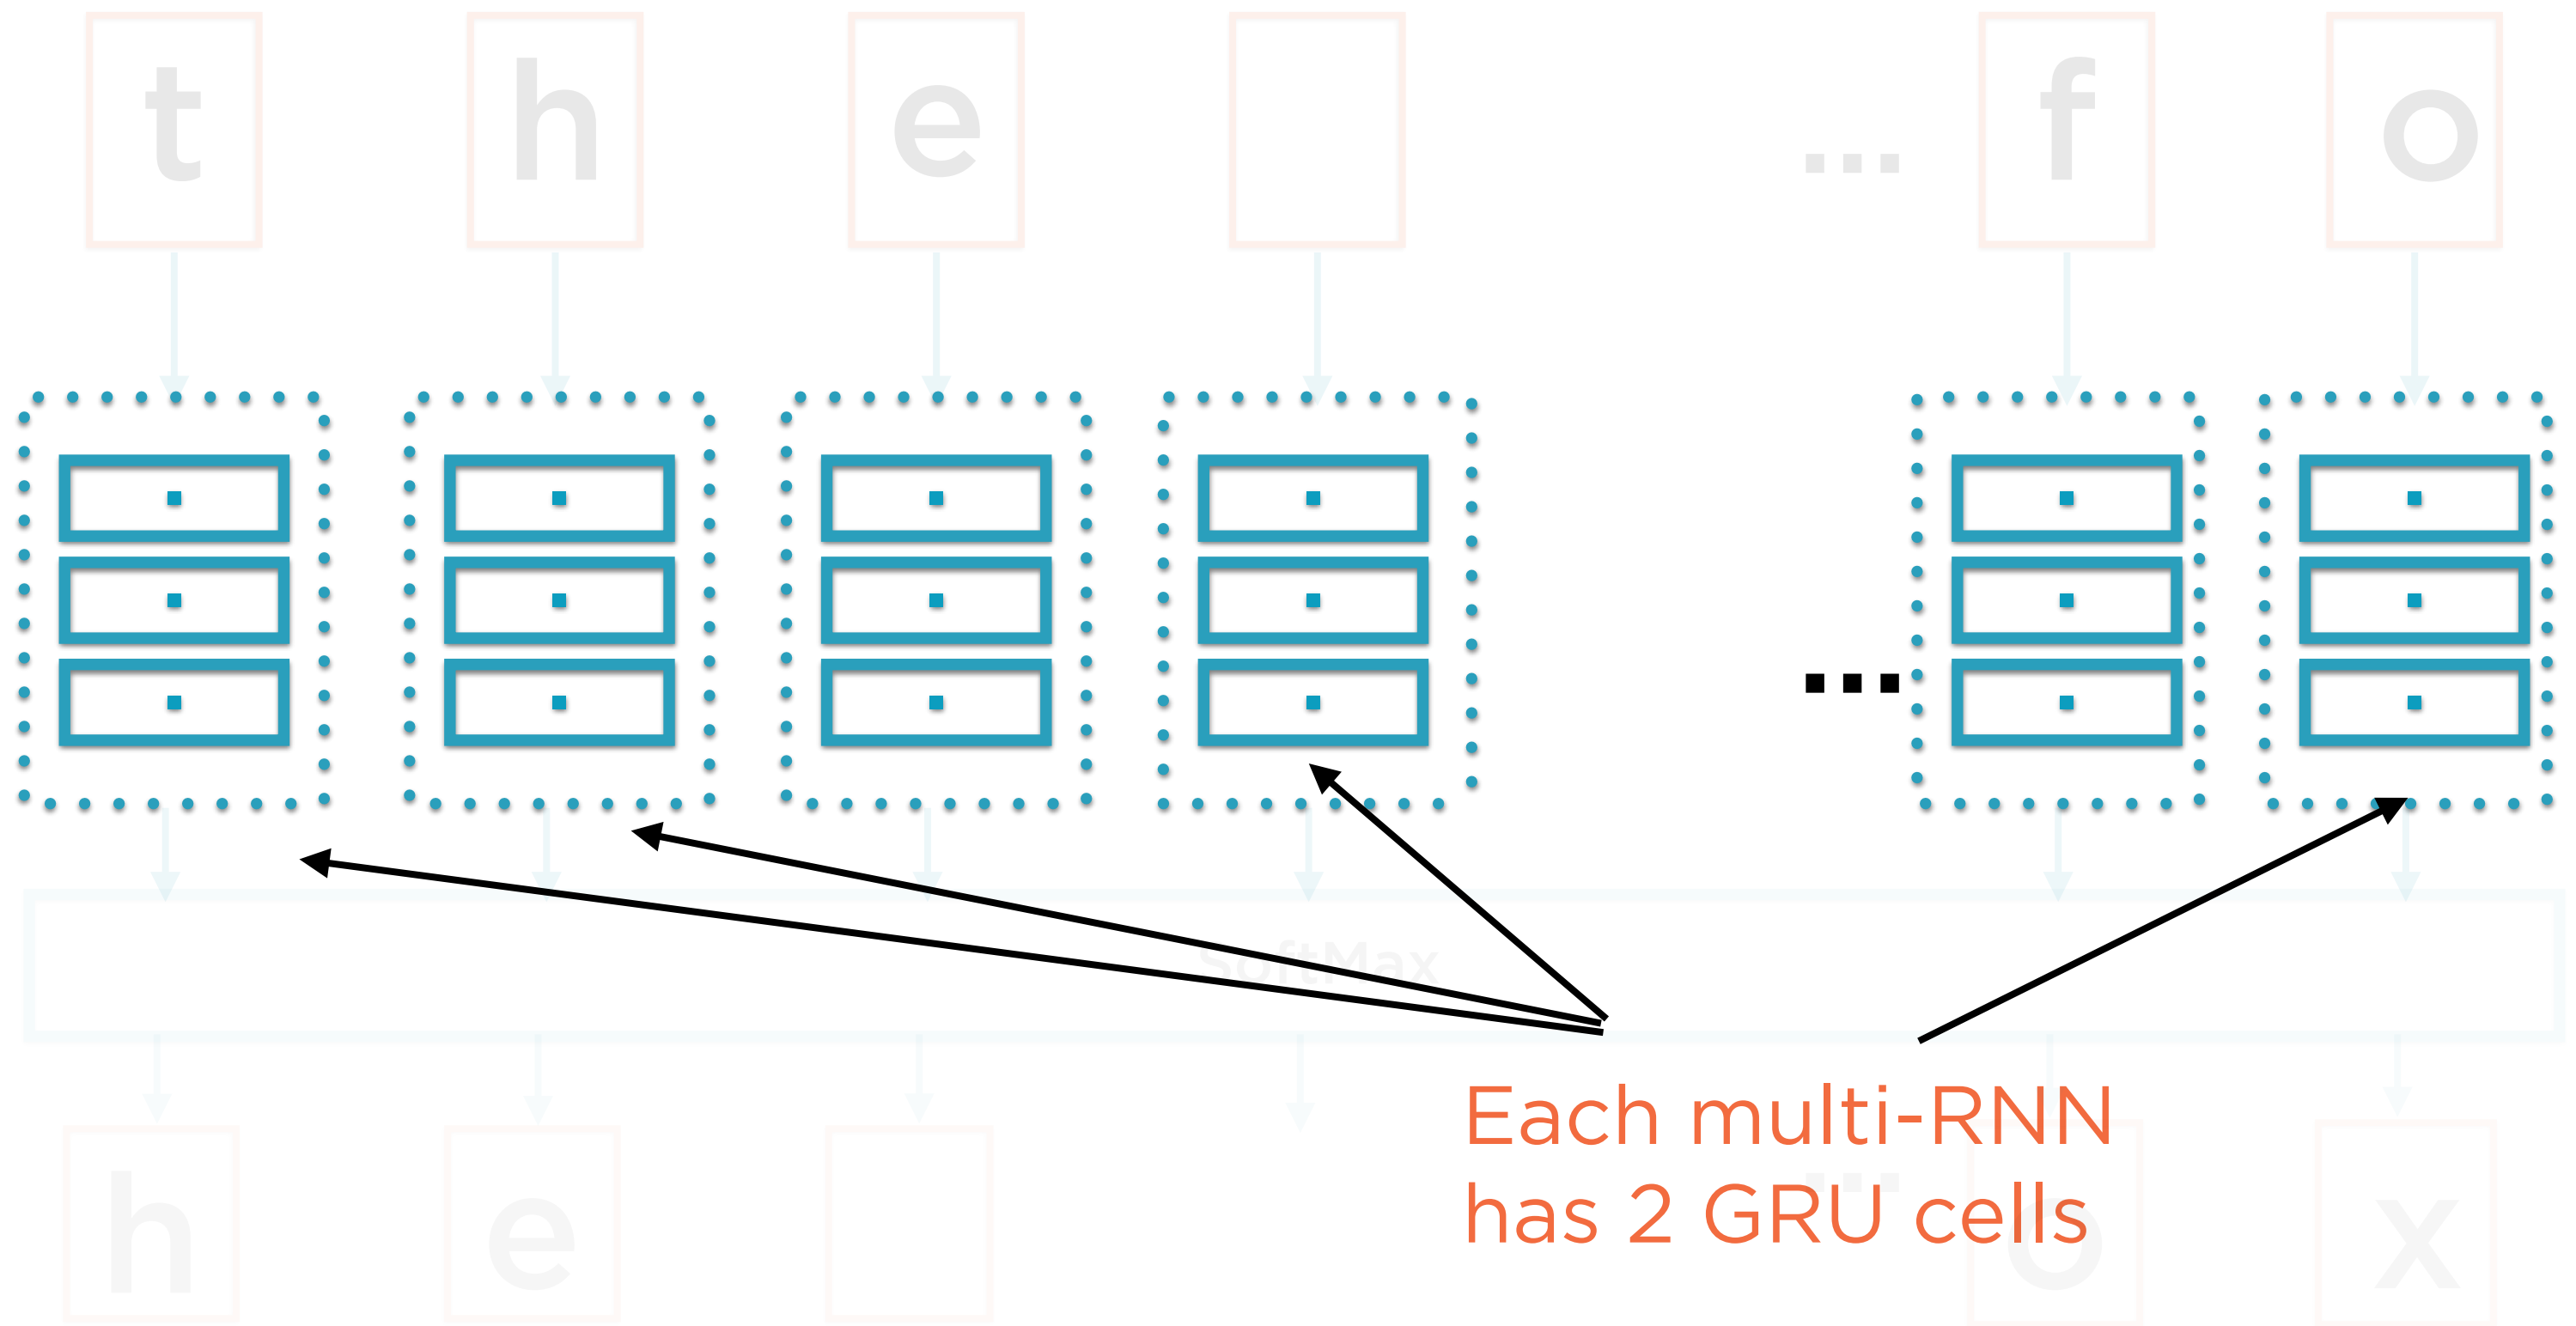
OCR: RNN Architecture



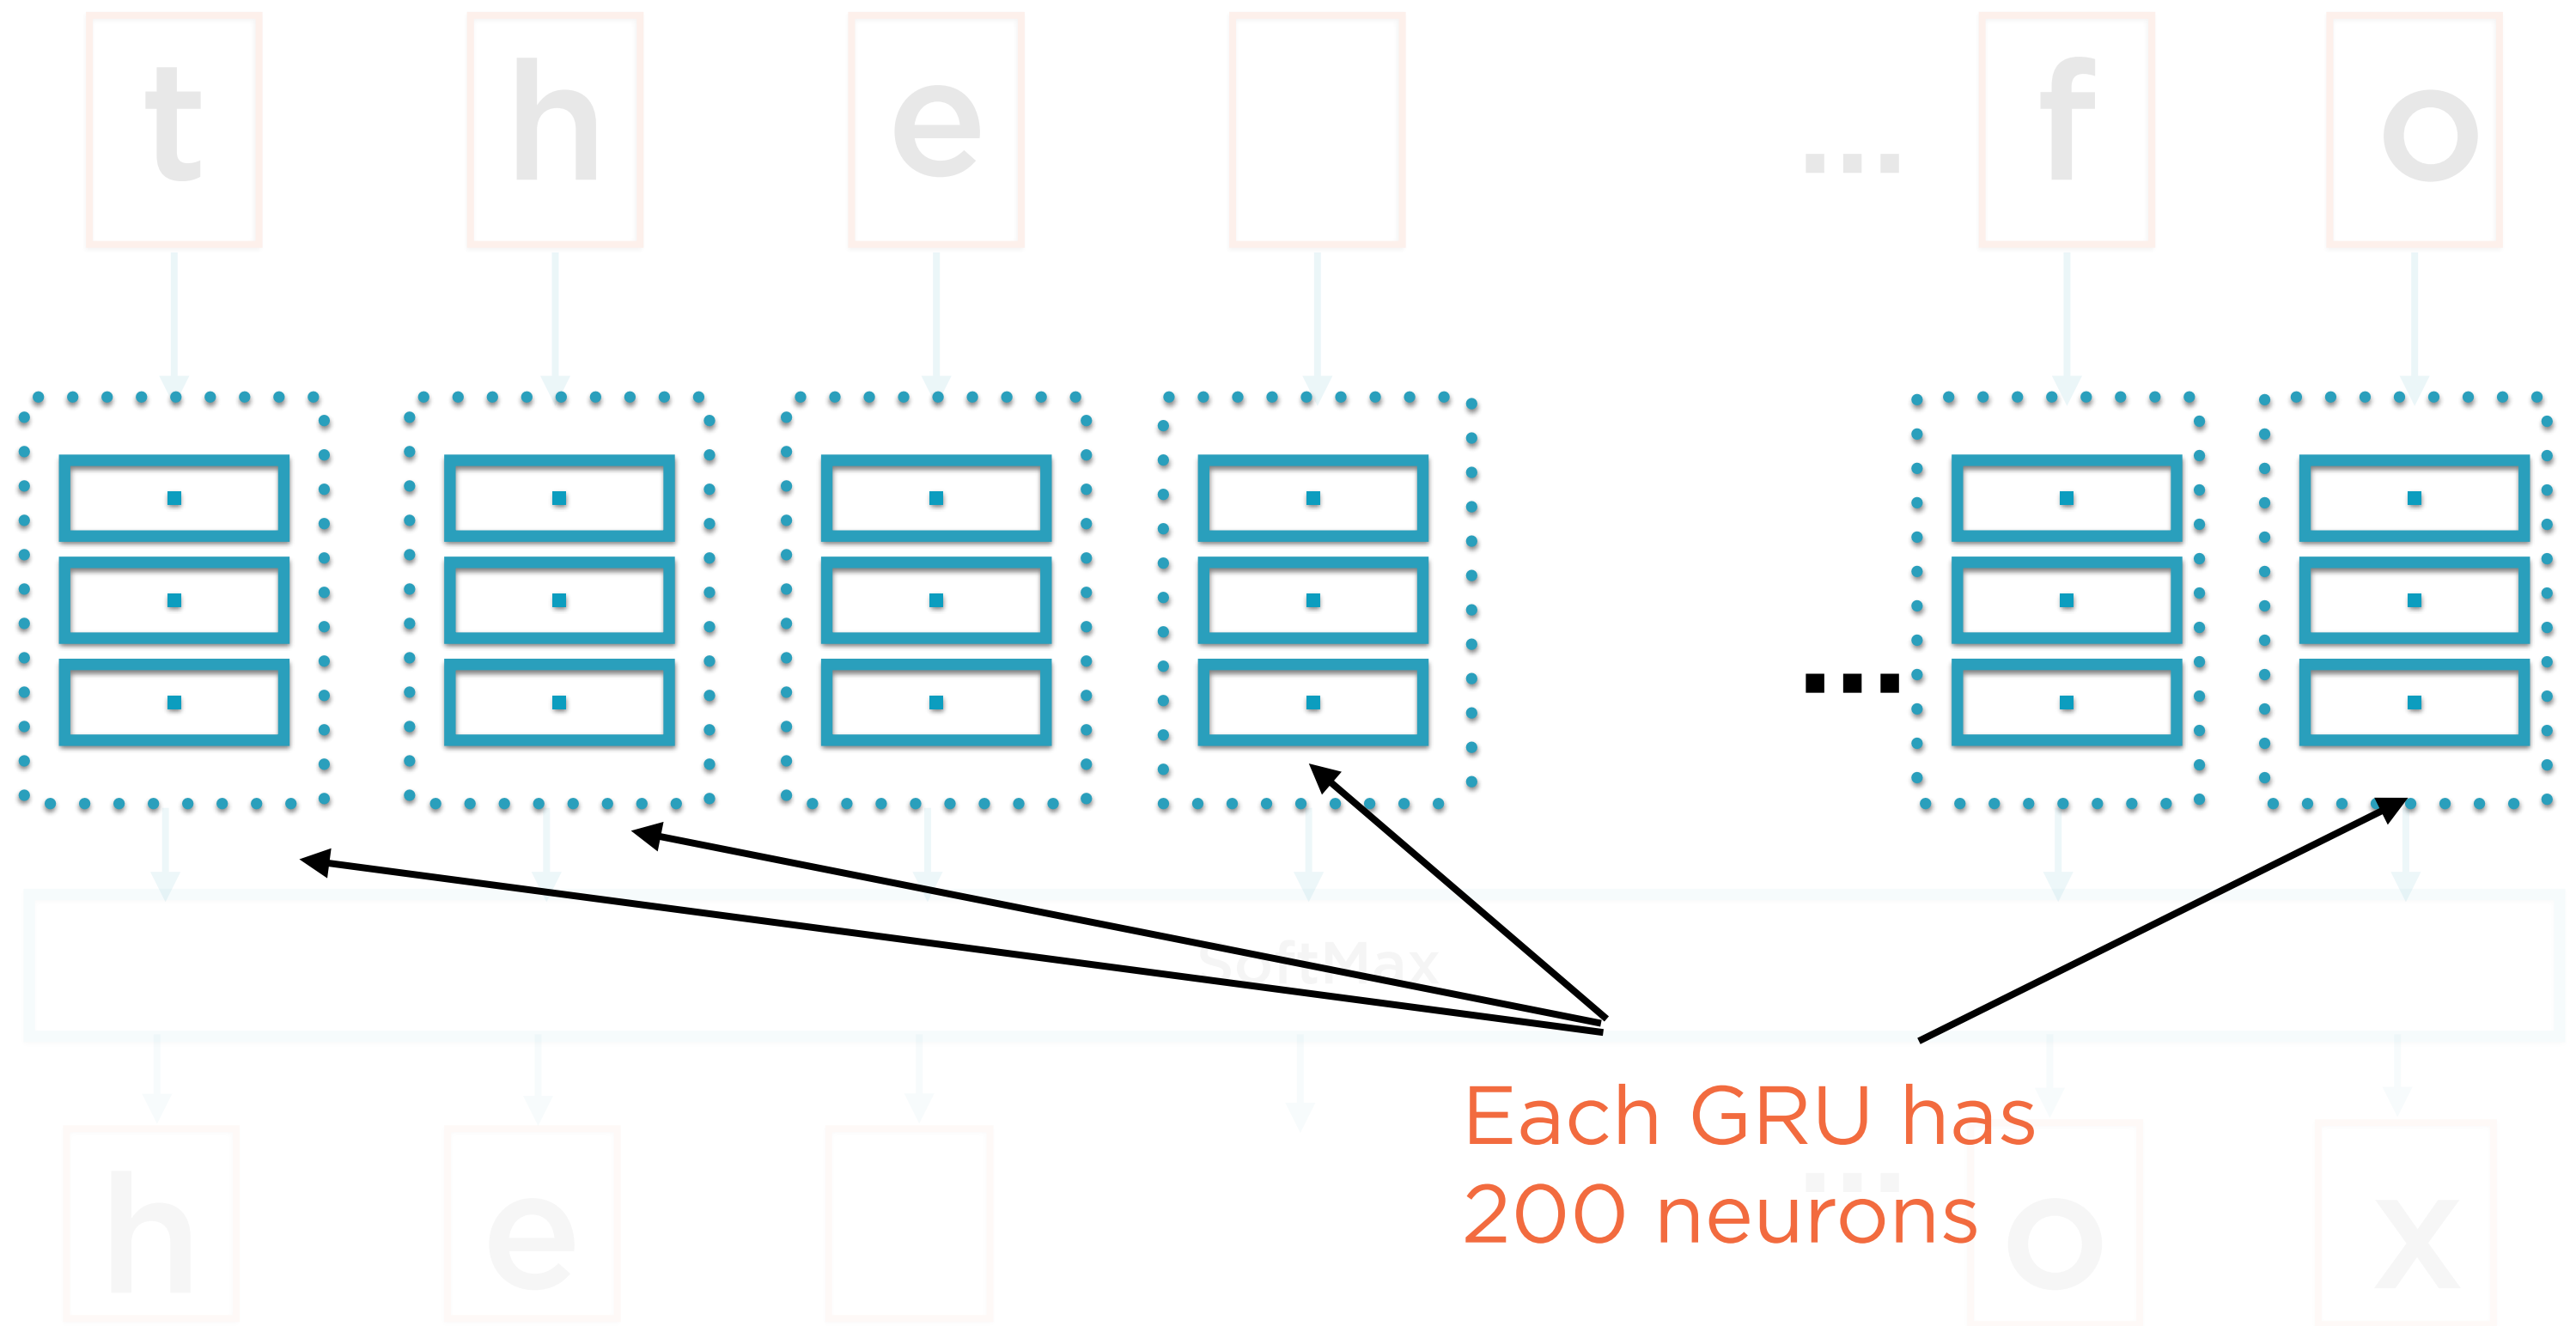
Text Prediction: RNN Architecture



Text Prediction: RNN Architecture



Text Prediction: RNN Architecture



Contrasting Architectures

OCR Classification

Shared Softmax layer

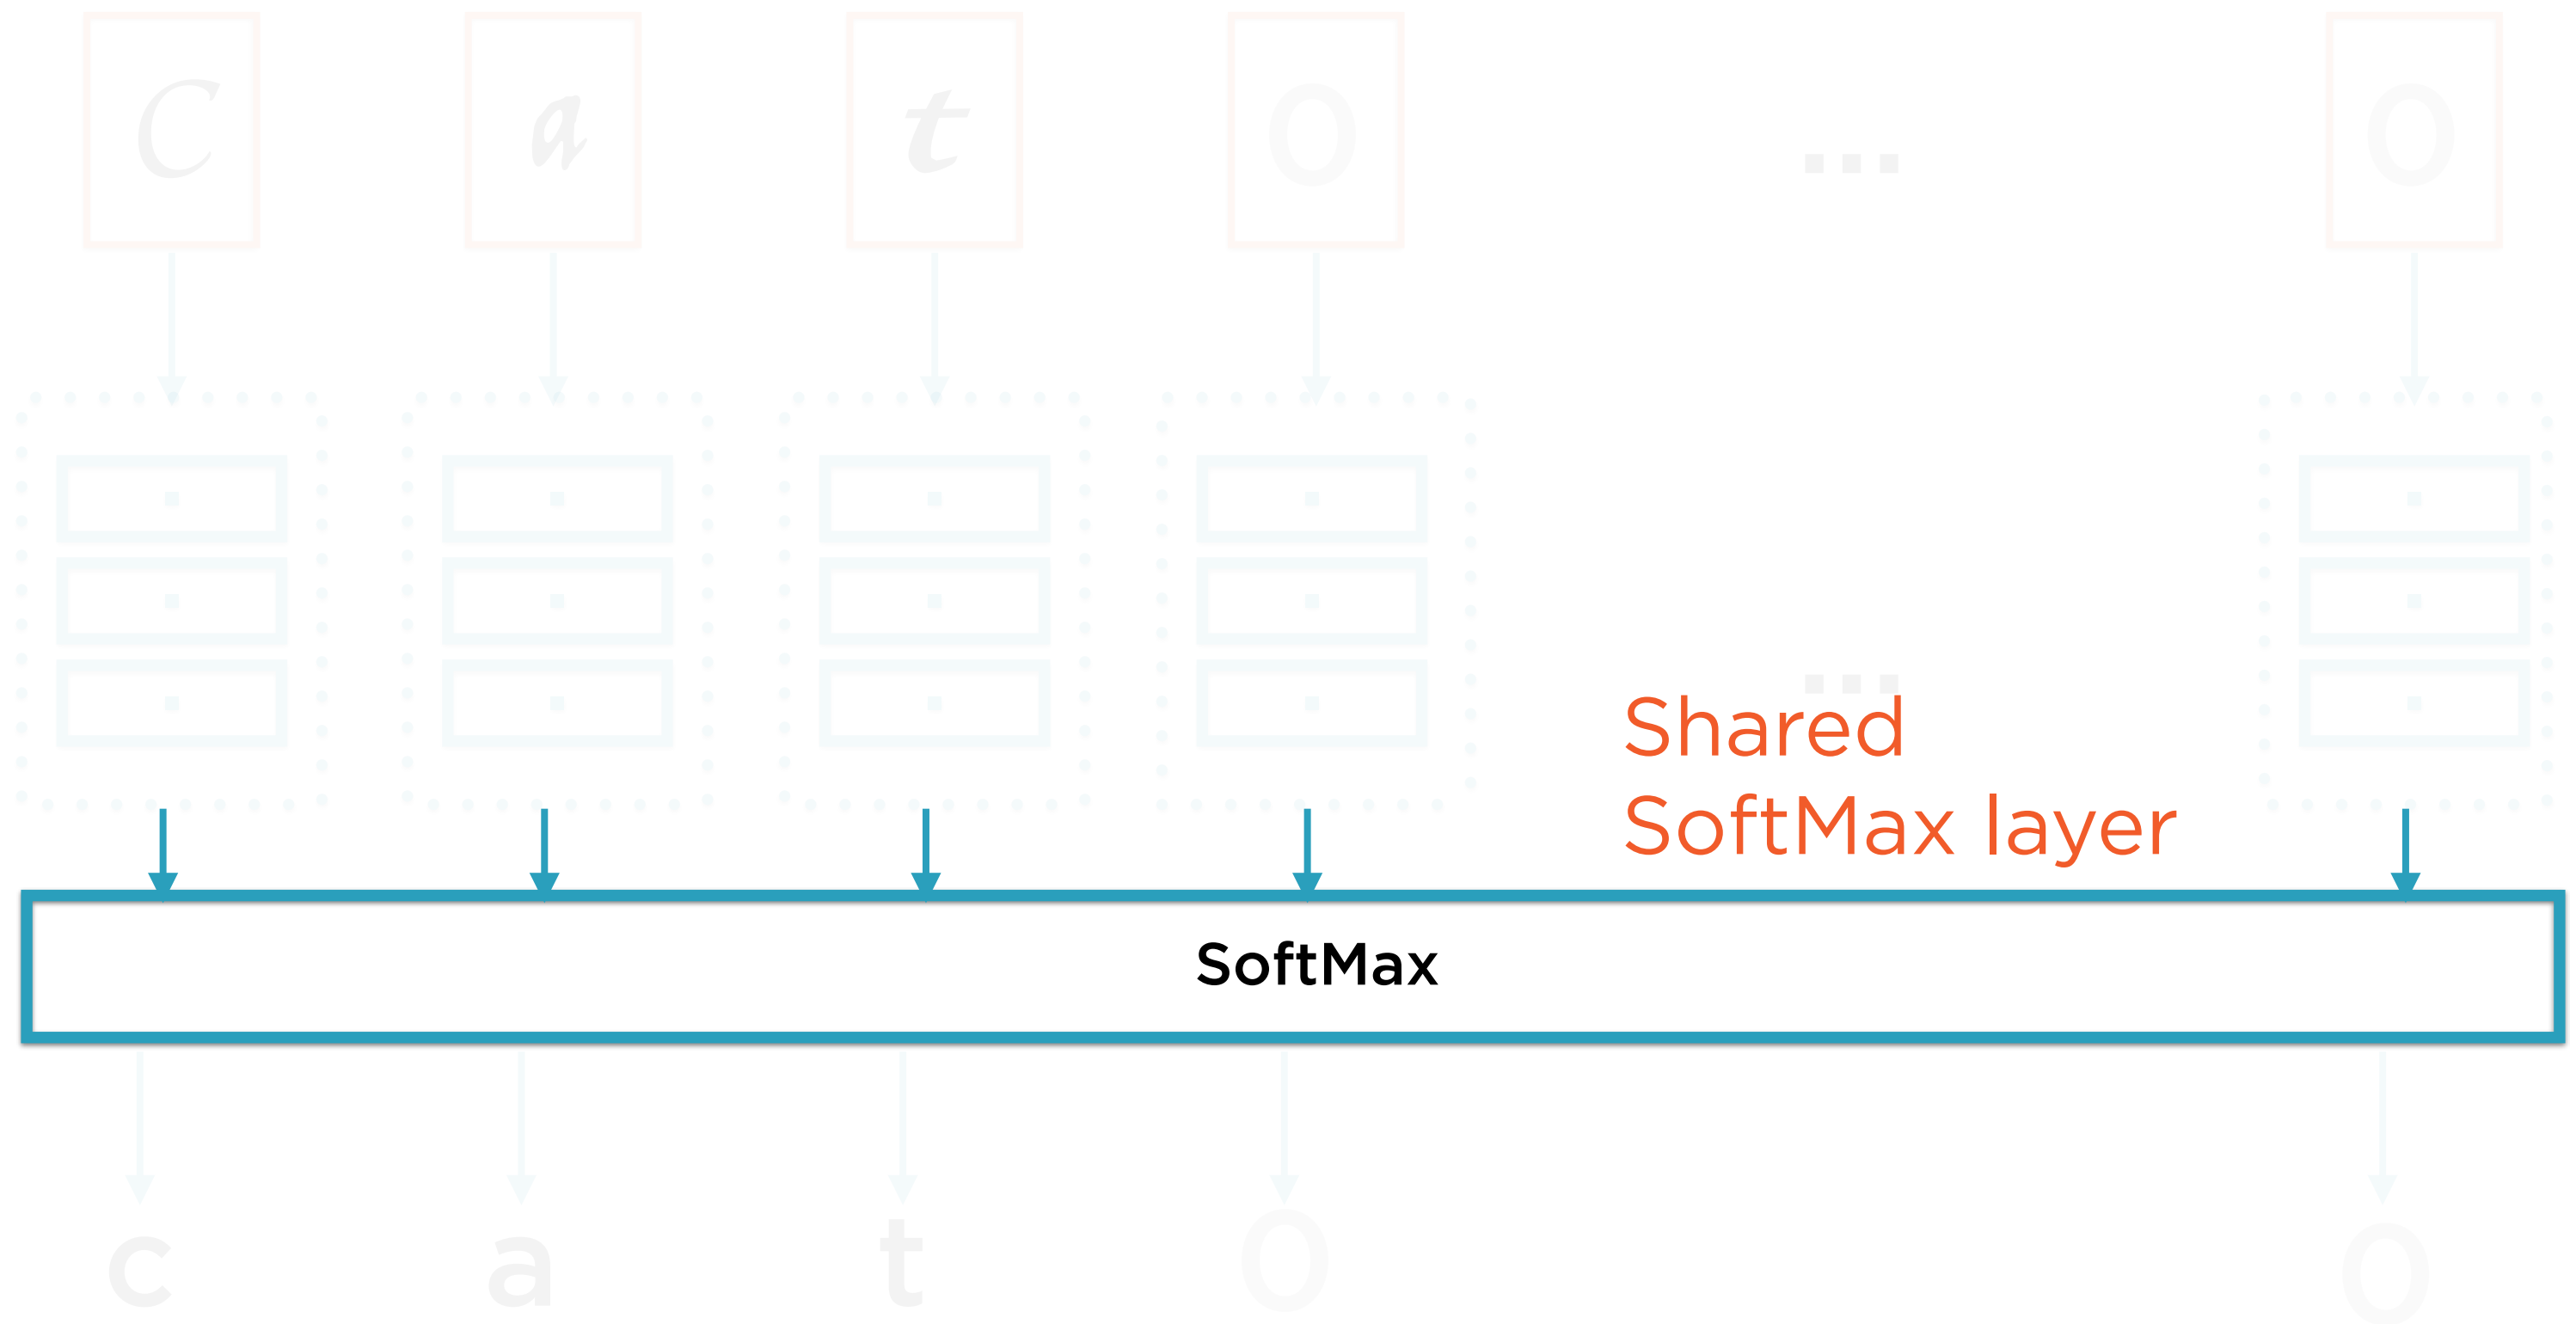
Output probabilities of 26 elements

Text Prediction

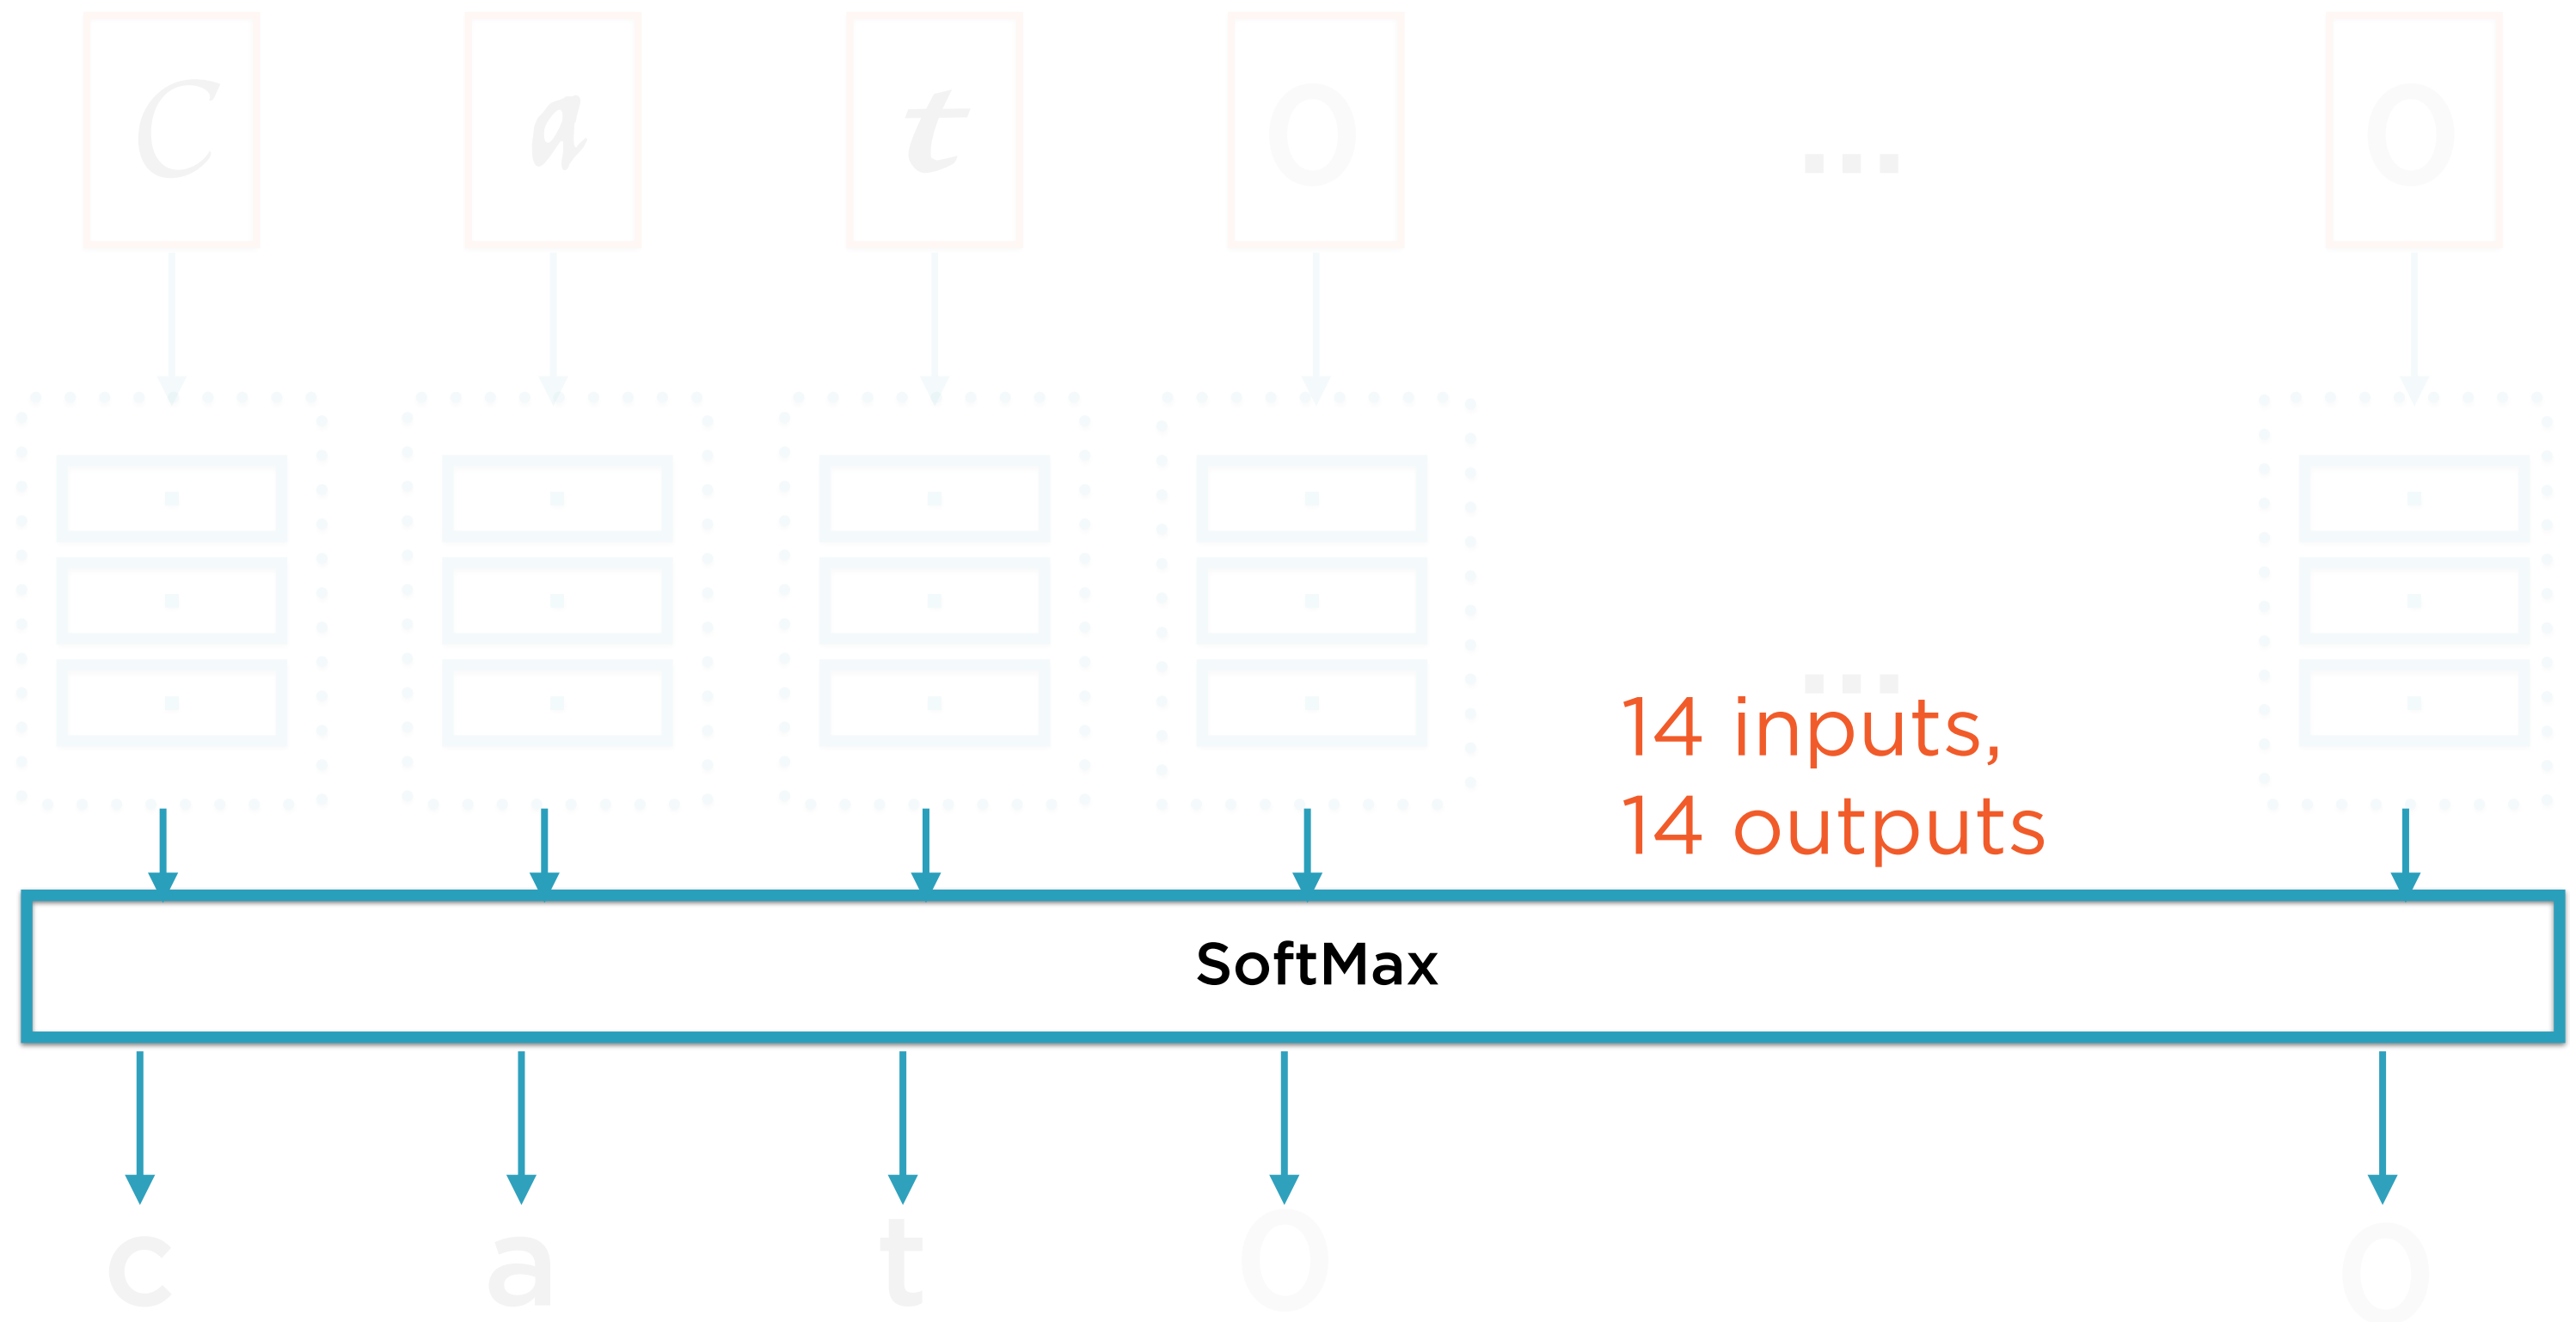
Shared Softmax layer

Output probabilities of 83 elements

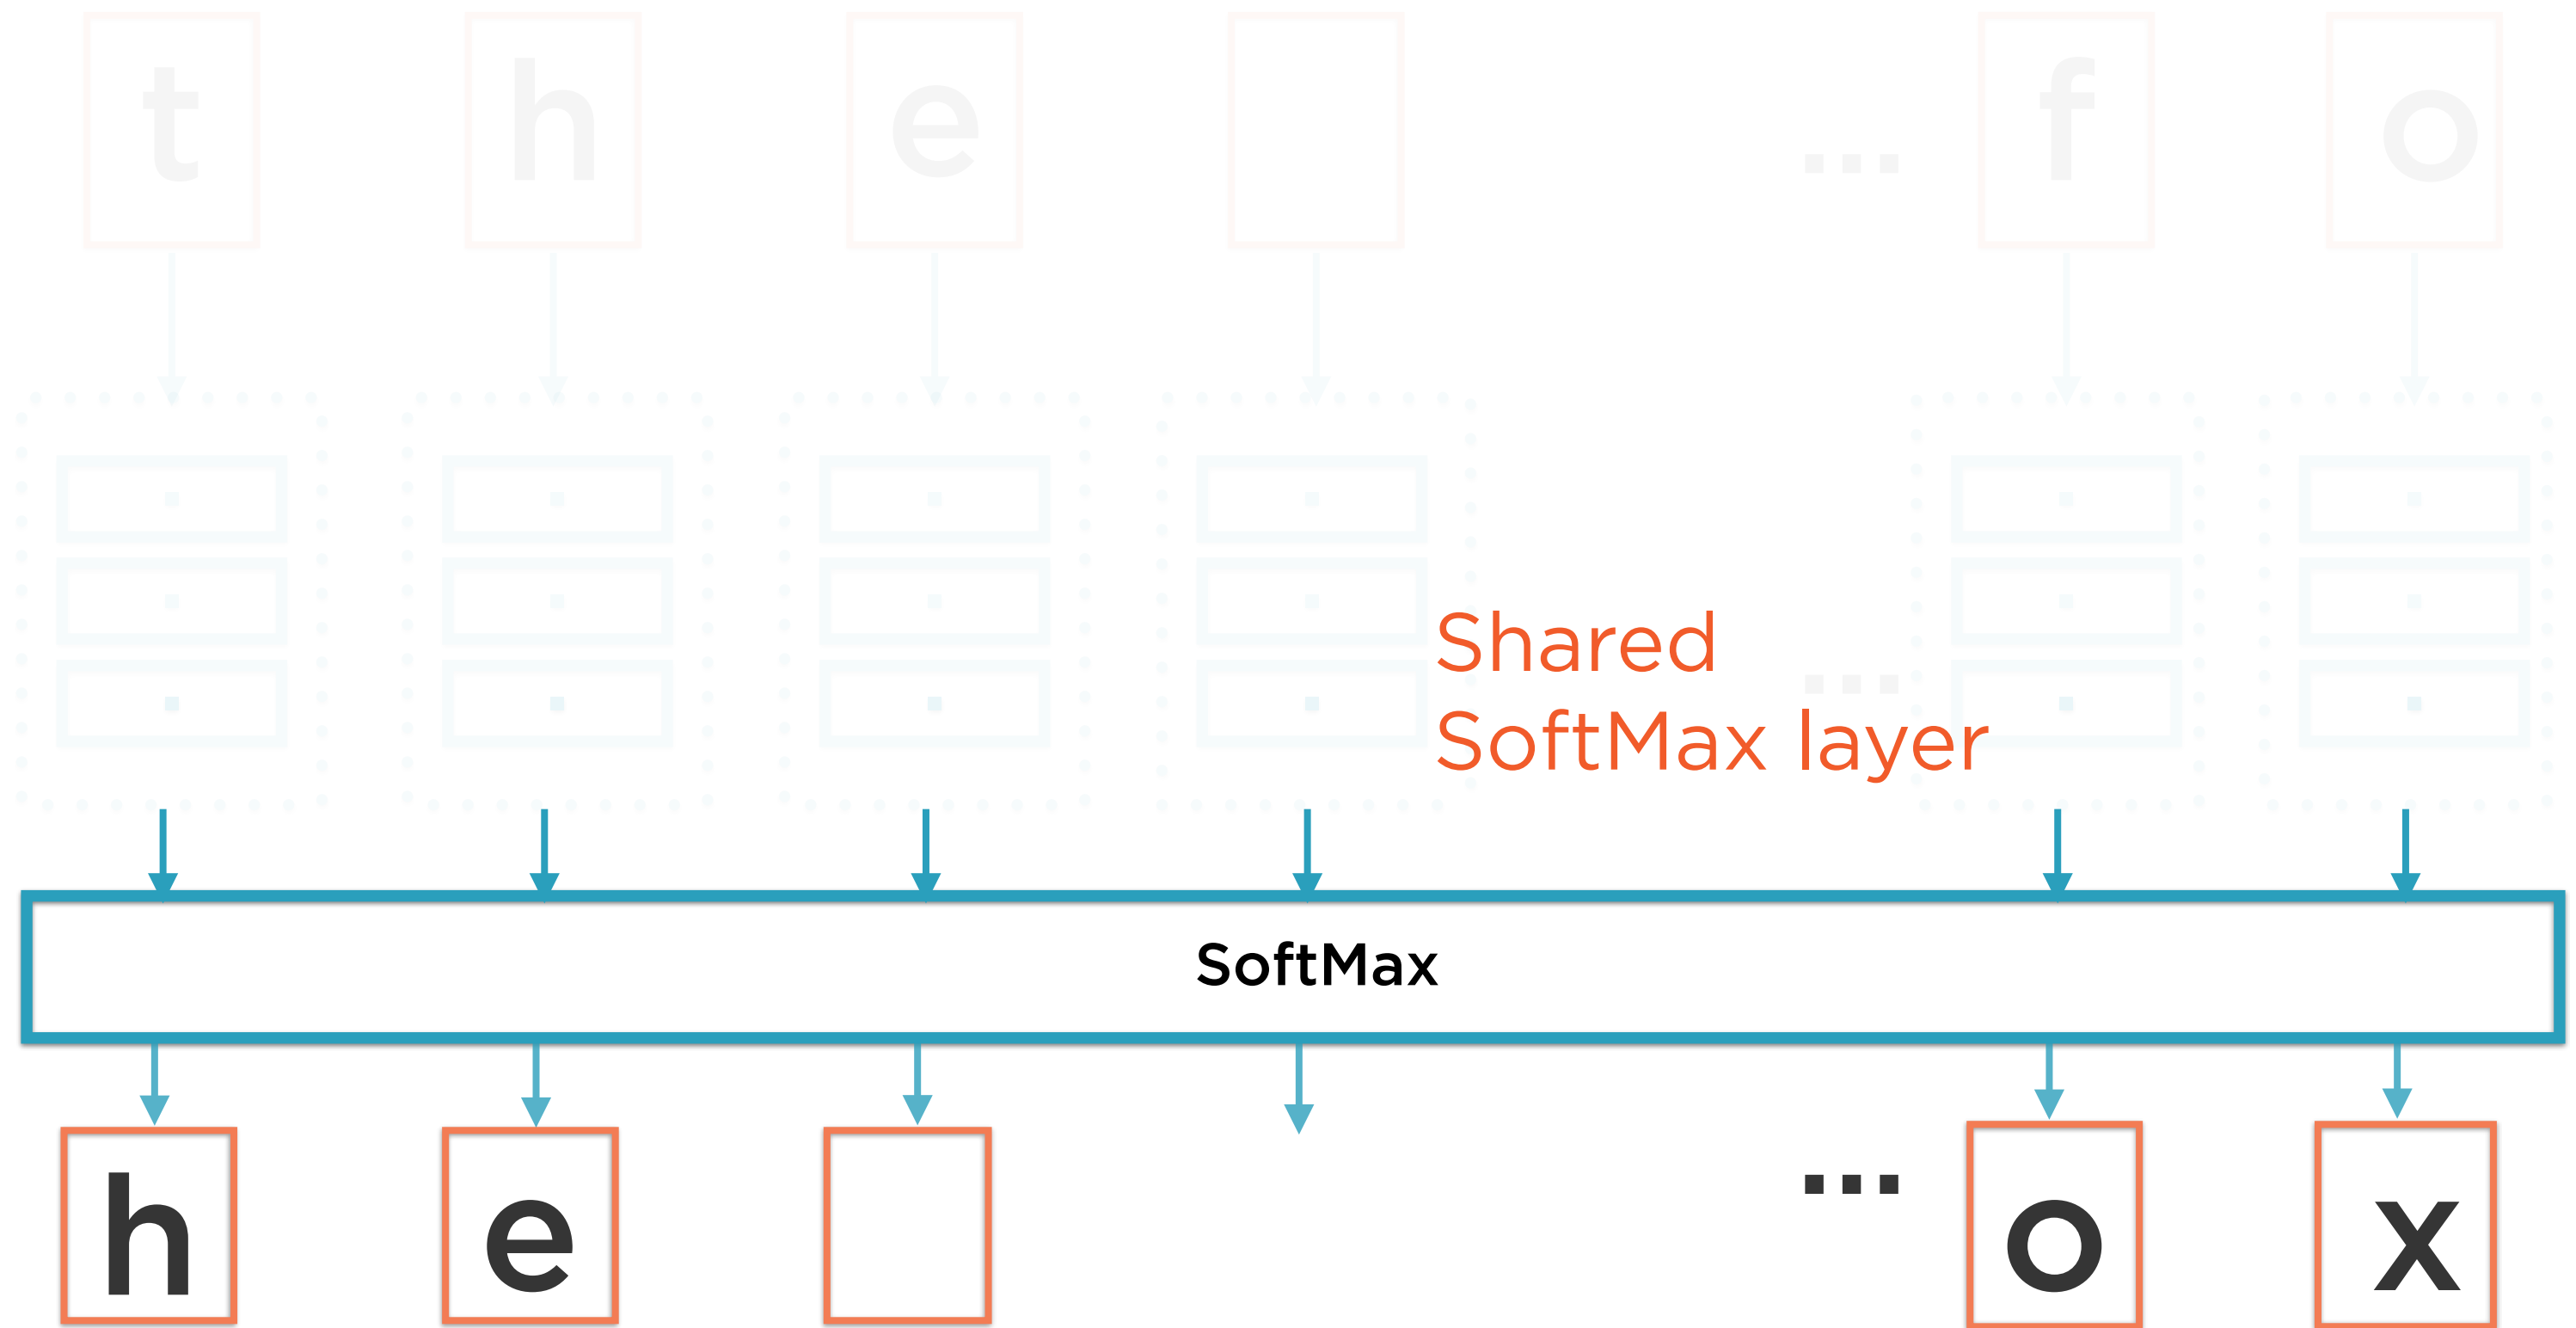
OCR: RNN Architecture



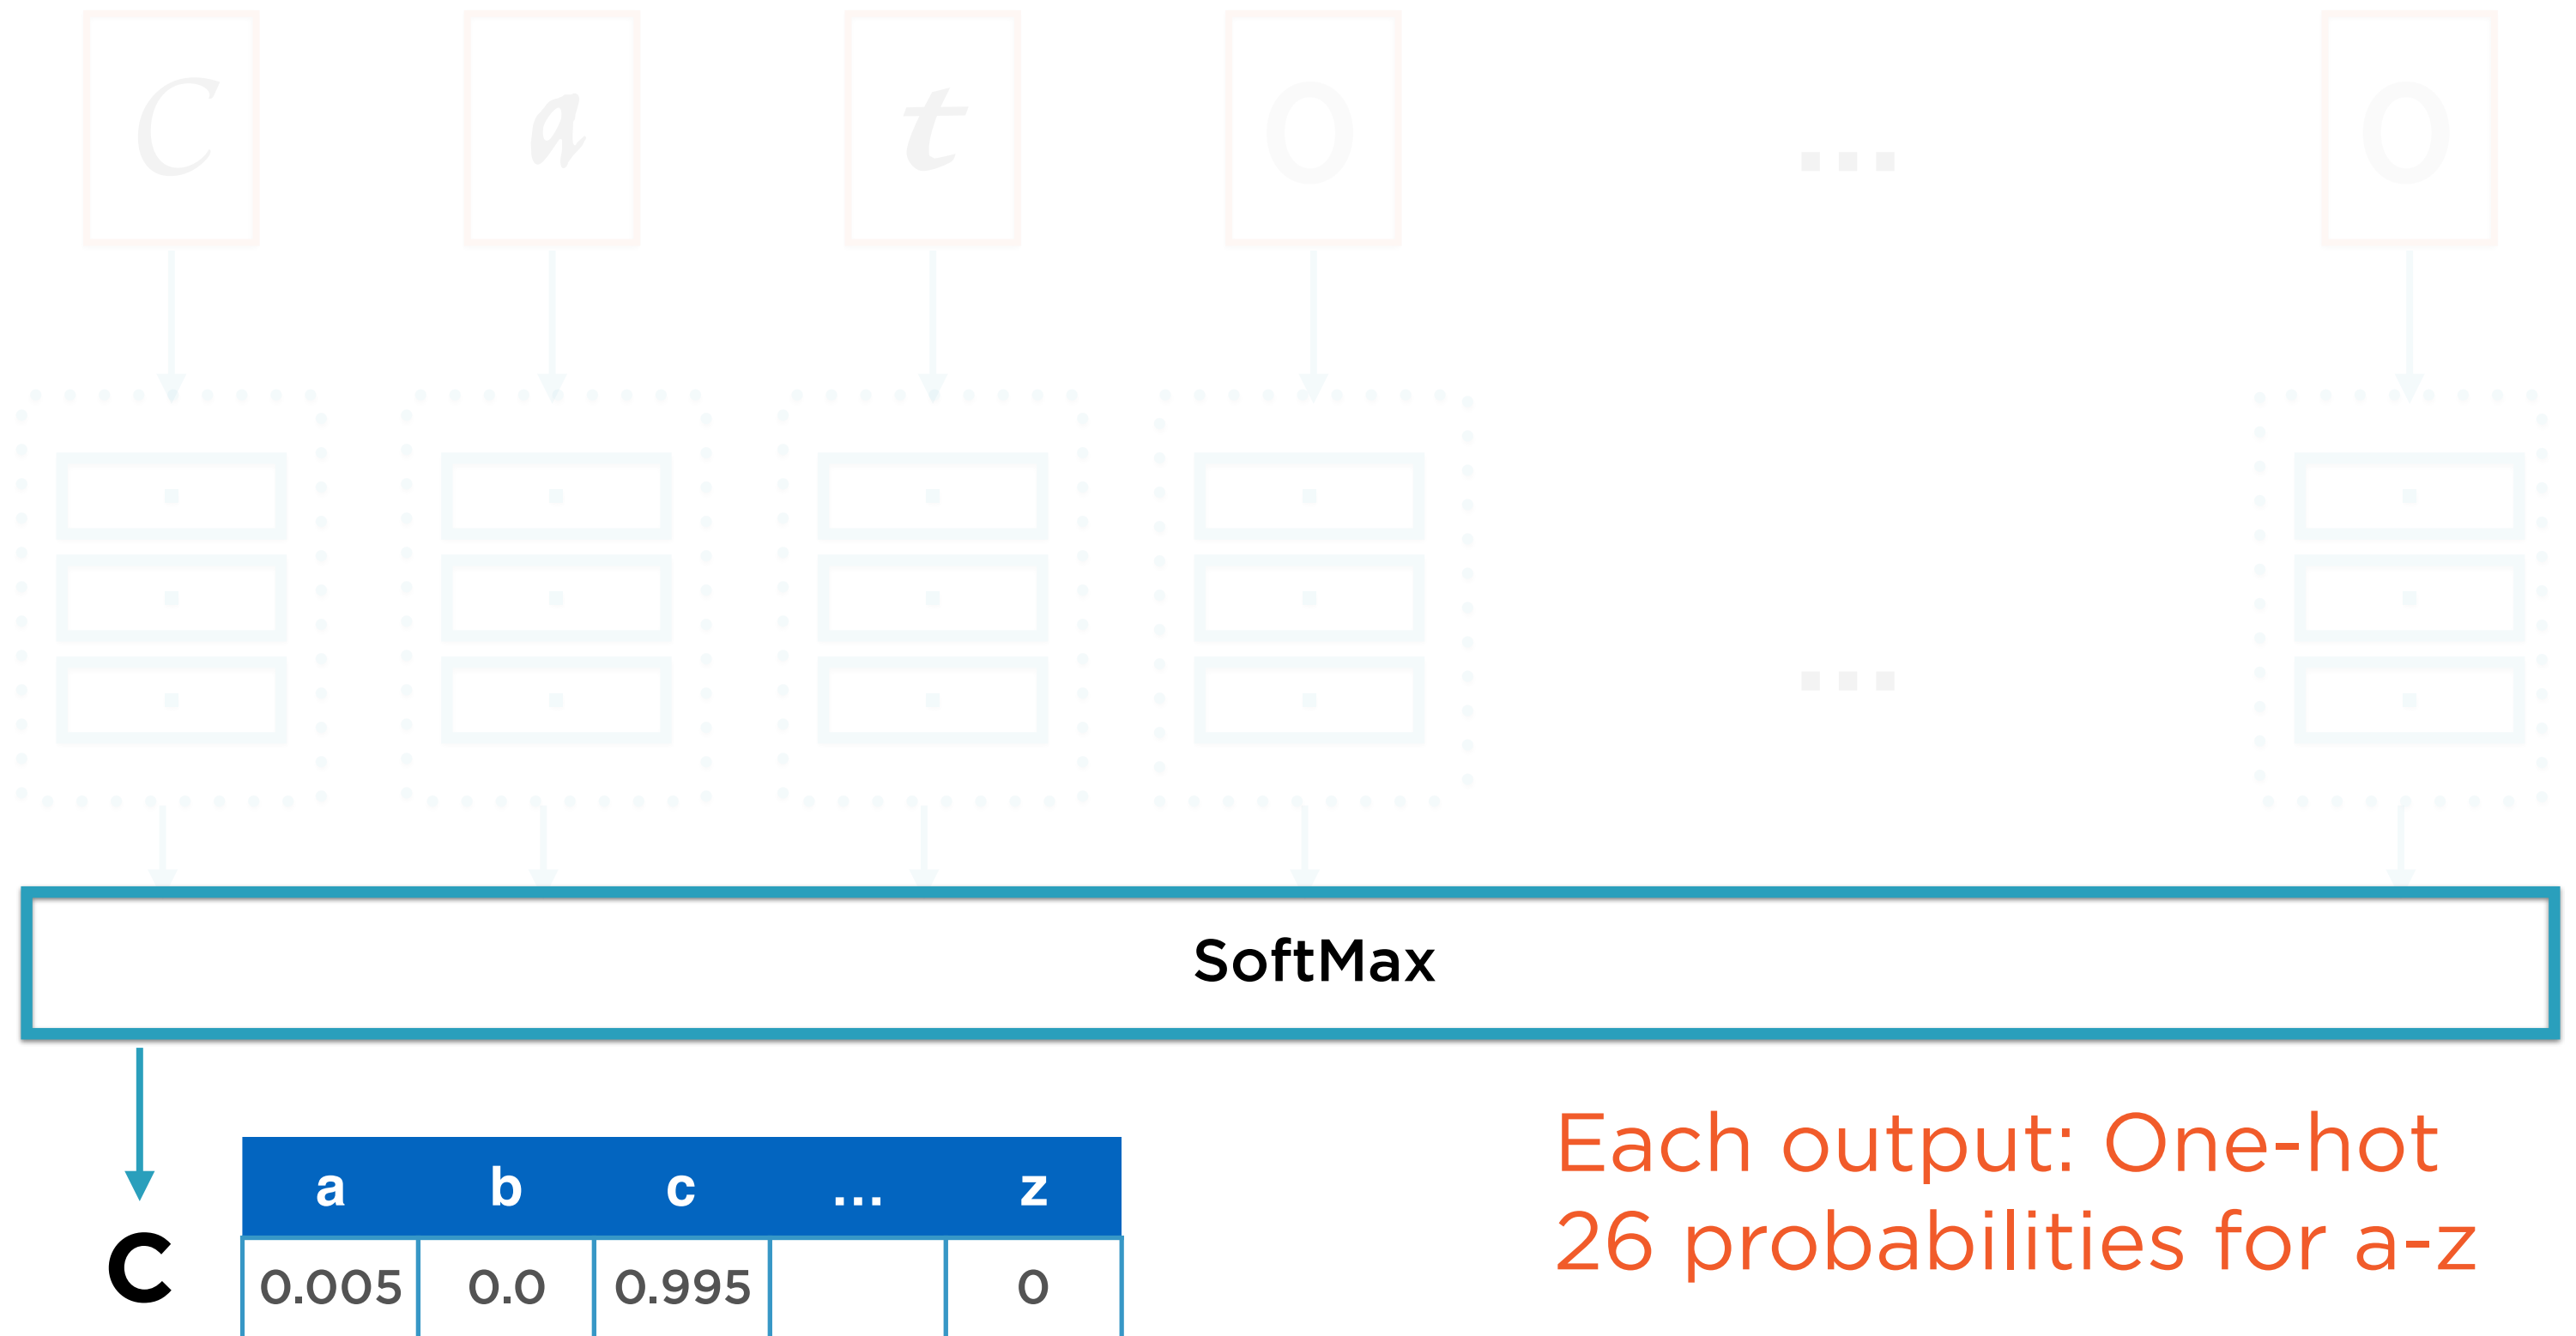
OCR: RNN Architecture




Text Prediction: RNN Architecture




OCR: RNN Architecture



Text Prediction: RNN Architecture



Text Prediction: RNN Architecture



Perplexity

Contrasting Architectures

OCR Classification

Cross-entropy as cost function

Accuracy as evaluation metric


Text Prediction

Cross-entropy as cost function


Perplexity as evaluation metric

The actual training of a neural network happens via Gradient Descent Optimization

Minimizing Cross-entropy




Minimizing Cross-entropy




Intuition: Low Cross-entropy


Y_{actual}



$Y_{\text{predicted}}$




Intuition: Low Cross-entropy



The labels of the two series are in-synch

Intuition: Low Cross-entropy




$-\text{Sum}(P(Y_{\text{actual}}) * \log [P(Y_{\text{predicted}})])$ will be small


Cross-entropy

Intuition: High Cross-entropy


Y_{actual}



$Y_{\text{predicted}}$




Intuition: High Cross-entropy



The labels of the two series are out-of-synch

Intuition: High Cross-entropy



$-\text{Sum}(P(Y_{\text{actual}}) * \log [P(Y_{\text{predicted}})])$ will be large

Cross-entropy

Perplexity

In information theory, perplexity is a measurement of how well a probability model predicts a sample. A low perplexity indicates the model is good at predicting the sample

Perplexity

In information theory, perplexity is a measurement of how well a probability model predicts a sample. A low perplexity indicates the model is good at predicting the sample

Perplexity

In information theory, perplexity is a measurement of how well a probability model predicts a sample. A low perplexity indicates the model is good at predicting the sample

Training a Model

Fact

Characters in our universe are drawn from distribution p

Training samples are drawn from probability distribution p

Model estimates some probability distribution q

Training process minimises cross-entropy between p and q

Explanation

This distribution p is unknown

Use these to train a model

q should be as close to p as possible

Cross-entropy is a measure of distance between two distributions

Evaluating Prediction

Fact

Draw test sample $x_1, x_2, x_3 \dots x_N$ from corpus

Use the model to predict what follows

Calculate perplexity to evaluate how well the predictor did

Explanation

The test sample follows distribution p

Model will predict using distribution q

Perplexity is just $2^{\text{cross-entropy}}$

Perplexity

$$\text{Cross-entropy} = - \sum (P(Y_{\text{actual}}) * \log [P(Y_{\text{predicted}})])$$

$$\text{Perplexity} = 2^{\text{cross-entropy}}$$

A perfect model has cross-entropy of 0, and
perplexity of 1

Perplexity

$$\text{Cross-entropy} = - \sum (P(Y_{\text{actual}}) * \log [P(Y_{\text{predicted}})])$$

$$\text{Perplexity} = 2^{\text{cross-entropy}}$$

Perplexity captures how many different options the model has to choose between

The more choices - the more perplexed (confused) the model is

A perfect model has cross-entropy
of 0, and perplexity of 1

Summary

Text prediction is one of several classic language modeling problems

Multi-RNNs are a specific type of RNN that work well in language modeling

A key technique is smartly re-initialising state of the RNN during prediction

An evaluation metric called perplexity is used to assess predictive performance